

SIP Servlet Specification, version 1.1

JSR 289 Expert Group

Specification Lead:

Mihir Kulkarni and Yannis Cosmadopoulos
Oracle Corporation

Please send comments to jsr-289-comments@jcp.org

License

BEA SYSTEMS, INC. ("BEA") IS WILLING TO LICENSE THIS SPECIFICATION TO YOU ONLY UPON THE CONDITION THAT YOU ACCEPT ALL OF THE TERMS CONTAINED IN THIS LICENSE AGREEMENT ("AGREEMENT"). PLEASE READ THE TERMS AND CONDITIONS OF THIS AGREEMENT CAREFULLY. BY DOWNLOADING THIS SPECIFICATION, YOU ACCEPT THE TERMS AND CONDITIONS OF THE AGREEMENT. IF YOU ARE NOT WILLING TO BE BOUND BY IT, SELECT THE "DECLINE" BUTTON AT THE BOTTOM OF THIS PAGE AND THE DOWNLOADING PROCESS WILL NOT CONTINUE.

JSR 289 Specification ("Specification")

Version: 1.1

Status: Final Release

Release: August 1, 2008

Copyright © 2008 BEA Systems, Inc.

2315 North First Street, San Jose, CA 95131, U.S.A.

All rights reserved.

NOTICE; LIMITED LICENSE GRANTS

1. License for Evaluation Purposes. BEA hereby grants you a fully-paid, non-exclusive, non-transferable, worldwide, limited license (without the right to sublicense), under BEA's applicable intellectual property rights to view, download, use and reproduce the Specification only for the purpose of internal evaluation, which shall be understood to include developing applications intended to run on an implementation of the Specification provided that such applications do not themselves implement any portion(s) of the Specification.

2. License for the Distribution of Compliant Implementations. BEA also grants you a perpetual, non-exclusive, non-transferable, worldwide, fully paid-up, royalty free, limited license (without the right to sublicense) under any applicable copyrights or, subject to the provisions of subsection 3 below, patent rights it may have covering the Specification to create and/or distribute an implementation of the Specification that: (a) fully implements the Specification including all its required interfaces and functionality, and (b) passes the Technology Compatibility Kit for such Specification ("Compliant Implementation").

3. Reciprocity Concerning Patent Licenses.

a. With respect to any patent claims covered by the license granted under subparagraph 2 above that would be infringed by all technically feasible implementations of the Specification, such license is conditioned upon your offering on fair, reasonable and non-discriminatory terms, to any party seeking it from You, a perpetual, non-exclusive, non-transferable, worldwide license under Your patent rights which are or would be infringed by all technically feasible implementations of the Specification to develop, distribute and use a Compliant Implementation.

b With respect to any patent claims owned by BEA and covered by the license granted under subparagraph 2, whether or not their infringement can be avoided in a technically feasible manner when implementing the Specification, such license shall terminate with respect to such claims if You initiate a claim against BEA that it has, in the course of performing its responsibilities as the Specification Lead, induced any other entity to infringe Your patent rights.

c Also with respect to any patent claims owned by BEA and covered by the license granted under subparagraph, where the infringement of such claims can be avoided in a technically feasible manner when implementing the Specification such license, with respect to such claims, shall terminate if You initiate a claim against BEA that its making, having made, using, offering to sell, selling or importing a Compliant Implementation infringes Your patent rights.

4. Definitions. For the purposes of this Agreement: “Technology Compatibility Kit” or “TCK” shall mean the test suite and accompanying documentation provided by BEA which corresponds to the particular version of the Specification being tested.

BEA shall have the right to terminate this Agreement immediately notice if you fail to comply with any material provision of or act outside the scope of the licenses granted above.

TRADEMARKS

No right, title, or interest in or to any trademarks, service marks, or trade names of BEA or BEA's licensors is granted hereunder. Java is a registered trademark of Sun Microsystems, Inc. in the United States and other countries.

DISCLAIMER OF WARRANTIES

THE SPECIFICATION IS PROVIDED "AS IS". BEA MAKES NO REPRESENTATIONS OR WARRANTIES, EITHER EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO, WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT (INCLUDING AS A CONSEQUENCE OF ANY PRACTICE OR IMPLEMENTATION OF THE SPECIFICATION), OR THAT THE CONTENTS OF THE SPECIFICATION ARE SUITABLE FOR ANY PURPOSE.

THE SPECIFICATION COULD INCLUDE TECHNICAL INACCURACIES OR TYPOGRAPHICAL ERRORS. CHANGES ARE PERIODICALLY ADDED TO THE INFORMATION THEREIN; THESE CHANGES WILL BE INCORPORATED INTO NEW VERSIONS OF THE SPECIFICATION, IF ANY. BEA MAY MAKE IMPROVEMENTS AND/OR CHANGES TO THE PRODUCT(S) AND/OR THE PROGRAM(S) DESCRIBED IN THE SPECIFICATION AT ANY TIME. Any use of such changes in the Specification will be governed by the then-current license for the applicable version of the Specification.

LIMITATION OF LIABILITY

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL BEA OR ITS LICENSORS BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION, LOST REVENUE, PROFITS OR DATA, OR FOR SPECIAL, INDIRECT, CONSEQUENTIAL, INCIDENTAL OR PUNITIVE DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF OR RELATED TO ANY FURNISHING, PRACTICING, MODIFYING OR ANY USE OF THE SPECIFICATION, EVEN IF BEA AND/OR ITS LICENSORS HAVE BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

You will indemnify, hold harmless, and defend BEA and its licensors from any claims arising or resulting from: (i) your use of the Specification; (ii) the use or distribution of your application or applet written to and/or Your implementation of the Specification; and/or (iii) any claims that later versions or releases of any Specification furnished to you are incompatible with the Specification provided to you under this license.

RESTRICTED RIGHTS LEGEND

U.S. Government: If this Specification is being acquired by or on behalf of the U.S. Government or by a U.S. Government prime contractor or subcontractor (at any tier), then the Government's rights in the Software and accompanying documentation shall be only as set forth in this license; this is in accordance with 48 C.F.R. 227.7201 through 227.7202-4 (for Department of Defense (DoD) acquisitions) and with 48 C.F.R. 2.101 and 12.212 (for non-DoD acquisitions).

REPORT

You may wish to report any ambiguities, inconsistencies or inaccuracies you may find in connection with your use of the Specification ("Feedback"). To the extent that you provide BEA with any Feedback, you hereby: (i) agree that such Feedback is provided on a non-proprietary and non-confidential basis, and (ii) grant BEA a perpetual, non-exclusive, worldwide, fully paid-up, irrevocable copyright license, with the right to sublicense through multiple levels of sublicensees, to incorporate, disclose, and use without limitation the Feedback for any purpose related to the Specification and future versions, implementations, and test suites thereof.

Contents

Preface

Overview

Goals of the SIP Servlet API	1-1
What is a SIP Servlet?	1-2
What is a SIP Servlet Container?	1-2
Relationship with the HTTP Servlet API	1-3
Differences from the HTTP Servlet API	1-3
Use of Servlet.service()	1-4
Asynchronicity	1-4
Application Composition	1-5
Converged Applications	1-5
Deployment Models	1-5
Standalone SIP Servlet Container	1-5
SIP and HTTP Converged Servlet Container	1-5
SIP and Java EE Convergence	1-6
Examples	1-6
A Location Service	1-6
A Multi-Protocol Servlet Application	1-7
A Java EE Converged Application	1-10

The Servlet Interface

Servlet Life Cycle	2-1
------------------------------	-----

Servlet Life Cycle Listener	2-2
Processing SIP Messages	2-3
SIP Specific Request Handling Methods	2-3
Receiving Requests	2-5
SIP Specific Response Handling Methods	2-6
Number of Instances	2-6

Servlet Context

The SipFactory	3-1
Extensions Supported	3-2
RFCs Supported	3-2
Context Path	3-2
Context Parameters	3-3

Addressing

The Address Interface	4-1
The Parameterable Interface	4-2
The From and To Header Fields.	4-3
The Contact Header Field	4-4
Transport Parameters in Contact Header	4-6
URIs	4-6
SipURI	4-6
TelURL.	4-7

Requests and Responses

The SipServletMessage Interface	5-1
Implicit Transaction State	5-2
Access to Message Content.	5-3
Character Encoding	5-4

Headers	5-4
Parameterable and Address Header Fields	5-5
System Headers	5-6
TLS Attributes	5-6
Transport Level Information	5-7
Requests	5-7
Parameters	5-7
Attributes	5-8
Popped Route Header	5-9
Responses	5-10
Reliable Provisional Responses	5-10
Buffering	5-11
Accessibility of SIP Servlet Messages	5-11
Internationalization	5-12
Indicating Preferred Language	5-12
Indicating Language of Message Content	5-13

Sessions

SipApplicationSession	6-2
Protocol Sessions	6-2
SipApplicationSession Lifetime	6-3
SipApplicationSession Timer Operation and SipApplicationSession Expiration	6-4
SipApplicationSession Invalidation	6-4
Binding Attributes into a SipApplicationSession	6-7
SipSession	6-7
Relationship to SIP Dialogs	6-7
Cancel Message Processing	6-11
Maintaining Dialog State in the SipSession	6-11

When in a Dialog	6-11
When in the INITIAL SipSession State	6-11
Creation of SipSessions	6-12
Extensions Creating Dialogs.	6-13
Derived SipSessions	6-13
SipSession Lifetime	6-14
SipSession Invalidation.	6-14
Important Semantics	6-18
The RequestDispatcher Interface	6-19
The SipSession Handler	6-19
Binding Attributes into a SipSession	6-19
Last Accessed Times	6-20
Important Session Semantics	6-20
Message Context	6-20
Threading Issues	6-21
Distributed Environments	6-21

SIP Servlet Applications

Relationship with HTTP Servlet Applications	7-1
Relationship to ServletContext	7-1
Elements of a SIP Application	7-2
Deployment Hierarchies	7-2
Directory Structure	7-2
Application Names	7-3
Example Application Directory Structure	7-4
Servlet Application Archive File.	7-5
SIP Application Configuration Descriptor	7-5
Dependencies On Extensions	7-5

Servlet Application Classloader	7-6
Replacing a Servlet Application	7-6
Servlet Application Environment	7-7

Application Listeners and Events

SIP Servlet Event Types and Listener Interfaces	8-1
---	-----

Timer Service

TimerService	9-1
ServletTimer	9-2
TimerListener	9-3

Proxying

Parameters	10-1
Operation	10-3
Proxy Branches	10-4
Pushing Route headers	10-6
Sending Responses	10-6
Receiving Responses	10-7
Handling 2xx Responses to INVITE	10-8
Correlating responses to proxy branches	10-9
Sequential Search Timeout	10-10
Handling 3xx responses	10-10
Sending CANCEL	10-11
Receiving CANCEL	10-11
Sending ACK	10-12
Receiving ACK	10-12
Handling Subsequent Requests	10-12
Max-Forwards Check	10-13

Proxying and Sessions.	10-13
Record-Route Parameters	10-14
Path Header and Path Parameters	10-15

Acting as a User Agent

Client Functions	11-1
Creating Initial Requests	11-1
Copying From and To Addresses	11-3
Creating Subsequent Requests	11-3
Pushing Route Header Field Values.	11-3
Sending a Request as a UAC	11-4
Receiving Responses	11-4
Handling Multiple Dialogs	11-5
Transaction Timeout.	11-5
Sending ACK	11-5
Sending PRACK	11-6
Sending CANCEL	11-6
Server Functions	11-6
Sending Responses.	11-6
Receiving ACK	11-7
Receiving CANCEL	11-7
Handling Error Conditions.	11-7
Receiving unimplemented subsequent requests	11-7
Handling Pending Invites	11-8

Back To Back User Agents

B2BUA Helper Class	12-1
Creating new B2BUA Request	12-2

Linked SipSessions And Linked Requests	12-3
Explicit Session Linkage	12-3
Implicit Session Linkage	12-5
Access to Un-Committed Messages	12-5
Original Request and Session Cloning	12-6
Cloning and Linking.	12-6

Converged Container and Applications

Converged Application	13-1
SIP and Java EE Converged Application	13-2
SIP and HTTP Converged Application	13-2
Accessing SIP Factory	13-3
Accessing SipApplicationSession By ID.	13-4
Encoding URLs	13-5
Association of HTTP Session With SipApplicationSession	13-5
Finding Parent SipApplicationsSession From A Protocol Session	13-6
Encoding HTTP URLs With HTTP Sessions In Non HTTP Scope	13-6

Container Functions

Division of Responsibility	14-1
Protocol Compliance	14-2
Multihomed Host Support	14-3
Application Composition on Multihomed Container.	14-4

Application Selection And Composition Model

Application Selection Process	15-2
The Role of the Application Router	15-2
The Role of Applications	15-3
The Role of the Container	15-3

Application Independence	15-4
Subscriber Identity and Routing Regions	15-5
Routing Directives	15-7
Application Environment and Behaviour	15-7
Receiving an Initial Request	15-7
Sending an Initial Request	15-7
Application Router Behavior	15-11
Order of Routing Regions	15-13
Inter-Container Application Routing	15-14
Container Behavior	15-14
Procedure for Routing an Initial Request.	15-15
Application Router Packaging and Deployment	15-17
Application Names	15-17
Responses, Subsequent Requests and Application Path	15-18
Transport Information	15-20
Popping of Top Route Header.	15-21
Container Behavior	15-21
Top Route Processing Examples	15-21
Request with two route headers arriving at the container.	15-21
Application pushes route pointing back at the container	15-22
Application pushes external route.	15-23
Examples	15-23
Example with Two Applications	15-24
Simple call with no modifications of requests	15-26
Modification of Headers	15-27
Reversing the Direction of the Call	15-28
Initiating a New Request	15-29
Loop Detection	15-29

Session Targeting.	15-29
Session Targeting and Application Selection.	15-30
Session Key Based Targeting Mechanism	15-31
The Encode URI Mechanism	15-33
Join and Replaces Targeting Mechanism	15-34
Resolving Session Targeting Conflicts.	15-36

Mapping Requests To Servlets

Multiple Servlets	16-1
Servlet Selection	16-1
Compatibility with v1.0 Specification	16-2

Security

Introduction	17-1
Declarative Security.	17-2
Programmatic Security	17-2
Roles	17-3
Authentication	17-4
Server Tracking of Authentication Information	17-4
Propagation of Security Identity in EJBTM Calls.	17-5
Specifying Security Constraints	17-5
Default Policies	17-7
Authentication of Servlet Initiated Requests	17-8
Description of the Procedure	17-9

Java Enterprise Edition 5 Container

Java 5	18-1
Annotations and Resource Injection	18-1
Servlet 2.5 alignment	18-2

@SipServlet Annotation	18-4
@SipApplication Annotation.	18-6
@SipListener Annotation.	18-9
@SipApplicationKey Annotation	18-10
Annotation for SipFactory Injection	18-11
Annotation for SipSessionsUtil Injection	18-11
Annotation for TimerService Injection	18-12
Annotation Parsing	18-13

Deployment Descriptor

Differences from the HTTP Servlet Deployment Descriptor	19-1
Converged SIP and HTTP Applications	19-2
Deployment Descriptor Elements	19-2
Rules for Processing the Deployment Descriptor	19-3
The SIP Servlet XSD.	19-4

Changes since v1.0

Backward Compatibility considerations	A-4
Changes in the API since v1.0	A-5

Definition of Initial Request

Retried Requests	B-2
REGISTER Requests	B-3

Default Application Router

The DAR Configuration File.	C-1
The DAR Operation	C-2

References

Glossary

Preface

This specification defines version 1.1 of the SIP Servlet API. The specification requires J2SE 5.0 and support for SIP as defined in [RFC 3261].

The following IETF specifications referenced in this specification provide information relevant to the development and implementation of the SIP Servlet API and SIP Servlet containers: [RFC 3265], [RFC 3262], [RFC 2976], [RFC 3311], [RFC 3515], [RFC 3903], [RFC 3841], [RFC 3966], [RFC 3327], [RFC 3725], [RFC 3856], [RFC 4028], [RFC 3326]

Online versions of these RFCs are available at <http://www.ietf.org/rfc>.

Who should read this Document

The intended audience for this specification includes the following groups:

- SIP application server vendors who want to provide servlet engines that conform to this standard.
- SIP servlet application developers.

Familiarity with SIP is assumed throughout.

Terminology

In this document, the key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" are to be interpreted as described in BCP 14, RFC 2119 [2] and indicate requirement levels for compliant JSR 289 implementations.

Typographic Conventions

Java classes, method names, parameters, literal values, and fragments of code are printed in constant width as are XML elements and documents.

Providing Feedback

Feedback on this specification is welcome. Please e-mail your comments to jsr-289-comments@jcp.org

Acknowledgements

This specification was developed under the Java Community Process 2.6 as JSR 289. The JSR 289 expert group consists of the following members:

Hira Agrawal (Telcordia Technologies Inc.),
Bartosz Baranowski (Red Hat Middleware LLC),
Diego Besposvan (Mailvision LTD),
Chris Boulton (Avaya, Inc),
Matthias Britsch (T-Mobile International AG & Co. KG),
Erik Burckart (IBM),
Eric Cheung (AT&T),
David Connelly,
Christophe Dejouhanet (Orange France SA),
Jean Deruelle (Red Hat Middleware LLC),
Simon Dutkowski (Fraunhofer-Gesellschaft Institute FIRST),
Sreeram Duvur (Sun Microsystems, Inc.),
Tomas Ericson (Oracle),
Jarek Wilkiewicz (Hewlett-Packard),
Goran Eriksson (Ericsson AB),
Ian Evans (Avaya, Inc),
Anat Fradin (IBM),
Kristoffer Gronowski (Ericsson AB),

Robert Handl (Ericsson AB),
Avshalom Hourì (IBM),
Ivelin Ivanov (Red Hat Middleware LLC),
Akinori Iwakawa (Fujitsu Limited),
Hakan Jonsson (Appium Technologies AB),
Nasir Khan,
Yueh Lee (Cisco Systems),
Thomas Leseney (Nexcom Systems),
Johan Liseborn (Oracle),
Stephane H. Maes (Oracle),
Roger N. Mahler (Cingular Wireless),
Gines Gomez Martinez (Voztelecom Sistemas S.L.),
Remy Maucherat (Apache Software Foundation),
Francesco Moggia (Telecom Italia),
Phelim O'Doherty (Oracle),
Amir Perlman (IBM),
Marc Petit-Huguenin (8x8),
Per Pettersson (Ericsson AB),
Jean-Pierre Pinal (Nexcom Systems),
Kermit Hal Purdy (AT&T),
Uri Segev (IBM),
James Steadman (Oracle),
Prasad Subramanian (Sun Microsystems, Inc.),
Atul Varshneya,
Lebing Xie (Fraunhofer-Gesellschaft Institute FIRST),
Sami Samandi (Netcentrex)
Damian O'Neill (Aepona Technologies AB)

We would like to thank Anders Kristensen (Cisco Systems Inc.) and the members of the JSR 116 Expert Group who developed the Sip Servlet Specification v1.0 that contributed immensely to the development of the v1.1 specification. This specification builds on the Sip Servlet Specification v1.0 and for this reason, this specification copies text from the v1.0 specification in many places [SIP Servlet API].

Additionally, Rhys Ulerich of IBM, Paul Devine, Michael Palmeter, Vinod Mehra and Anno Langen of BEA Systems made enormous contributions to this specification. We would like to thank Paul Sydell of Oracle for his help with Java EE schema alignment, Adrian Kalaveshi of Oracle for his help with administrating the collaborative site used by the JSR 289 expert group and David Yozie of Oracle for his help in drafting this specification as a PDF.

An important goal of the SIP Servlet API is to build on the work of the well-established (HTTP) Servlet API. For this reason, this specification copies text from the Java Servlet Specification in many places [Servlet API]. The glossary includes a large number of entries copied from the SIP specification [RFC 3261].

This specification draws and adapts the the principles of application independence, composition, and selection from Distributed Feature Composition [DFC1998], [CN2004], [SE2005]. This specification makes use of Java annotations introduced in J2SE 5.0 and re-uses common annotations for Java platform [JSR250].

1 Overview

The Session Initiation Protocol (SIP) is used to establish, modify, and tear down IP multimedia sessions including IP telephony, presence, instant messaging besides other SIP applications. An important aspect of any communication infrastructure is programmability and the purpose of the SIP Servlet API is to standardize the platform for delivering SIP based services. The term platform is used here to include the Java API itself as well as standards covering the packaging and deployment of applications.

1.1 Goals of the SIP Servlet API

The following summarizes some important properties of the SIP Servlet API:

- **SIP signaling:** It allows applications to perform a fairly complete set of SIP signaling actions, including support for acting as user agent client (UAC), user agent server (UAS), and proxy.
- **Simplicity:** Containers handle “non-essential” complexity such as managing network listen points, retransmissions, CSeq, Call-ID and Via headers, routes, etc.
- **Converged applications:** It is possible for containers to support converged applications, that is, applications that span multiple protocols and interfaces, for example, Web, telephony, and other Java EE interfaces.
- **Third party application development:** The servlet model supports third party application development. An XML deployment descriptor is used to communicate application information from the application developer to deployers.

- **Application composition:** It is possible for several applications to execute on the same incoming or outgoing request or response. Each application has its own set of rules and executes independently of other applications in a well-defined and orderly fashion.
- **Carrier grade:** Servlets store application data in container managed session objects. Implementations may persist and/or replicate this data to achieve high availability.

1.2 What is a SIP Servlet?

A SIP servlet is a Java-based application component which is managed by a SIP servlet container and which performs SIP signaling. Like other Java-based components, servlets are platform independent Java classes that are compiled to platform neutral bytecode that can be loaded dynamically into and run by a Java-enabled SIP application server. Containers, sometimes called servlet engines, are server extensions that provide servlet functionality. Servlets interact with (SIP) clients by exchanging request and response messages through the servlet container.

1.3 What is a SIP Servlet Container?

The servlet container is a part of an application server that provides the network services over which requests and responses are received and sent. It decides which applications to invoke and in what order. A servlet container also contains and manages servlets through their lifecycle.

A servlet container can be built into a host SIP server, or installed as an add-on component to a SIP Server via that server's native extension API. Servlet containers can also be built into or possibly installed into servlet-enabled application servers.

A SIP servlet container manages the network listen points on which it listens for incoming SIP traffic (a listen point being a combination of transport protocol, IP address and port number). The SIP specification requires all SIP elements to support both UDP and TCP, and optionally TLS, SCTP, and potentially other transports.

A servlet container may place security restrictions on the environment in which a servlet executes. In a Java 5 Platform Standard Edition (J2SE 5.0) or Java™ Platform, Enterprise Edition 5 (Java EE 5) environment, these restrictions should be placed using the permission architecture defined by the Java Platform. For example, high-end application servers may limit the creation of a Thread object, to ensure that other components of the container are not negatively impacted.

1.4 Relationship with the HTTP Servlet API

The Java Servlet API is defined in the *Java Servlet Specification* [Servlet API]. It consists of a generic part defined as package `javax.servlet` and an HTTP specific part in package `javax.servlet.http`. This specification refers to the generic part using the unqualified term *Servlet API* and to the HTTP specific API as the *HTTP Servlet API*. The SIP Servlet API builds on the generic servlet API in much the same way as the HTTP Servlet API does, and is defined as package `javax.servlet.sip`. As such, a SIP servlet container **MUST** support the packages `javax.servlet` and `javax.servlet.sip`.

This specification is structured along the lines of the *Java Servlet Specification*, and like it, includes text not specific to the SIP Servlet API. Parts of this text has been copied from that other document, albeit modified to reflect the non-HTTP nature of SIP servlets, for example to use the broader term *servlet application* instead of *web application* when the context applies equally to SIP and HTTP applications.

1.4.1 Differences from the HTTP Servlet API

SIP was to some extent derived from HTTP and so the two protocols have much in common. Both are request-response protocols and messages have very similar structure and formats. However, in terms of providing services, there are important differences between the two protocols:

- HTTP services (including HTTP servlet applications) are almost exclusively hosted on HTTP origin servers, that is, on the Web server generating the final response to requests (as opposed to a proxy server). In contrast, an important function of SIP applications is intelligent request routing and the ability to act as a proxy is crucial in this context.
- HTTP is not a peer-to-peer protocol as is SIP and web applications never originate requests. SIP applications, on the other hand, need the ability to initiate requests of their own. An application that accepts an incoming call may have to terminate it by subsequently sending a BYE request towards the caller, a wakeup-call application may have to send the initial INVITE establishing a dialog, and a presence server application needs to be able to initiate NOTIFY requests. A back-to-back user agent (B2BUA) is a type of network based application that achieves a level of control over calls not attainable through proxying, and that requires client functionality, also. These examples demonstrate why client-side functionality is necessarily part of a SIP service infrastructure and explains the presence of such functionality in the SIP Servlet API.

It follows that, in addition to the ability inherited from the Servlet API of responding to incoming requests, the SIP Servlet API **MUST** support the following capabilities:

- generate multiple response (for example, one or more 1xx followed by a final response)
- proxying requests, possibly to multiple destinations
- initiate requests
- receive responses as well as requests

1.4.1.1 Use of Servlet.service()

In order to allow for these features, the SIP Servlet API uses the original Servlet API interfaces in a manner that differs from the HTTP Servlet API. SIP servlet applications are invoked when events occur in which they have registered an interest. These events can be either incoming requests or responses and they are delivered to applications through the service method of the `javax.servlet.Servlet` interface:

```
void service(ServletRequest request, ServletResponse response);
```

This is the application entry point used by the HTTP Servlet API also, but it is used slightly differently here. When used to process SIP traffic only one of the request and response objects is non-null. When invoked to handle incoming requests, the response argument will be null and vice versa, when invoked to handle incoming responses the request argument will be null.

Note: This caters for the fact that there is not necessarily a one-to-one correspondence between requests and responses in SIP applications.

1.4.1.2 Asynchronicity

Another important difference is the fact that the SIP Servlet API event model is asynchronous rather than synchronous as in the HTTP API. This means that applications are not obliged to respond to incoming requests in the upcall reporting the event. They may initiate some other action, return control to the container, and then respond to the request at some later point in time. The container relies on timeout of the application instance as a whole in order to guarantee that resources are always recovered eventually, even if an application fails to respond to the request in a timely manner.

Asynchronicity simplifies the programming of event driven services and allows an application such as a B2BUA not to hog threads while waiting for a potentially long-lived transaction to complete.

1.4.1.3 Application Composition

While the model of the HTTP Servlet API is to select a single application to process each incoming request, it is often desirable to apply multiple services to incoming SIP requests. This specification describes an application composition model that defines the conditions that a SIP servlet container must ensure are satisfied when it chooses to invoke multiple applications. This model is discussed in detail in the chapter [15 Application Selection And Composition Model](#).

1.4.2 Converged Applications

While the SIP Servlet API can certainly be implemented independently of the HTTP Servlet API, it is expected that many interesting services will combine multiple modes of communication, for example, telephony, Web, email, instant messaging and other server side components like Web services and other Java EE interfaces. A converged SIP container enables the deployment of applications that use SIP, HTTP Servlet API and other Java EE components like EJBs, webservices, messaging etc. [1.5 Deployment Models](#) describes different deployment models for a SIP Servlet container.

1.5 Deployment Models

The following three major deployment models are envisaged where the SIP Servlet container could be used. Each of these models have their uses in different real life scenarios and it is expected that the SIP Servlet container implementations shall allow the container to operate in all these three deployment models.

1.5.1 Standalone SIP Servlet Container

A standalone SIP Servlet container provides only SIP interface and hosts only SIP Servlets as its applications. This could be a lightweight container which could be used in situations where the application only has a SIP interface. This could also be used as an embeddable SIP Servlet container into a Java application.

1.5.2 SIP and HTTP Converged Servlet Container

SIP Servlets and HTTP Servlets are both derived from the same base Servlet API and in a converged SIP and HTTP application the SIP and HTTP Servlets shall share the same context. A combined SIP/HTTP servlet application will contain deployment descriptors pertaining to both the SIP and HTTP parts and will share state through the notion of application session objects representing instances of the application (see [13 Converged Container and Applications](#)). For this

reason, it is an important goal that it be possible to implement a servlet container that supports both the SIP and HTTP Servlet APIs at the same time and in a manner that allows applications to include both SIP and HTTP components.

1.5.3 SIP and Java EE Convergence

SIP Servlet Container will usually be part of a larger application server providing not just a HTTP Servlet container but also other Java EE applications like EJBs, Webservices etc. It is an important goal of this specification to facilitate the use of SIP Servlet technology in conjunction with the larger Java EE deployment model. There are several such features in this specification that shall allow the deployment of an application that has Java EE, HTTP and SIP components and there is a seamless access from one context to the other. It is expected that this shall enable development of very powerful applications.

These are just reference models of deployment. In order to be compliant with this specification the container **MUST** implement all the minimum features as specified in this specification.

1.6 Examples

1.6.1 A Location Service

Routing is an integral part of SIP and is a common function of SIP services. This example is a Location Service SIP application that performs a database lookup on the request URI of incoming requests and proxies the request to a set of destination addresses associated with that URI. The steps performed by the application and container are as follows:

1. Alice makes a call to sip:bob@example.com. The INVITE is received by the servlet container which invokes the Location Service.
2. The Location Service application determines, using non-SIP means, that the callee (Bob) is registered with two locations, identified by, say, two SIP URIs.
3. The service proxies to those two destinations in parallel, without record-routing, and in unsupervised mode. One of the destinations return 200 (OK) response and the other branch is cancelled by the container.
4. The 200 response is forwarded upstream to Alice and the call setup is completed as usual.

In this example, the application (and the host application server) is involved only in establishing the SIP dialog and will not be involved in subsequent signaling within that dialog.

1.6.2 A Multi-Protocol Servlet Application


The following is an example of a converged application. It consists of both SIP and HTTP components with SIP being used both in its capacity to establish media sessions and as a presence subscription and notification protocol [simple].

The service is called Call Schedule on Busy or No Answer (CSBNA). As in the previous example, it resides on a network server situated on the path between a caller and callee. When a call setup attempt fails because the callee is busy or not available, the application returns a URL pointing to a Web page allowing the caller to ask for a call to be set up the moment the callee becomes available. The caller may then elect to have the application establish a phone call when the callee becomes available again. The application knows about the availability of the callee by subscribing to his or her presence status. When a notification is received that the callee has become available, the application establishes a session between the original caller and callee using third party call control techniques [3pcc].

Figure 1-1 illustrates the call flow. The steps are as follows:

1. Alice makes a call to Bob. The INVITE is routed to the converged servlet container.
2. The CSBNA SIP servlet acts as B2BUA and sends an INVITE to Bob's SIP phone.
3. The converged servlet container also sends a 100 (Trying) informational response upstream to stop retransmissions.
4. Bob's phone returns a 486 (Busy Here) as Bob is currently in a phone call.
5. The CSBNA servlet returns a 302 (Temporarily Unavailable) response containing a single Contact with an HTTP URL pointing back to an HTTP servlet that is part of the CSBNA application.
6. The SIP application server sends an ACK for the 486 error response to Bob's phone as per the SIP specification.
7. Alice's SIP UA sends an ACK for the 302 error response to the application server.
8. Alice's SIP UA launches a Web browser to retrieve the HTTP Contact URL from the 302 response.
9. The HTTP CSBNA servlet returns an HTML page containing a form allowing Alice to have a call automatically setup based on Bob's availability. The form is pre-populated by the servlet with all required information, in particular Alice and Bob's SIP addresses.

10. Alice gasps at the sight of this amazing application and hits the Submit button on the HTML page thus causing an HTTP request to be sent to the application server.
11. The service sends a SIP SUBSCRIBE request to Bob's SIP UA. This will allow the service to know when Bob becomes available again.
12. Bob's UA accepts the subscription.
13. The HTTP CSBNA servlet returns a Web page to Alice which she can subsequently use to modify or cancel the scheduled call.
14. At some later point in time, Bob hangs up and his UA sends a notification of his new availability status to the CSBNA subscriber servlet in a SIP NOTIFY request.
15. The CSBNA application responds to the NOTIFY. (The application will unsubscribe from Bob's status at this point—this is not shown in the diagram.)

Figure 1-1 Call Flow for the CSBNA Application

In steps 16 through 24, the application establishes a call between Alice and Bob using the third party call control mechanisms described in [3pcc]. In this case, a call is established and media is exchanged over RTP.

In steps 25 through 28, the call is terminated.

In addition to illustrating the functioning of a converged container, this example demonstrates the need for the concept of application sessions. Over the lifetime of an instance of the CSBNA application, some 5 point-to-point SIP signaling relationships are created. Each of these correspond to a `SipSession` instance. Additionally, one `HttpSession` is created—this also corresponds to a point-to-point “signaling” relationship between the container and the HTTP client. Obviously, all of these protocol sessions are related and need to share state. The application session represents an instance of the application and binds together the protocol sessions representing individual signaling relationships and allows for sharing of state between them. The session model is further discussed in chapter 6 [Sessions](#).

1.6.3 A Java EE Converged Application

The following example demonstrates SIP and Java EE application convergence. It shows how a SIP UA (User-Agent) is notified of UA Profile change for which it had subscribed previously (The general mechanism of UA Profile propagation is described in the draft <http://tools.ietf.org/wg/sipping/draft-ietf-sipping-config-framework/>.)

In this example, there is a SIP UA which could be a SIP capable device like a SIP phone or SIP set-top box. This SIP UA is served by the SIP application server which is a SIP servlet container. The interaction uses the SIP Event Notification mechanism defined in [RFC 3265]. The SIP UA has its configuration data, also known as UA profile data which is maintained somewhere on the network. The SIP UA uses a converged application deployed on the SIP Servlet Container to access its profile data. As the data can also be changed by mechanisms other than SIP (out of band mechanisms) and the SIP UA can further subscribe (using SIP SUBSCRIBE) for these changes, the converged application can notify the SIP UA of the profile changes (using SIP NOTIFY).

As the data is supposed to be stored outside of the application context, the application knows about these out of band profile data changes via JMS messages from an abstract profile database. (where the profile actually resides is beyond the scope of this example)

This example is an enterprise application consisting of a SIP servlet application and an EJB. The EJB is actually a Message Driven Bean (MDB).

The SipServlet handles SUBSCRIBE and responses to the NOTIFY messages from the external SIP UA.

This application uses the `javax.servlet.sip.SipSessionsUtil` utility class described in [13 Converged Container and Applications](#). The `SipSessionsUtil` utility is a container provided lookup interface which can be used to access the reference to a particular `SipApplicationSession` given its id.

The steps involved in the working of this example are as follows:

1. The SIP Servlet application using the annotation `@EJB` or `ejb-ref` element in the deployment descriptor gets the reference to the EJB.
2. The EJB gets the reference to the `SipSessionsUtil` utility through an annotation defined in [18.2.7 Annotation for SipSessionsUtil Injection](#).
3. The SIP UA sends a SIP SUBSCRIBE to the application. In the `doSubscribe()` processing the `SipServlet` stores the `sip-application-session-id` and `sip-session-id` using a business method defined on the EJB. (This is the process of installing the subscription with the bean).

4. The SIP Servlet sends a 200 OK to the SUBSCRIBE and a NOTIFY according to the SIP procedures defined in [RFC 3265].
5. At some later time the user profile changes. The UA profile is managed by an abstract database which sends a JMS message to the application server whenever the profile is modified. The EJB (which is actually a MDB) in the application gets a JMS message on its queue. The message contains an identifier which tells us which UA this profile update is for.
6. The EJB maintains a map of UAs and the current sip-application-session-id. On getting the JMS message for a certain UA, it retrieves the sip-application-session-id from its map.
7. The MDB, which has access to the SipSessionsUtil utility, retrieves the `SipApplicationSession` reference using the sip-application-session-id.
8. Since the MDB also knows the `SipSession` id (previously set by the `SipServlet` and stored alongside the sip-application-session-id), it gets to the right SIP dialog from the `SipApplicationSession`.
9. Using normal in-session request creating mechanism the bean creates a new SIP NOTIFY request and sends the profile changes as payload on the dialog to the SIP UA
10. The SIP UA gets this SIP NOTIFY with profile data and sends a 200 OK. This 200 OK gets delivered to the `SIP Servlet` and the process completes.

Overview

2 The Servlet Interface

The `Servlet` interface is the central abstraction of the Servlet API and hence of the SIP Servlet API. All servlets implement this interface either directly or, more commonly, by extending a class that implements the interface. The generic Servlet API defines a class `GenericServlet` that is an implementation of the `Servlet` interface. The SIP Servlet API defines a class `SipServlet` that extends the `GenericServlet` interface and performs dispatching based on the type of message received. For most purposes, developers will extend `SipServlet` to implement their servlets.

2.1 Servlet Life Cycle

SIP servlets follow the lifecycle of servlets defined in [Servlet API, section 2.3]. All provisions in that section apply directly to SIP servlets. The lifecycle is illustrated in [Figure 2-1 Servlet lifecycle](#).

Figure 2-1 Servlet lifecycle.



Very briefly, the servlet container loads the servlet class, instantiates it and invokes the `init()` method on it, passing along configuration information in the form of a `ServletConfig` object. Having been successfully initialized, the servlet is available for service, and the container repeatedly invokes the service by calling the `service()` method with arguments representing

incoming requests and responses. When the container decides to deactivate the servlet instance it invokes the `destroy()` method – the servlet frees up resources allocated in `init()` and becomes garbage collected.

The failure to initialize any of the servlet within the application **MUST** be taken as failure of the entire application and the application **MUST** be taken out of service. If the failed servlet was a `load-on-startup` servlet then the Application Router **MUST NOT** be notified of this application's deployment using `SipApplicationRouter.applicationDeployed()`. The Application Router component is described in detail in [15 Application Selection And Composition Model](#). If the servlet happened to not be a `load-on-startup` but the initialization of it failed then the Application Router **MUST** be informed of this application's undeployment using `SipApplicationRouter.applicationUndeployed()`.

2.1.1 Servlet Life Cycle Listener

The Servlet `init()` method is called by the container to allow the Servlet to initialize one time setup tasks that are not specific to any application instance. The Servlet is not usable until the `init()` method successfully returns. A SIP Servlet in addition to receiving requests can also initiate requests (acting as a UAC) or create timers using the `TimerService` interface. Until the initialization of the Servlet is complete the SIP Servlet **SHOULD NOT** perform any of the signaling related tasks including sending SIP messages or setting timers. A listener for SIP Servlet lifecycle if present **MUST** be invoked by the container. The listener defines a single method indicating that the initialization of the Servlet is now complete and it can now receive messages as well as perform any other tasks.

```
void servletInitialized(javax.servlet.ServletContextEvent event);
```

The following sequence of initialization **MUST** be followed for a `load-on-startup` servlet.

1. Deploy the application.
2. Invoke `Servlet.init()`, the initialization method on the servlet. Invoke the `init()` on all the `load-on-startup` servlets in the application.
3. Invoke `SipApplicationRouter.applicationDeployed()` for this application.
4. If present, invoke `SipServletListener.servletInitialized()` on each of initialized servlet's listeners.

For applications without any servlets declared as `load-on-startup`, the `SipApplicationRouter.applicationDeployed()` **MUST** be invoked right after the

deployment has succeeded. The `init()` method on these servlets and `servletInitialized()` callback methods on their listener's would be called just before getting the first request.

2.2 Processing SIP Messages

The basic `Servlet` interface defines a `service` method for handling client requests. This method is called for each message that the servlet container routes to an instance of a servlet. The handling of concurrent messages in a servlet application generally requires the developer to design servlets that can deal with multiple threads executing within the `service` method at a particular time. Generally, the servlet container handles concurrent requests to the same servlet by concurrent execution of the `service` method on different threads.

SIP servlets are invoked to handle both incoming requests and responses. In either case, the message is delivered through the `service` method of the `Servlet` interface:

```
void service(ServletRequest req, ServletResponse res)
    throws ServletException, java.io.IOException;
```

If the message is a request the response argument **MUST** be null, and if the message is a response, the request argument **MUST** be null. When invoked to process a SIP event, the arguments must implement the `SipServletRequest` or `SipServletResponse` interfaces as the case may be. The `SipServlet` implementation of the `service` method dispatches incoming messages to methods `doRequest` and `doResponse` for requests and responses, respectively:

```
protected void doRequest(SipServletRequest req);
protected void doResponse(SipServletResponse resp);
```

These methods then dispatch further as described in the following sections.

Note: Section [18.2.2 @SipServlet Annotation](#) describes an alternative mechanism of declaring the Servlet class by using annotations.

2.3 SIP Specific Request Handling Methods

The `SipServlet` abstract subclass defines a number of methods beyond what is available in the basic `Servlet` interface. These methods are automatically called by the `doRequest` method (and indirectly from `service`) in the `SipServlet` class to aid in processing SIP based requests. These methods are:

- `doInvite` for handling SIP INVITE requests

- `doAck` for handling SIP ACK requests
- `doOptions` for handling SIP OPTIONS requests
- `doBye` for handling SIP BYE requests
- `doCancel` for handling SIP CANCEL requests
- `doRegister` for handling SIP REGISTER requests
- `doPrack` for handling SIP PRACK requests
- `doSubscribe` for handling SIP SUBSCRIBE requests
- `doNotify` for handling SIP NOTIFY requests
- `doMessage` for handling SIP MESSAGE requests
- `doInfo` for handling SIP INFO requests
- `doUpdate` for handling SIP UPDATE requests
- `doRefer` for handling SIP REFER requests
- `doPublish` for handling SIP PUBLISH requests

The first six Java methods correspond to request methods defined in the baseline SIP specification [RFC 3261]. The following methods correspond to request methods defined in various SIP extensions. PRACK is defined in [RFC 3262] and is discussed in [5.7.1 Reliable Provisional Responses](#). The SUBSCRIBE and NOTIFY methods are defined in the SIP event notification framework [RFC 3265] upon which the SIP presence framework is defined [simple]. The MESSAGE method supports instant messaging [IM], the INFO method is a general purpose mid-dialog signaling transport mechanism [RFC 2976], the UPDATE method is used as a mechanism to update SIP session parameters without affecting dialog state [RFC 3311], the REFER method is used for transfer of SIP calls [RFC 3515] and the PUBLISH method introduced in [RFC 3903] is a mechanism for SIP specific event state publication.

The `SipServlet` implementation of these methods is as follows. The `doAck` and `doCancel` methods do nothing. All other methods check whether the request is an initial request, as described in [10.2.9 Handling Subsequent Requests](#). If the request is initial, it is rejected with status code 501; otherwise the method does nothing (if the application proxied the initial request, the container will proxy the subsequent request when the method call returns). A servlet will typically override only those methods that are relevant for the service it is providing.

Note: The handling of incoming requests is asynchronous in the sense that servlets are not required to fully process incoming requests in the container's invocation of the service method. The request may be stored in a `SipSession` or `SipApplicationSession` object to be retrieved and responded to later—typically triggered by some other event. The container will not generate a response of its own if a servlet returns control to the container without having responded to an incoming request.

Applications wishing to handle SIP methods unknown to `SipServlet.doRequest` can override this method and invoke `super` as follows:

```
protected void doRequest(SipServletRequest request)
    throws ServletException, IOException
{
    if ("STORE".equals(request.getMethod())) {
        doStore(request);
    } else {
        super.doRequest(request);
    }
}
```

2.4 Receiving Requests

SIP servlets are invoked to process incoming requests in the following cases:

- It's an initial request, the container chooses to trigger an application based on application selection process described in [15 Application Selection And Composition Model](#).
- It's a subsequent request in a dialog for which the application is a UA or in which it is a proxy and it record-routed on the initial request.
- It's an ACK for a 2xx response to an INVITE which the application either responded to as a UAS or proxied with record-routing enabled.
- It's a CANCEL for an INVITE the application has received but not yet generated a final response for.

In all cases the servlet is invoked by the container through the `Servlet.service` method with a `SipServletRequest` object and `null` for the response argument.

If a servlet throws an exception when invoked to process a request other than ACK and CANCEL, the servlet container **MUST** generate a 500 response to that request. The header or body of the response may contain additional information that can be of use in identifying the cause of the problem.

2.5 SIP Specific Response Handling Methods

The `doResponse` method dispatches to one of the following methods based on the class of the status code of the response:

- `doProvisionalResponse` for handling SIP 1xx informational responses
- `doSuccessResponse` for handling SIP 2xx responses
- `doRedirectResponse` for handling SIP 3xx responses
- `doErrorResponse` for handling SIP 4xx, 5xx, and 6xx responses

Chapters 10 and 11 describe, in detail, the rules surrounding the invocation of these methods.

2.6 Number of Instances

Refer to [Servlet API, section 2.2].

The number of servlet instances created by a container is different from what is called application instances in this specification. This latter term effectively refers to an application session together with its contained protocol sessions as described in chapter [6 Sessions](#). Servlet objects are independent of any particular application instance and will typically process requests belonging to many different application instances.

3 Servlet Context

The `ServletContext` defines a servlet's view of the SIP application within which the servlet is running. Chapter 3 of the Java Servlet Specification describes the `ServletContext` [Servlet API] and it applies to the SIP servlet API, also. The following sections address issues specific to the SIP Servlet API.

3.1 The SipFactory

The `SipFactory` interface is used by servlets to create instances of various interfaces:

- requests: the `createRequest` methods create instances of the `SipServletRequest` interface and is used by UAC applications when creating new requests that do not belong to existing `SipSessions`. When creating subsequent requests in an existing dialog, `SipSession.createRequest` is used instead. [11.1.1 Creating Initial Requests](#) discusses requirements of the `createRequest` methods.
- address objects: ability to create `URI`, `SipURI`, `Address` and `Parameterable` instances.
- application sessions: ability to create new application sessions.

All servlet containers **MUST** make an instance of the `javax.servlet.sip.SipFactory` interface available to servlets via the context attribute of the same name, `javax.servlet.sip.SipFactory`.

With this specification the `SipFactory` instance can be injected using the Java Metadata annotations in applications in addition to context lookup. Additionally the `SipFactory` instance can also be injected into Java EE applications which do not have access to `ServletContext`.

Containers compliant with this specification **MUST** make the `SipFactory` instance available via this annotation as described in [13.2 Accessing SIP Factory](#).

3.2 Extensions Supported

SIP servlet containers **MUST** make an immutable instance of the `java.util.List` interface available as a `ServletContext` parameter with name `javax.servlet.sip.supported`. This `List` contains the the option tags of the SIP extensions registered with IANA, as supported by the container. This can be used by applications to determine whether the container supports a particular extension. For an example use see [5.7.1 Reliable Provisional Responses](#).

3.3 RFCs Supported

SIP servlet containers **MUST** make an immutable instance of the `java.util.List` interface available as a `ServletContext` parameter with name `javax.servlet.sip.supportedRfc`s. This `List` contains the RFC numbers of the SIP RFCs supported by the container. Using this `List`, applications will be able to determine whether the container supports a particular RFC and either adapt the call-flow accordingly or fail deployment of the application if a RFC which it relies on is not supported by the container. This context parameter shall help in portability of applications across containers.

3.4 Context Path

The Servlet API defines the notion of a *context path*. This is a URL path prefix with which a web application is associated. All requests with an HTTP URL starting with the context path of a web application will be routed to the corresponding servlet context. As SIP URIs do not have a notion of paths, the following `ServletContext` methods have no meaning for SIP-only servlet applications/containers and must return null:

```
ServletContext getContext(String uripath);  
String getRealPath(String path);  
RequestDispatcher getRequestDispatcher(String path);
```

As far as resource loading is concerned, the context path of a SIP-only servlet is always `"/`". For a combined HTTP and SIP application executing in an HTTP Servlet capable container, the context path is defined by the HTTP Servlet API and resource loading proceeds according to the Java Servlet Specification [Servlet API].

3.5 Context Parameters

Containers compliant with this specification **MUST** make the following references available to the applications as the `ServletContext` parameters.

Table 3-1 Context Parameters

Parameter Name	Description
<code>javax.servlet.sip.supported</code>	An immutable instance of the <code>java.util.List</code> interface containing the String names of SIP extensions supported by the container.
<code>javax.servlet.sip.supportedRfc</code> <code>cs</code>	An immutable instance of the <code>java.util.List</code> interface containing the RFC numbers represented as Strings of SIP RFCs supported by the container.
<code>javax.servlet.sip.100rel</code>	Parameter whose value suggests whether the container supports the 100rel extension i.e. RFC 3262. This parameter has been deprecated in this specification in favor of the <code>javax.servlet.sip.supported</code> parameter.
<code>javax.servlet.sip.SipSessions</code> <code>Util</code>	A container class <code>SipSessionsUtil</code> for ID based lookup of <code>SipApplicationSession</code> instances.
<code>javax.servlet.sip.SipFactory</code>	Instance of the applications <code>SipFactory</code> .
<code>javax.servlet.sip.outboundInt</code> <code>erfaces</code>	An immutable instance of the <code>java.util.List</code> interface containing the <code>SipURI</code> representation of IP addresses which are used by the container to send out the messages.
<code>javax.servlet.sip.TimerServic</code> <code>e</code>	Instance of the <code>TimerService</code> class.


Servlet Context

4 Addressing

Addressing plays a role in many SIP functions and so, in addition to the request URI, many header fields are defined to carry one or more addresses, for example, From, To, and Contact. The format of these addresses is generally the same. They consist of a URI with an optional display name and an optional set of parameters.

Figure 4-1 shows the addressing abstractions of the SIP Servlet API and the relationship between them.

Figure 4-1 Addressing Abstractions



The `SipFactory` interface described in [3.1 The SipFactory](#) is used to construct instances of these interfaces.

4.1 The Address Interface

The `Address` interface is used to represent values of header fields which conform to the following address BNF rule:

```
address= (name-addr | addr-spec) *(SEMI generic-param)
```

The constituent non-terminals are defined in RFC 3261, chapter 25, and are (incompletely) given below for ease of reference.

```
name-addr= [ display-name ] LAQUOT addr-spec RAQUOT
addr-spec      = SIP-URI / SIPS-URI / absoluteURI
display-name   = *(token LWS)/ quoted-string
```

The baseline SIP specification defines the following set of header fields that conform to this grammar: From, To, Contact, Route, Record-Route, Reply-To, Alert-Info, Call-Info, and Error-Info [RFC 3261]. The `SipServletMessage` interface defines a set of methods which operate on any address header field (see [5.4.1 Parameterable and Address Header Fields](#)). This includes the RFC 3261 defined header fields listed above as well as extension headers such as Refer-To [refer] and P-Asserted-Identity [privacy].

It is highly likely that future SIP extensions will define more such header fields and so it is useful to have a set of generic methods that can be used by applications to access header fields as `Address` objects rather than as `Strings`.

SIP Servlet containers will typically have built in knowledge of certain headers but are required to be able to handle unknown headers as well. When an application attempts to access an unknown header as an address header (by calling one of the methods discussed in [5.4.1 Parameterable and Address Header Fields](#)), the container MUST try to parse all values of that header field as `Address` objects according to the address rule.

The actual definition of address-based header fields differs in small ways:

- some define specific parameters and possibly limit the set of values they can take, for example, the From and To headers define a tag parameter which, when present, must be a token.
- some may not have parameters at all, for example, Reply-To.
- some may not have a display name

The address BNF rule can be thought of as defining a superset of all legal values for a large number of header fields. If the SIP Servlet container knows the actual, possibly more restrictive, definition for a particular address-based header it should enforce any such constraint when serializing a message.

4.1.1 The Parameterable Interface

Some SIP headers follow the following form described in RFC 3261, section 7.3.1:

```
field-name: field-value * (;parameter-name [=parameter-value])
```

where the field-value may be in `name-addr` or `addr-spec` format as defined in RFC 3261 or may be any sequence of tokens till the first semicolon.

The list of headers in RFC 3261 that follow this form are Accept, Accept-Encoding, Accept-Language, Alert-Info, Call-Info, Content-Disposition, Content-Type, Error-Info, Retry-After and Via. The `Address` API manifests a specialized form of this BNF for Address headers like Contact, From, To, Route, Record-Route and Reply-To. Besides these there are a number of other extensions where this form is applicable like Accept-Contact and Reject-Contact from RFC 3841. A generic `Parameterable` interface caters to all the SIP headers of this form.

The `Address` interface extends this interface.

Methods are available in `SipFactory` to create a `Parameterable` object from a `String` object and `SipServletMessage` to read/write `Parameterable` objects. In general any SIP header having parameters as in form described above, MUST be available as `Parameterable` to the application.

4.1.2 The From and To Header Fields

The From and To header fields contain the addresses of the UAC and UAS respectively. [11.1.1.1 Copying From and To Addresses](#) discusses how the container guarantees the integrity of From and To headers through cloning and immutability.

The previous version of this specification defined the From and To headers as system headers, thus restricting any modification to these headers. This specification relaxes that restriction to allow support for [RFC 4916]. This RFC extends the use of the From header by allowing the From header field URI to change during a dialog to reflect the connected identity. It also deprecates the mandatory reflection of the original To and From URIs in mid-dialog requests and their responses. This extension is compatible for containers compliant with [RFC 3261] as it requires the use of only tags from the From and To headers for dialog identification [RFC 3261 Section 12.2.1.1] and not the entire URIs as in the case for [RFC 2543].

This specification allows modifications to all parts of the From and To headers except the tag parameter. From and To headers are still considered system headers with respect to the tags (tag parameters). Containers should preserve any changes to the From and To headers in mid-dialog and use the updated ones for any subsequent requests in the dialog.

From and To header modification is supported through the following API methods:

```
SipServletMessage.getAddressHeader("From");
```

```
SipServletMessage.getAddressHeader("To");  
SipServletMessage.getFrom();  
SipServletMessage.getTo();
```

Applications can change all parts of the `Address` returned by the above methods except the tag parameter. The tag parameter **MUST NOT** be modified by applications and containers **MUST** throw an `IllegalStateException` if an attempt is made to set the tag parameter for these headers. The following methods operating on the `From` and `To` headers **MUST** continue to throw `IllegalArgumentException`:

```
SipServletMessage.setAddressHeader("From", fromHdr)  
SipServletMessage.addAddressHeader("To", toHdr)  
SipServletMessage.removeAddressHeader("From", fromHdr)
```

Note: Containers that need to support [RFC 2543] **MUST NOT** allow modification of the `From` and `To` headers as that RFC requires the entire URI for dialog identification. Container support for Connected Identity [RFC 4916] is optional in this specification and is indicated by the presence of the "from-change" option tag in the `javax.servlet.sip.supported` list. (see [3.2 Extensions Supported](#)).

4.1.3 The Contact Header Field

The `Contact` header field is another address header that warrants extra comment.

The `Contact` header field specifies one or more locations (URIs) where a user is reachable. This header field plays two different roles in SIP. One is as a mechanism for a UA to specify, for example in `INVITE` and `2xx` responses to `INVITE`, to its peer UA the exact address to which the peer should address subsequent messages within that dialog. In this case, the `Contact` header field will always have a single value. **Servlets must not set the `Contact` header in these cases.**

Containers know which network interfaces they listen on and are responsible for choosing and adding the `Contact` header in these cases. Containers should throw an `IllegalArgumentException` on application attempts to set the `Contact` header field in these cases.

The other use of `Contact` is in `REGISTER` requests and responses, as well as `3xx` and `485` responses. The value of `Contact` header fields in these messages specify alternate addresses for a user, and there may be more than one. These are the uses where SIP servlets may legitimately set `Contact` addresses. The `Contact` header field defines two “well-known” parameters, `q` and `expires` and the `Address` interface includes methods for those parameters.

The special Contact value “*” is used in REGISTER requests when the UAC wishes to remove all bindings associated with an address-of-record without knowing their precise values. The `isWildcard` method returns true for `Address` objects representing the wildcard Contact value and `SipFactory.createAddress` will return a wildcard `Address` given a value of “*”. Note that wildcard `Address` objects are legal only in Contact header fields.

Contact header is a system header which means that it is managed by the container and cannot be modified or set by the applications except for the following messages:

1. REGISTER requests and responses
2. 3xx responses
3. 485 responses
4. 200/OPTIONS responses

Unlike other system headers defined in this specification some of the Contact header constituents are modifiable by applications as described below :

1. A UA application MUST be able to set parameters in the Contact header.
2. A UA application MUST be able to set the user part of the Contact header.
3. Also as per Sec 11.2 of RFC 3261 *"Contact header fields MAY be present in a 200 (OK) response and have the same semantics as in a 3xx response. That is, they may list a set of alternative names and methods of reaching the user."* So UAs responding to OPTIONS must be capable of setting Contact header(s) like in a 3xx response.

The UA applications can perform the above 3 operations on the Contact URI as retrieved using the following API:

```
SipServletMessage.getAddressHeader("Contact");
```

Note that the `Address` returned should be good only for specifying a set of parameters that the application can set on `Address` and/or `URI` and set the user part in the `URI`. The host component of the `URI` and `URI Scheme` are irrelevant and cannot be trusted to reflect the actual values that the container will be using when inserting a Contact header into Request or Response. The container MUST ignore or overwrite any host/port set on the Contact `URI` accessed as above. This is because the container is responsible for managing the actual network listen points and uses these to create Contact headers' host/port.

Further following list of `URI` parameters MUST NOT be modified by the applications and containers MUST ignore any value for these parameters if set by applications.

- method
- ttl
- maddr
- lr

The applications may set any other SIP URI parameter or Contact header parameter relevant for the message.

The following method **MUST** continue to throw `IllegalArgumentException` wherever the Contact header is defined as system header by this specification:

```
SipServletMessage.setAddressHeader("Contact", contactHdr);
```

In case the application modifies the Contact as specified above but decides to proxy the request subsequently, the containers **MUST** ignore any modification made to the Contact header.

4.1.3.1 Transport Parameters in Contact Header

While setting the Contact header for an outgoing message the containers **MUST** set the transport parameter to the transport that is used to send this message.

4.2 URIs

SIP entities are addressed by URI. When initiating or proxying a request, SIP servlets identify the destination by specifying a URI. SIP defines its own URI scheme that SIP containers are required to support. Particular implementations may know how to handle other URI schemes, e.g. tel URLs [RFC 3966]. Implementations are required to be able to represent URIs of any scheme, so that if, for example, a 3xx response contains a Contact header with a mailto or http URL, the container is able to construct `Address` objects containing `URIs` representing those Contact URIs. Also, `SipFactory.createURI` should return a `URI` instance given any valid URI string. The container may not be able to route SIP requests based on such URIs but must be able to present them to applications.

4.2.1 SipURI

This interface represents SIP and SIPS URIs. Implementations are required to be able to route requests based on SIP URIs.

SIP and SIPS URIs are similar to email addresses in that they are of the form `user@host` where `user` is either a user name or telephone number, and `host` is a host or domain name, or a numeric IP address. Additionally, SIP/SIPS URIs may contain parameters and headers. See RFC 3261, section 19.1.1 for restrictions on the contexts in which various parameters are allowed. Headers are allowed only in SIP/SIPS URIs appearing in Contact headers or in external URIs, for example when being used as a link on a Web page.

As an example, the following SIP URI:

```
sip:alice@example.com;transport=tcp?Subject=SIP%20Servlets
```

contains a transport parameter with value “tcp” and a Subject header with value “SIP Servlets”.

The string form of SIP/SIPS URIs may contain escaped characters. The SIP servlet container is responsible for unescaping those characters before presenting URIs to servlets. Likewise, string values passed to setters for various SIP/SIPS URI components may contain reserved or excluded characters that need escaping before being used. The container is responsible for escaping those values as necessary. Syntactically, SIP and SIPS URIs are identical except for the name of the URI scheme. The semantics differ in that the SIPS scheme implies that the identified resource is to be contacted using TLS. Quoting from RFC 3261:

A SIPS URI specifies that the resource be contacted securely. This means, in particular, that TLS is to be used between the UAC and the domain that owns the URI. From there, secure communications are used to reach the user, where the specific security mechanism depends on the policy of the domain. Any resource described by a SIP URI can be “upgraded” to a SIPS URI by just changing the scheme, if it is desired to communicate with that resource securely.

SIP and SIPS URIs are both represented by the `SipURI` interface as they’re syntactically identical and are used the same way. The `isSecure` method can be used to test whether a `SipURI` represents a SIP or a SIPS URI and the `setSecure` method can be used to change the scheme.

4.2.2 TelURL

This interface represents tel URLs as defined in [RFC 3966]. The tel URL scheme is used to represent addresses of terminals in the telephone network, i.e. telephone numbers. SIP servlet containers may or may not be able to route SIP requests based on tel URLs, but must be able to parse and represent them. `SipFactory.createURI` must return an instance of the `TelURL` interface when presented with a valid tel URL.

Addressing

5 Requests and Responses


This chapter describes the structure of SIP servlet message objects. Chapters 10 and 11 describe operational aspects of message processing.

5.1 The SipServletMessage Interface

The generic Servlet API is defined with an implicit assumption that servlets receive requests from clients, inspect various aspects of the corresponding `ServletRequest` object, and generate a response by setting various attributes of a `ServletResponse` object. As HTTP servlets reside only on origin servers and always generate responses to incoming requests, this model is a good fit for HTTP.

The requirement that SIP services must be able to initiate and proxy requests implies that SIP request and response classes need to be more symmetric, that is, requests must be writable as well as readable, and likewise, responses must be readable as well as writable.

The `SipServletMessage` interface defines a number of methods which are common to `SipServletRequest` and `SipServletResponse`, for example setters and getters for message headers and content. Figure 5-1 illustrates how the SIP request and response interfaces extend the generic `javax.servlet` interfaces the same way the HTTP request and response interfaces does, but additionally implement the `SipServletMessage` interface.

Figure 5-1 Request-response Hierarchy

5.2 Implicit Transaction State

`SipServletRequest` and `SipServletResponse` objects always implicitly belong to a SIP transaction, and the transaction state machine (as defined by the SIP specification) constrains which messages can legally be sent at various points of processing. If a servlet attempts to send a message which would violate the transaction state machine, the container throws an `IllegalStateException`.

A `SipServletMessage` for which one of the following conditions is true is called committed:

- the message is an incoming request for which a final response has been generated
- the message is an outgoing request which has been sent
- the message is an incoming non-reliable provisional response received by a servlet acting as a UAC
- the message is an incoming reliable provisional response for which PRACK has already been generated. (Note that this scenario applies to containers that support the 100rel extension.)
- the message is an incoming final response received by a servlet acting as a UAC for a Non INVITE transaction
- the message is a response which has been forwarded upstream
- message is an incoming final response to an INVITE transaction and an ACK has been generated
- message is an outgoing request, the client transaction has timed out and no response was received from the UAS and the container generates a 408 response locally

The semantics of committed message is that it cannot be further modified or sent in any way. Also see [5.8 Accessibility of SIP Servlet Messages](#).

5.3 Access to Message Content

The HTTP Servlet API provides access to message content through a stream metaphor.

Applications have access to posted data through an `InputStream` or a `Reader` and can generate content through either an `OutputStream` or a `Writer`.

SIP is more like email than HTTP with regard to message content. Messages are generally smaller and using a chunked encoding for generated content is not practical due to the SIP requirement that a `Content-Length` header be included in all messages. Also, SIP applications will often need access to parsed representations of message content.

For these reasons there are no stream based content accessors defined for SIP messages, and the `ServletRequest` methods `getInputStream` and `getReader` as well as `ServletResponse` methods `getOutputStream` and `getWriter` must return null.

The `SipServletMessage` interface defines the following setters and getters for message content (exceptions omitted for clarity):

```
int getContentLength();
void setContentLength(int len);

String getContentType();
void setContentType(String type);

Object getContent();
byte[] getRawContent();
void setContent(Object obj, String type);
```

The interface parallels how the JavaMail API defines access to message content in the `javax.mail.Part` interface [JavaMail]. The `getRawContent` method is not present in JavaMail but is useful when writing back-to-back user agents (B2BUA) as these may wish to copy content from one message to another without incurring the overhead of parsing it (as they may not actually care what is in the body). Which type of `Object` is returned by `getContent` depends on the message's `Content-Type`. It is required to return a `String` object for MIME type `text/plain` as well as for other text MIME media types for which the container does not have specific knowledge. It is encouraged that the object returned for multipart MIME content is a `javax.mail.Multipart` object. For unknown content types other than text, the container must return a `byte[]`. Likewise, `setContent` is required to accept `byte[]` content with any MIME type, and `String` content when used with a text content type. When invoked with non-`String` objects and a text content type, containers should invoke `toString()` on the content `Object` in

order to obtain the body's character data. Again, it is recommended that implementations know how to handle `javax.mail.Multipart` content when used together with multipart MIME types.

5.3.1 Character Encoding

Several of the message content accessors discussed above need to be able to convert between raw eight-bit bytes and sixteen-bit Unicode characters. The character encoding attribute of `SipServletMessage` specifies which mapping to use:

```
String getCharacterEncoding();
void setCharacterEncoding(String enc)
    throws UnsupportedOperationException;
```

Character encodings are identified by strings and generally follow the conventions documented in [RFC 2278]. The character encoding may affect the behavior of methods `getContent` and `setContent`. A message's character encoding may be changed by calls to `setCharacterEncoding` and `setContentType`. For incoming messages, the character encoding is specified by the `charset` parameter of the Content-Type header field, if such a parameter is present.

Note: Because of evolution of servlet spec 2.4, `SipServletResponse.setCharacterEncoding()` which extends both `SipServletMessage` and `ServletResponse` now does not throw the `UnsupportedEncodingException` because it inherits a more generic method from `ServletResponse`. This is a binary compatible change, meaning that the compiled applications shall run on the v1.1 containers without change but it is not a source compatible change if the application is explicitly catching the `UnsupportedEncodingException` on `SipServletResponse.setCharacterEncoding()`.

5.4 Headers

A servlet can access the headers of a SIP message through the following methods of the `SipServletMessage` interface (see also [Servlet API, sections 4.3 and 5.2]):

```
String getHeader(String name);
ListIterator getHeaders(String name);
Iterator getHeaderNames();
void setHeader(String name, String value);
void addHeader(String name, String value);
```


The `getHeader` method allows access to the value of a named header field. Some header fields, for example `Warning`, may have multiple values in a SIP message. In this case `getHeader` returns the first value of the header field. The `getHeaders` method allow access to all values of a specified header field by returning an iterator over `String` objects representing those values.

The `setHeader` method sets a header with a given name and value. If a previous header exists, it is replaced by the new header. In the case where a set of header values exist for the given name, all values are cleared and replaced with the new value.

The `addHeader` method adds a header with the specified name and value to the message. If one or more headers with the given name already exists, the new value is appended to the existing list.

5.4.1 Parameterable and Address Header Fields

The methods listed above treat SIP header field values as `Strings`. As discussed in [4.1 The Address Interface](#) many SIP header fields carry `Parameterables` (and hence `Addresses`) and having the ability to access such `Parameterables` (and hence `Addresses`) in a parsed form is both more convenient and allows for better performance than accessing those header field values as `Strings`. The following methods on the `SipServletMessage` interface are defined in terms of the `Parameterable` and `Address` interfaces:

```
Parameterable getParameterableHeader(java.lang.String name);
java.util.ListIterator getParameterableHeaders(java.lang.String name)
void setParameterableHeader(java.lang.String name, Parameterable param)
void addParameterableHeader(java.lang.String name, Parameterable param,
boolean first)

Address getAddressHeader(String name);
ListIterator getAddressHeaders(String name);
void setAddressHeader(String name, Address addr);
void addAddressHeader(String name, Address addr, boolean first);
```

If a header field has multiple `Parameterable` values in a message, the `getParameterableHeader` method returns the first value. The `getParameterableHeaders` method allows access to all values of the specified header field returning an iterator over `Parameterable` objects. The same explanation applies to the `getAddressHeader` and `getAddressHeaders` methods if a header field has multiple `Address` values.

`Parameterable` and `Address` objects obtained from, or added to messages are live in the sense that modifications made to them cause the corresponding header field value of the underlying message or messages to be modified accordingly.

`Parameterable` and `Address` objects may belong to more than one `SipServletRequest` at a time. In this case modification of any `Parameterable` or `Address` object will result in modification of the value of the underlying header field for both messages. This sort of aliasing can result in bugs when not intended. Application writers may practice defensive programming by deep cloning the `Parameterable` and `Address` objects to avoid sharing.

5.4.2 System Headers

The term system header is used to refer to those headers that are managed by the servlet container and which servlets must not attempt to modify directly via calls to `setHeader` or `addHeader`. This includes the headers `Call-ID`, `From`, `To`, `CSeq`, `Via`, `Record-Route`, `Route`, `Path` as well as `Contact` when used to confer a session signaling address, that is, in messages other than `REGISTER` requests and responses, `200` response to `OPTIONS` as described by section 11.2 [RFC3261], and `3xx` and `485` responses. For the `contact` header only the parameters and user part can be modified/added by the application as described in [4.1.3 The Contact Header Field](#). For the `From` and `To` headers all parts of the headers except the tags (tag parameter) can be modified as described in [4.1.2 The From and To Header Fields](#). There is no need for services to set the system headers directly and disallowing it makes it easier for containers to ensure correct protocol behavior. SIP servlet containers throw an `IllegalArgumentException` on attempts to modify system headers.

Containers are required to enforce this immutability requirement also when system headers are accessed through `Address` objects or through the `ListIterator` returned by the `getAddressHeaders` method.

5.4.3 TLS Attributes

If an incoming request or response has been transmitted over a secure protocol, such as TLS, this information must be exposed via the `isSecure` method of the `SipServletRequest` interface.

Note: The `isSecure` method indicates whether a message was received over a secure transport. However, since secure protocols are frequently used in a hop-by-hop manner in SIP, it does not follow that the message was sent over a secure transport in all hops before reaching the container. It may well have passed over an insecure link at some point. If there is a TLS certificate associated with the message, it **MUST** be exposed to the servlet programmer as an array of objects of type `java.security.cert.X509Certificate` and accessible via a `SipServletRequest` or `SipServletResponse` attribute of `javax.servlet.request.X509Certificate` for requests and `javax.servlet.response.X509Certificate` for responses. Since the TLS

connection is hop-by-hop this TLS certificate is not available beyond the first application in case of application composition chain as described in [15.7 Transport Information](#).

5.5 Transport Level Information

The following `SipServletMessage` methods allow applications to obtain transport level information from incoming messages:

```
String getLocalAddr();
int getLocalPort();
String getRemoteAddr();
int getRemotePort();
String getTransport();
String getInitialRemoteAddr();
int getInitialRemotePort();
String getInitialTransport();
```

These methods have meaning for incoming messages only, in which case they return the IP address/port number on which the message was received locally (`getLocalAddr`, `getLocalPort`), or the IP address/port number of the sender of the message (`getRemoteAddr`, `getRemotePort`, `getInitialRemoteAddr`, `getInitialRemotePort`), as well as the transport protocol used, e.g. UDP, TCP, TLS, or SCTP.

5.6 Requests

The `SipServletRequest` object encapsulates information of a SIP request, including headers and the message body.

5.6.1 Parameters

Request parameters are strings sent by the client to a servlet container as part of a request. The parameters are made available to applications as a set of name-value pairs. Multiple parameter values can exist for any given parameter name. The following `ServletRequest` methods are available to access parameters:

```
String getParameter(String name);
Enumeration getParameterNames();
String[] getParameterValues(String name);
```

The `getParameterValues` method returns an array of `String` objects containing all the parameter values associated with a parameter name. The value returned from the `getParameter` method must always equal the first value in the array of `String` objects returned by `getParameterValues`.

Note: Support for multi-valued parameters is defined mainly for HTTP because HTML forms may contain multi-valued parameters in form submissions.

When passing an incoming `SipServletRequest` to an application, the container populates the parameter set in a manner that depends on the nature of the request itself:

- For initial requests where the application is invoked the parameters are those present on the request URI, if this is a SIP or a SIPS URI. For other URI schemes, the parameter set is undefined.
- For initial requests where a preloaded Route header specified the application to be invoked, the parameters are those of the SIP or SIPS URI in that Route header.
- For subsequent requests in a dialog, the parameters presented to the application are those that the application itself set on the Record-Route header for the initial request or response (see [10.4 Record-Route Parameters](#)). These will typically be the URI parameters of the top Route header field but if the upstream SIP element is a “strict router” they may be returned in the request URI (see RFC 3261). It is the containers responsibility to recognize whether the upstream element is a strict router and determine the right parameter set accordingly.

Note: In addition, this specification introduces the access to popped Route header which is popped by the container when the Route header points towards the container. This new mechanism allows the application to have complete access to the popped route header and all parameters. See [5.6.3 Popped Route Header](#). It is recommended that applications access the Route parameters from the popped route rather than from the request parameters, in future this mechanism may be deprecated.

5.6.2 Attributes

Attributes are objects associated with a message. Attributes may be set by the container to express information that otherwise could not be expressed via the API, or may be set by a servlet to communicate information to another servlet (via `RequestDispatcher`). Attributes set on a message by a servlet **MUST** be available to other servlets within the same application only when `RequestDispatcher` is used to forward the message.

Attributes are accessed with the following methods of the `ServletRequest` interface:

```
Object getAttribute(String name);
```

```
Enumeration getAttributeNames();
void setAttribute(String name, Object o);
void removeAttribute(String name);
```

Only one attribute value may be associated with an attribute name.

Attribute names beginning with the prefix `javax.servlet.sip.` are reserved for definition by this specification.

It is suggested that all attributes placed into the attribute set be named in accordance with the reverse package name convention suggested by the Java Programming Language Specification for package naming [JLS].

5.6.3 Popped Route Header

On receiving an initial request that contains a SIP Route header (preloaded) or receiving a subsequent request with Route header (converted from a Record-Route header), a SIP Servlet container will determine if the request is intended for itself (this is based on local policy, for example IP addresses of interfaces or representative DNS entries). **If an initial request is identified as being intended for the SIP Servlet container it MUST remove the Route header before passing it to any application or the Application Router.**

A side effect of removing a SIP Route message header before presenting the request to applications (and Application Router) is that applications do not have access to the SIP Route message header and its associated information. Certain architectures utilize the SIP Route header for transporting application and other related information. The following methods return the Route header popped by the container.

```
Address getPoppedRoute();
Address getInitialPoppedRoute();
```

If application composition is being used, the values returned by these methods may differ. The first method (`getPoppedRoute`) returns the route popped before current application invocation in the composition chain. The second one (`getInitialPoppedRoute`) returns the route popped by the container when it first received the request.

If no header has been popped by the SIP Servlet container on an initial request, then both methods return null.

Note that the above methods return the Route header as an `Address`. So, parameters added to the Record-Route header using the `Proxy.getRecordRouteURI()` API should be retrieved not from the popped route `Address` directly but from the URI of the popped route `Address`.

5.7 Responses

Response objects encapsulate headers and content sent from UA servers upstream to the client. Unlike the case for HTTP, a single request may result in multiple responses. The

`SipServletResponse.getRequest()` method returns the request object associated with a response. For UAC and UAS applications this is the original `SipServletRequest`. For proxying applications it is an object representing the request that was sent on the branch on which the response was received - `getProxy().getOriginalRequest()` can be used to obtain the original request object (see also [10.2.4.2 Correlating responses to proxy branches](#) and [6.2.3 Creation of SipSessions](#)).

5.7.1 Reliable Provisional Responses

Provisional responses (1xx's) are not sent reliably in baseline SIP. However, an extension has been defined that allows transmission of 1xx's other than 100 with guarantees concerning reliability and ordering [RFC 3262]. This is useful in cases where a provisional response carries information critical to providing a desired service, often related to PSTN interworking.

Support for the 100rel extension is optional in SIP servlet containers ("100rel" is the name of the option tag defined for the provisional response extension). If implemented, the container must include the string "100rel" in the list of supported extensions available to applications through the `ServletContext`, see [3.2 Extensions Supported](#).

Applications can determine at runtime whether the container supports the 100rel extension by testing whether this `String` is present in the supported list.

When wishing to send a 1xx reliably, the application invokes

`SipServletResponse.sendReliably()`. If this method call is successful, the container will send the response reliably. A `Rel100Exception` is thrown if the response is not within 101-199 range, if the request is not an INVITE, if the UAC didn't indicate support for the 100rel extension in a Supported or Required header, or if the container itself doesn't support the extension.

For containers that do support the 100rel extension, the RSeq and RACK headers are system headers (see [5.4.2 System Headers](#)), that is, they are handled by the container and may not be added, modified, or deleted by applications. PRACK requests are treated as other subsequent requests, meaning they will be associated with the same `SipSession` as the corresponding INVITE is, and will get delivered to `Servlet.service()` whose `SipServlet` implementation dispatches to the `doPrack()` method.

Following method has been defined on `SipServletResponse` to create PRACK requests :

```
SipServletRequest createPrack();
```

Since different PRACKs can be generated on different reliable responses and since RACK is system header, this method must be used to correctly to create a PRACK request.

If no PRACK is received for a reliable provisional response within the time specified by RFC 3262, the container will inform the application through the `noPrackReceived` method of the `SipErrorListener` interface if this is implemented by the application. It is then up to the application to generate the 5xx response recommended by RFC 3262 for the INVITE transaction. The original INVITE request as well as the unacknowledged reliable response is available from the `SipErrorEvent` passed to the `SipErrorListener`.

When a container supporting 100rel receives a retransmission of a reliable provisional response, it does not invoke the `application(s)` again.

Also, containers supporting 100rel are responsible for guaranteeing that UAC applications receive incoming reliable provisional responses in the order defined by the RSeq header field.

5.7.2 Buffering

The Servlet API defines a number of `ServletResponse` methods related to buffering of content returned by a server in responses [Servlet API, section 5.1]:

- `getBufferSize`
- `setBufferSize`
- `reset`
- `flushBuffer`

These methods can improve performance of HTTP servlets significantly. However, unlike HTTP, SIP is not intended as a content transfer protocol and buffering is not usually an issue of concern. Therefore, it is not expected that these methods will yield any useful result and implementations may simply do nothing. It is recommended that `getBufferSize` return 0.

The `isCommitted()` method returns true if the message is committed in the sense of [5.2 Implicit Transaction State](#).

5.8 Accessibility of SIP Servlet Messages

Access to SIP Servlet messages is not limited to the scope of the Servlet's `service` method. Applications can handle a SIP message within the service method but also from other threads

(either timer events or any other thread in the system) outside the scope of the Servlet's `service` method. E.g. An application can respond to a request based on an event delivered via JMS.

A `SipServletMessage` cannot be sent again after it has been committed as defined in [5.2 Implicit Transaction State](#). Further since the committed message belongs to a context which has completed its state machine or lifecycle, any modification of the message is meaningless. Containers SHOULD throw an `IllegalStateException` for any mutation attempt on a committed message. However, the URI and other headers of the committed messages can be used to construct a new message subject to following restriction: In case the URI or any other header is modified for use it MUST be cloned by the application as it is not guaranteed that containers will return a deep copy on access, even if a message is committed the container may still access them for handling a retransmission.

5.9 Internationalization

Language identification tags are used in several places in SIP, notably `Accept-Language` and `Content-Language` headers and the `charset` parameter of various `Content-Type` header values.

In the SIP Servlet API, languages are identified by instances of `java.util.Locale`.

Note: While the `javax.servlet` interfaces `ServletRequest` and `ServletResponse` contains methods related to internationalization, these assume that servlets only respond to incoming requests and are insufficient for the SIP Servlet API.

5.9.1 Indicating Preferred Language

User agents may optionally indicate to proxies and peer UAs in which natural language(s) it prefers to receive content, reason phrases, warnings, etc. This information can be communicated from the UA using the `Accept-Language` header.

The following methods are provided in the `SipServletMessage` interface to allow the sender of a message to indicate preferred locale(s):

```
void setAcceptLanguage(Locale locale);  
void addAcceptLanguage(Locale locale);
```

The `setAcceptLanguage` method sets the preferred language of the `Accept-Language` header and removes any existing `Accept-Language` header, while `addAcceptLanguage` adds another (least preferred) locale to the list of acceptable locales.

The following `SipServletMessage` methods are used to determine the preferred locale of the sender of the message:


```
Locale getAcceptLanguage();
Iterator getAcceptLanguages();
```

The `getAcceptLanguage` method will return the preferred locale that the client will accept content in. See section 14.4 of RFC 2616 (HTTP/1.1) for more information about how the Accept-Language header must be interpreted to determine the preferred language of the client. If no preferred locale is specified by the client, `getAcceptLanguage()` must return null and `getAcceptLanguages()` must return an empty `Iterator`. (Note that this behavior has changed from JSR 116 and is noted in Appendix A.)

The `getLocales` method will return an `Iterator` over the set of `Locale` objects indicating, in decreasing order starting with the preferred locale, the locales that are acceptable to the UA originating the message. If no preferred locale is specified by the client, `getLocale` must return the default locale for the servlet container and `getLocales` must return an `Iterator` over a single `Locale` element of the default locale.

5.9.2 Indicating Language of Message Content

When sending a message containing a body, SIP servlets may indicate the language of the body by calling the `setContentLanguage` method of the `SipServletMessage` interface:

```
void setContentLanguage(Locale locale);
```

This method must correctly set the Content-Language header (along with other mechanisms described in the SIP specification), to accurately communicate the `Locale` to the client.

Note that a call to the `setContent` or `setContentTyped` methods with a charset component for a particular content type, will set the message's character encoding.

The default encoding of message content is “UTF-8” if none has been specified by the servlet programmer. Upon receiving messages, servlets can obtain information regarding the locale of the message content using the following `SipServletMessage` method:

```
Locale getContentLanguage();
```

Requests and Responses

6 Sessions

SIP applications typically must process multiple messages in order to provide the intended service. As servlets themselves are stateless (or rather, contain no per-dialog or per-transaction data), the API provides a mechanism that allows messages to be correlated and specify how containers associate application data with subsequent messages processed in an “application instance”.

The HTTP Servlet API provides such a mechanism in the form of HTTP sessions. The `HttpSession` interface allows servlets to correlate a series of HTTP requests from a particular client and also acts as a store for application data.

The `SipSession` interface is the SIP Servlet API equivalent of `HttpSession`. It represents a point-to-point relationship between two user agents and roughly corresponds to a SIP dialog [RFC 3261]. However, SIP applications are typically more complicated than Web applications:

- many services involve multiple dialogs, for example conferencing applications and applications acting as back-to-back user agents and third-party controllers
- converged applications communicate with other network elements using multiple protocols, for example SIP, HTTP, email, etc.
- application composition allows for many applications active on one call.

This implies that there may be more than one application invoked on a single call and any one application instance may consist of multiple point-to-point relationships, and that these relationships may employ different protocols. This is reflected in the SIP Servlet API through the notions of protocol sessions and application sessions. A protocol session is a protocol specific

session object typically representing a point-to-point relationship. The `SipSession` and `HttpSession` interfaces are both examples of protocol sessions.

An application session in a sense represents an application instance. It contains a number of protocol sessions and is also used as a container for application data. All protocol sessions belonging to the same application instance belong to the same `SipApplicationSession`. For example, a SIP servlet acting as a back-to-back user agent will consist of two `SipSessions` and one `SipApplicationSession` for each application instance.

Containers may, as an optimization, create application session and SIP session objects lazily, for example postpone creation until requested by the application. The result should be indistinguishable from an implementation that always creates the session objects.

6.1 SipApplicationSession

The application session serves two purposes: it provides storage for application data and correlates a number of protocol sessions. Note: The application session is not SIP specific, but for practical reasons, this version of the SIP Servlet API defines the application session as `javax.servlet.sip.SipApplicationSession`, whereas the more logical choice would be `javax.servlet.ApplicationSession`. It is our hope that a future version of the Servlet API will adopt the notion of application sessions, in which case the `SipApplicationSession` will be deprecated and/or refactored to extend `ApplicationSession`.

6.1.1 Protocol Sessions

The following `SipApplicationSession` methods allow servlets to obtain contained protocol sessions:

- `getSessions()` iterates over all valid child protocol session objects
- `getSessions(String protocol)` iterates over all valid child session objects that are of the specified protocol, for example “SIP” to get all `SipSessions`, and “HTTP” to get all `HttpSessions`.
- `getSipSession(java.lang.String id)` returns a certain `SipSession` by its id (since v1.1)
- `getSession(java.lang.String id, Protocol protocol)` returns a child session associated with the specified protocol by its id (since v1.1)

6.1.2 SipApplicationSession Lifetime

Containers manage session data and so need a mechanism for knowing when `SipApplicationSession` objects are no longer in use and are therefore eligible for garbage collection. The mechanism provided is known as `SipApplicationSession` invalidation. An application session becomes invalidated in one of three ways:

1. The `SipApplicationSession` expires and the container subsequently invalidates it.
2. A servlet explicitly invalidates it by invoking the `invalidate()` method.
3. A servlet marks the `SipApplicationSession` to be invalidated and the container invalidates it when the `SipApplicationSession` is in the ready-to-invalidate state as described in [6.1.2.2.2 Invalidate When Ready Mechanism](#).

For reasons of performance, it is recommended that applications explicitly invalidate `SipApplicationSession` and `SipSession` objects as soon as possible. Note: It would be unfortunate if applications were to force creation of an application session just so that they can invalidate it. The `getApplicationSession(boolean create)` method can be used with a false argument to avoid forcing creation of the session object.

When used, an application session expiration timer ensures that application sessions will eventually become eligible for container invalidation regardless of whether an application explicitly invalidates them. Servlets can register for application session timeout notifications using the `SipApplicationSessionListener` interface. In the `sessionExpired()` callback method, the application may request an extension of the application session lifetime by invoking `setExpires()` on the timed out `SipApplicationSession` giving as an argument the number of minutes until the session expires again. The container may grant the extension of session lifetime, grant it with a modified timeout value, or reject it. The ability to accept with a different timeout value allow containers to apply their own policies on application session lifetime and timeouts. A container might for example choose to enforce a maximum total lifetime for application sessions. A `SipApplicationSession` object is said to be expired when its expiration timer is no longer active. An expired `SipApplicationSession` is not invalid until the container explicitly invalidates it.

The ability to extend session lifetime is useful to applications because it allows them to not use an unrealistically high expiration timer value in cases where application lifetime depends on some “external” event, that is, an event unknown to the servlet container.

6.1.2.1 SipApplicationSession Timer Operation and SipApplicationSession Expiration

When a `SipApplicationSession` is created, the `SipApplicationSession` timer starts if the timeout value specified by the `session-timeout` parameter in the deployment descriptor or `@SipApplication(sessionTimeout)` annotation is set to a positive number. If not specified, the default timeout value is set to 3 minutes. Explicit invalidation of the `SipApplicationSession` leads to cancellation of the application session timer. If the session timeout value is 0 or less, then an application session timer never starts for the `SipApplicationSession` object and the container does not consider the object to ever have expired. However, if a the session timeout value was set to 0 or less, and a servlet subsequently calls the `setExpires()` method on a `SipApplicationSession` object, it is left up to container implementation whether to accept that request and start an expiration timer for the `SipApplicationSession` or to reject the `setExpires()` call by returning 0. If the container accepts the `setExpires()` request thereby starting an expiration timer, then it becomes the container's responsibility to invalidate the `SipApplicationSession` when it expires if the application neglects to do so.

A `SipApplicationSession` expires in one of two ways:

1. Container rejects an application session lifetime extension by returning 0 to a `setExpires()` call.
2. The application's `SipApplicationSessionListener` implementation chooses not to call `setExpires()` on the `SipApplicationSession` object inside the `sessionExpired()` callback.

If the application session timer is active and the expiration timeout is reached, then the `sessionExpired()` method of any `SipApplicationSessionListener` implementation is called.

6.1.2.2 SipApplicationSession Invalidation

There are two ways in which a `SipApplicationSession` can be invalidated:

1. Explicit invalidation Mechanism
2. Invalidate When Ready Mechanism

Once a `SipApplicationSession` object is invalidated by either the application or the container, it may no longer be used. All references to the object should be removed by the container and applications as soon as possible thus enabling invalidated `SipApplicationSession` objects to

be garbage collected. The container MUST invoke the listener callback `sessionDestroyed()` if a listener exists for both `SipSessions` and `SipApplicationSessions` when they are destroyed.

6.1.2.2.1 Explicit Invalidation Mechanism

An application may invalidate a `SipApplicationSession` at any time using the `invalidate()` method. On explicit invalidation, the container MUST purge all state for that `SipApplicationSession` from its memory. This includes the application state stored in the `SipApplicationSession` as well as all the contained protocol session objects.

The `invalidate()` method will throw an `IllegalStateException` if the `SipApplicationSession` object has already been invalidated. Invalidating a `SipApplicationSession` using the `invalidate()` method causes all the protocol sessions contained within it to be explicitly invalidated by the container. Explicit invalidation of `SipSession` objects is described in [6.2.4.1 SipSession Invalidation](#).

6.1.2.2.2 Invalidate When Ready Mechanism

The explicit invalidation mechanism described above causes containers to invalidate `SipSessions` immediately, which could result in partially invalidated application paths and orphaned sessions in other network entities participating in the dialog.

This specification provides a new mechanism for invalidation that applications can use to circumvent the above drawbacks called the "Invalidate When Ready" mechanism.

A `SipApplicationSession` is said to be in a ready-to-invalidate state if the following conditions are met:

1. All the contained `SipSessions` are in the ready-to-invalidate state.
2. None of the `ServletTimers` associated with the `SipApplicationSession` are active.

A `SipApplicationSession` transitions into the ready-to-invalidate state when the following conditions are met:

1. The last protocol session belonging to the `SipApplicationSession` is invalidated
2. The last `ServletTimer` associated with the `SipApplicationSession` expires.

A `SipSession` is in the ready-to-invalidate state if it can be explicitly invalidated such that the SIP state is terminated cleanly across all the SIP network entities participating in the dialog. Refer to [6.2.4.2 Important Semantics](#) for details on when container determines a `SipSession` is in the ready-to-invalidate state.

This specification introduces new methods to help applications invalidate `SipApplicationSessions` cleanly. The methods introduced are:

1. `isReadyToInvalidate()` - returns true if the `SipApplicationSession` is in the ready-to-invalidate state and false otherwise.
2. `setInvalidateWhenReady(boolean flag)` - allows applications to indicate to the container to notify it when the `SipApplicationSession` is in the ready-to-invalidate state. The container notifies the application using the `SipApplicationSessionListener.sessionReadyToInvalidate(SipApplicationSessionEvent se)` callback method.
3. `getInvalidateWhenReady()` - returns true if the container will notify the application when the `SipApplicationSession` is in the ready-to-invalidate state.

An application willing to invalidate a `SipApplicationSession` cleanly could use the callback mechanism to perform any application clean up before the `SipApplicationSession` gets invalidated by the container.

Servlets can register for `sessionReadyToInvalidate` notifications on the `SipApplicationSessionListener` interface. In the `sessionReadyToInvalidate` callback method, an application may choose to invalidate the `SipApplicationSession` or perform any other cleanup activities. If the application does not explicitly invalidate the `SipApplicationSession` in the callback or has not defined a listener, the container will invalidate the `SipApplicationSession`.

Applications may also use the callback to call `setInvalidateWhenReady(false)` to indicate to the container to not observe this `SipApplicationSession` anymore. In this case, the containers MUST not invalidate the `SipApplicationSession` after the callback. Applications could then either rely on explicit invalidation mechanism or again call `setInvalidateWhenReady(true)`. This parallels the expiry callback mechanism defined above in [6.1.2.1 SipApplicationSession Timer Operation and SipApplicationSession Expiration](#).

The firing of the `SipApplicationSession` expiry timer influences the lifetime of a `SipApplicationSession` and overrides the behavior of a `SipApplicationSession` marked with `invalidateWhenReady(true)`. If the `SipApplicationSession` times out when it is not yet ready to be invalidated state, an application could detect it in the `sessionExpired` callback of the `SipApplicationSessionListener` and extend the lifetime of the `SipApplicationSession` using `setExpires`. Failing to do so will cause the `SipApplicationSession` expiry to explicitly invalidate all the contained sessions and itself. Conversely, a `SipApplicationSession` that transitions to the ready-to-invalidate state may

have an active expiry timer. The container **MUST** cancel the expiry timer before it invalidates the `SipApplicationSession`.

6.1.3 Binding Attributes into a `SipApplicationSession`

A servlet can bind an object attribute into a `SipApplicationSession` by name. Any object bound into an application session is available to any other servlet that belongs to the same `ServletContext` and that handles a request identified as being a part of the same application session.

Some objects may require notification when they are placed into, or removed from, an application session. This information can be obtained by having the object implement the `SipApplicationSessionBindingListener` interface. This interface defines the following methods that will signal an object being bound into, or being unbound from, an application session.

- `valueBound`
- `valueUnbound`

The `valueBound` method must be called before the object is made available via the `getAttribute` method of the `SipApplicationSession` interface. The `valueUnbound` method must be called after the object is no longer available via the `getAttribute` method of the `SipApplicationSession` interface.

6.2 SipSession

`SipSession` objects represent point-to-point SIP relationships, either as established dialogs or in the stage before a dialog is actually established. The `SipSession` can be obtained from a `SipServletRequest` by calling the `getSession` method.


6.2.1 Relationship to SIP Dialogs

A `SipSession` represents either an actual SIP dialog in its early, confirmed, or terminated state defined in RFC 3261, or else represents a pseudo dialog. The notion of pseudo dialogs extend the definition of dialogs to have a certain well-defined meaning before a dialog is established in the RFC 3261 sense and after it has transitioned away from the early state because of a non-2xx final response being received. The `SipSession` interface embodies this notion of pseudo dialogs and because of this, some `SipSession` instances do not correspond to SIP dialogs.

Containers enforce SIP protocol restraints based on the dialog and transaction state. If it is illegal to send a message in a given state, an `IllegalStateException` is thrown by the container.


The SIP dialog state machine is shown in [Figure 6-1](#).

Figure 6-1 The SIP Dialog State Machine



The “undefined” state in Figure 6-1 is not a real state. For dialogs created as the result of an INVITE, the dialog springs into existence on receipt of a 1xx or 2xx with a To tag.

Figure 6-2 The SipSession State Machine



`SipSession` objects are in one of the four states: `INITIAL`, `EARLY`, `CONFIRMED`, or `TERMINATED`. These states represent the state of a dialog that may be associated with the `SipSession` object. A new method on the `SipSession` class is defined that allows the application to have access to the dialog state. The `SipSession` state depends not only the state of an underlying SIP dialog but also on whether the servlet has acted as a UAC, UAS or a proxy. The rules governing the state of a `SipSession` object are given below. Note that for any state transition caused by the receipt of a SIP message, the state change must be accomplished by the container before calling the `service()` method of any `SipServlet` to handle the incoming message.

1. Before a dialog creating request is received or sent on a `SipSession`, the `SipSession` state is defined to be `INITIAL`.
2. In general, whenever a non-dialog creating request is sent or received, the `SipSession` state remains unchanged. Similarly, a response received for a non-dialog creating request also

leaves the `SipSession` state unchanged. The exception to the general rule is that it does not apply to requests (e.g. BYE, CANCEL) that are dialog terminating according to the appropriate RFC rules relating to the kind of dialog.

3. If the servlet acts as a UAC and sends a dialog creating request, then the `SipSession` state tracks directly the SIP dialog state except that non-2XX final responses received in the EARLY or INITIAL states cause the `SipSession` state to return to the INITIAL state rather than going to TERMINATED.
4. If the servlet acts as a UAS and receives a dialog creating request, then the `SipSession` state directly tracks the SIP dialog state. Unlike a UAC, a non-2XX final response sent by the UAS in the EARLY or INITIAL states causes the `SipSession` state to go directly to the TERMINATED state.
5. If the servlet acts as a proxy for a dialog creating request then the `SipSession` state tracks the SIP dialog state except that non-2XX final responses received from downstream in the EARLY or INITIAL states cause the `SipSession` state to return to the INITIAL state rather than going to the TERMINATED state. This enables proxy servlets to proxy requests to additional destinations when called by the container in the `doResponse()` method for a tentative non-2XX best response. After all such additional proxy branches have been responded to and after considering any servlet created responses, the container eventually arrives at the overall best response and forwards this response upstream. If this best response is a non-2XX final response, then when the forwarding takes place, the state of the `SipSession` object becomes TERMINATED. If this best response is a 2XX final response, then the `SipSession` state becomes CONFIRMED.
6. Because setting the supervised flag to false affects only whether responses are seen for the transaction associated with the current request, the value of the supervised flag has no effect on the `SipSession` state.

An enum that defines the possible SIP dialog states is defined for use with the `SipSession` interface:

```
public enum SipSession.State {INITIAL, EARLY, CONFIRMED, TERMINATED }
```

The following new method is introduced on the `SipSession` interface to return the current SIP dialog state:

```
public SipSession.State getState()
```

This method returns one of the `SipSession.State` enum values - `INITIAL`, `EARLY`, `CONFIRMED` or `TERMINATED`. These values represent the SIP dialog related state of the `SipSession` when the method is called.

The standard SIP rules governing when a second or subsequent response cause a single request to establish multiple dialogs hold unmodified for `SipSessions`. If, for example, two 200 responses are received for an initial INVITE, the container will create a second `SipSession` on receipt of the second 200 response. This second `SipSession` will be derived from the one in which the INVITE was generated as described in [6.2.3.2 Derived SipSessions](#) below. Both `SipSessions` will then represent dialogs in the `CONFIRMED` state.

The `INITIAL` state is introduced to allow a UAC to generate multiple requests with the same Call-ID, From (including tag), and To (excluding tag), and within the same CSeq space. This is useful, for example, in the following situations:

- When a UAC receives a 3xx for a request initiated outside of a dialog and decides to retry with the Contact addresses received in the 3xx, it is recommended to reuse the same To, From and Call-ID for the new request [RFC 3261, section 8.1.3.4].
- When a UAC receives certain “non-failure” 4xx responses indicating that the request can be retried, e.g. 401, 407, 413, 415, 416, and 420 [RFC 3261, section 8.1.3.5].
- REGISTER requests sent from a UAC to the same registrar should all use the same Call-ID header field value and should increment the CSeq by one for each request sent [RFC 3261, section 10.2].
- When a UAC using the session timer extension receives a 422 response to an initial INVITE it retries with the same Call-ID and a higher Min-SE value [timer].

These examples have in common, a need to create similar requests without an established dialog being in place. There may well be other scenarios where it’s desirable to correlate non-dialog requests by Call-ID and ensuring proper sequencing by using the CSeq header field. A new request can be created from the `INITIAL` state only when there is no ongoing transaction. The request URI should be changed if the request is to be sent to different target than previous request.

Note that the “pseudo dialog” semantics presented here is defined for use in UAC applications only. Containers treat incoming requests as subsequent requests, i.e., routes to existing sessions, only if those requests belong to an actual established SIP dialog. There is no expectation, for example, that a container treat an incoming INVITE as a subsequent request after it has previously sent a 3xx response to another INVITE with the same Call-ID, CSeq, and From (including tag).

6.2.1.1 Cancel Message Processing

Whether a servlet acts as a proxy or as a UAS, receiving a CANCEL request does not in itself cause a `SipSession` state change. However, since receiving a CANCEL request causes the UAS to respond to an ongoing INVITE transaction with a non-2XX (specifically, 487) response, the `SipSession` state normally becomes TERMINATED as a result of the non-2XX final response sent back to the UAC.

6.2.2 Maintaining Dialog State in the SipSession

UAs use the `SipSession` method `createRequest` to create subsequent requests in a dialog. This implies that containers must associate SIP dialog state with `SipSessions` when the application acts as a UA. The dialog state is defined in RFC 3261 and consists of the following pieces of data: local/remote URIs, local/remote tags, local/remote sequence numbers, route set, remote target URI, and the secure flag.

6.2.2.1 When in a Dialog

For applications acting as UAs, the `SipSession` must track dialog state according to RFC 3261. Proxies cannot generate new requests of their own and so `SipSession.createRequest` results in an `IllegalStateException`. Similarly, requests created by applications acting as UAs cannot be proxied. However, the methods `getCallId`, `getLocalParty` and `getRemoteParty` must be implemented. `getLocalParty` returns the address of the caller (value of the From header of the initial request in the dialog) and `getRemoteParty` returns the address of the callee. This implies that Call-ID, From, and To must be associated with “proxy” `SipSessions` along with any application state.

6.2.2.2 When in the INITIAL SipSession State

When a UAC or UAS transitions from the early state to the initial state in Figure 6-2, the dialog state maintained in the `SipSession` is updated as follows (see RFC 3261 for the definition of these dialog state components):

- the remote target is reset to the remote URI
- the remote tag component is cleared
- the remote sequence number is cleared (e.g. set to -1)
- the route set is cleared
- the “secure” flag is set to false

As a consequence of these rules, requests generated in the same `SipSession` have the following characteristics:

- they all have the same Call-ID
- they all have the same From header field value including the same non-empty tag
- the CSeq of requests generated in the same `SipSession` will monotonically increase by one for each request created, regardless of the state of the `SipSession`
- all requests generated in the same `SipSession` in the initial state have the same To header field value which will not have a tag parameter
- `SipSession` objects in the initial state have no route set, the remote sequence number is undefined, the “secure” flag is false, and the remote target URI is the URI of the To header field value.

6.2.3 Creation of SipSessions

A big difference between dialogs and `SipSessions` is that whereas SIP requests may exist outside of a dialog (for example, OPTIONS and REGISTER), in the SIP Servlet API all messages belong to a `SipSession`.

SIP dialogs get created as a result of non-failure responses being received to requests with methods capable of setting up dialogs. The baseline SIP specification defines one such method, namely INVITE. Dialog handling is complicated by the fact that proxies may fork. This means a single request, for example an INVITE, can be accepted at multiple UASs, thus causing multiple dialogs to be created. In such cases, the UAC SIP servlet will see multiple `SipSessions`.

The relationship between `SipSessions` and SIP dialogs can be summed up as follows:

When an initial request results in one dialog being set up, the `SipSession` of the initial `SipServletRequest` will correspond to that dialog. When more than one dialog is established, the first one will correspond to the existing `SipSession` and for each subsequent dialog a new `SipSession` is created in the manner defined in [6.2.3.2 Derived SipSessions](#).

The rules for when `SipSessions` are created are:

- Locally generated initial requests created via `SipFactory.createRequest` belong to a new `SipSession`.
- Incoming initial requests belong to a new `SipSession`.

- Responses to an initial request that do not establish a dialog, belong to the “original” `SipSession` of the request.
- The first message that establishes a SIP dialog (for example, a 2xx to an initial INVITE request) is associated with the original `SipSession` as are all other messages belonging to that dialog.
- Subsequent messages that establish dialogs are associated with new `SipSessions` derived from the original `SipSession`, see [6.2.3.2 Derived SipSessions](#) below.

When a proxying application receives a response that is associated with a new `SipSession` (say, because it is a second 2xx response to an INVITE request), the `getSession` method of that response returns a new derived `SipSession`. The original `SipSession` continues to be available through the original request object — itself available through the `getOriginalRequest` method on the `Proxy` interface.

6.2.3.1 Extensions Creating Dialogs

Extensions to the baseline SIP specification may define additional methods capable of establishing dialogs. The SIP event framework is one such extension [RFC 3265]. In the event framework, a 2xx response to an initial SUBSCRIBE establishes a dialog. In fact, a matching NOTIFY request that arrives before the 2xx response for SUBSCRIBE will also establish a dialog.

The other SIP extension introducing a dialog creating method as of this writing is [RFC 3515] SIP REFER.

A SIP servlet container that “understands” an extension capable of establishing dialogs should create new derived `SipSession` objects for the second and subsequent dialogs created as a result of sending a single request capable of establishing multiple dialogs.

6.2.3.2 Derived SipSessions

A derived `SipSession` is essentially a copy of the `SipSession` associated with the original request. It is constructed at the time the message creating the new dialog is passed to the application. The new `SipSession` differs only in the values for the tag parameter of the address of the callee (this is the value used for the To header in subsequent outgoing requests) and possibly the route set. These values are derived from the dialog-establishing message as defined by the SIP specification. The set of attributes in the cloned `SipSession` is the same as that of the original—in particular, the values are not cloned.

New `SipSessions` corresponding to the second and subsequent 2xx responses (or 1xx responses with To tags) are available through the `getSession` method on the `SipServletResponse`. The

“original” `SipSession` of the request continues to be available through the original request object.

6.2.4 SipSession Lifetime

`SipSessions` do not have a timeout value of their own separate from that of the parent `SipApplicationSession`. Rather, when the parent session times out, all child protocol sessions time out with it.

In general, the lifetime of a `SipSession` can be controlled in the following ways:

1. The parent `SipApplicationSession` times out or is invalidated explicitly and this invalidates all child protocol sessions.
2. Application explicitly invalidates the `SipSession` using the `invalidate()` API.
3. Application marks the `SipSession` to be invalidated and the container invalidates it when the session is in the ready-to-invalidate state.

Any attempt to retrieve or store data on an invalidated `SipSession` causes an `IllegalStateException` to be thrown by the container as does any call to `createRequest`.

When a `SipSession` terminates, either because the parent application session timed out or because the `SipSession` was explicitly invalidated, the container **MUST** purge all state of that `SipSession` from its memory. In such a case, if a subsequent request or response belonging to the corresponding dialog is received, the container is free to handle it in either of the following ways:

1. Reject the request by sending a 481 error response or
2. Route the request or response without application involvement

It is quite possible that different application instances on the same application path have different lifetimes. Containers handling subsequent requests and responses on a dialog corresponding to a partially invalidated application path would invoke all applications up to the first invalidated session instance and reject them if the session acts as an UA or proxy them statelessly if the session acts as a Proxy.

6.2.4.1 SipSession Invalidation

There are two ways in which a `SipSession` can be invalidated:

1. Explicit Invalidation Mechanism

2. Invalidate When Ready Mechanism

Once a `SipSession` object is invalidated either by the application or the container, it may no longer be used. On invalidation, the container **MUST** invoke the `sessionDestroyed()` callback on implementations of the `SipSessionListener` interface, if any exist.

6.2.4.1.1 Explicit Invalidation Mechanism

An application may invalidate a `SipSession` at any time using the `invalidate()` method. On explicit invalidation, the container **MUST** purge all state of that `SipSession` from its memory. The state would include both application state as well as container state such as the underlying transactions.

The `invalidate()` method called on a `SipSession` object will throw an `IllegalStateException` if that `SipSession` object has already been invalidated. Invalidating a `SipSession` using the `invalidate()` method causes the parent `SipApplicationSession` to have one less session.

6.2.4.1.2 Invalidate When Ready Mechanism

The explicit invalidation mechanism described above causes containers to clean up session state immediately and may suffer from the following drawbacks:

1. A broken or partially invalidated application path.
2. Orphaned sessions in network entities participating in the dialog.

This specification provides a new mechanism for invalidation that applications can use to circumvent the above drawbacks called the “Invalidate When Ready” mechanism.

A `SipSession` is said to be in a ready-to-invalidate state only when it can be explicitly invalidated such that the SIP state is terminated cleanly across all the SIP network entities participating in the dialog. A `SipSession` is invalidated cleanly in this state, therefore application developers invalidating `SipSessions` in the ready-to-invalidate state will not cause sessions in the application path to be orphaned.

This specification introduces new methods to help applications invalidate `SipSessions` cleanly across all network entities participating in a dialog. The methods introduced are:

1. `isReadyToInvalidate()` - returns true if the `SipSession` is in the ready-to-invalidate state and false otherwise. Applications can use this method to extend the lifetime of the parent `SipApplicationSession` if it times out before the `SipSession` is in the ready-to-invalidate state.

2. `setInvalidateWhenReady(boolean flag)` - allows applications to indicate to the container to notify it when the `SipSession` is in the ready-to-invalidate state. The container notifies the application using the `SipSessionListener.sessionReadyToInvalidate(SipSessionEvent se)` callback.
3. `getInvalidateWhenReady()` - returns true if the container will notify the application when the `SipSession` is in the ready-to-invalidate state.

An application willing to invalidate a `SipSession` cleanly could use the callback mechanism to perform any other application clean up before the `SipSession` gets invalidated by the container.

Servlets can register for `sessionReadyToInvalidate` notifications on the `SipSessionListener` interface. In the `sessionReadyToInvalidate` callback method, an application may choose to invalidate the session and perform any other cleanup activities. If the application does not explicitly invalidate the session in the callback or has not defined a listener, the container will invalidate the session.

Applications may call `setInvalidateWhenReady(false)` at any time to indicate to the container to not observe this session anymore. In such a case, the containers MUST not invalidate the session after the callback. Applications could then either rely on explicit invalidation mechanism or again call `setInvalidateWhenReady(true)`. This parallels the expiry callback mechanism defined above in [6.1.2.1 SipApplicationSession Timer Operation and SipApplicationSession Expiration](#).

For example: An application registers a `SipSession` for a callback using `session.setInvalidateWhenReady(true)` and handles the callback in the `SipSessionListener` as follows:

```
@SipListener
public class MySessionListener implements SipSessionListener {
    ...
    sessionReadyToInvalidate(SipSessionEvent se) {
        Session sess = se.getSession();

        // directs the container to stop observing this session and
        // not invalidate it when this callback returns.
        sess.setInvalidateWhenReady(false);
    }
    ...
}
```

The container determines the `SipSession` to be in the ready-to-invalidate state under any of the following conditions:

1. A dialog corresponding to a `SipSession` terminates when the `SipSession` transitions to the `TERMINATED` state. If derived sessions (sub-dialogs) were created, it is up to container implementations to declare the original `SipSession` to be in ready-to-invalidate state either when the original `SipSession` (primary dialog) terminates or when all the derived sessions (sub-dialogs) terminate.
2. A `SipSession` transitions to the `CONFIRMED` state when it is acting as a non-record-routing proxy.
3. A `SipSession` acting as a UAC transitions from the `EARLY` state back to the `INITIAL` state on account of receiving a non-2xx final response ([6.2.1 Relationship to SIP Dialogs](#), point 4) and has not initiated any new requests (does not have any pending transactions).

The container **MUST NOT** treat `SipSession` objects to be in the ready-to-invalidate state if the session has any underlying transactions in progress.

The firing of the parent `SipApplicationSession` timer influences the lifetime of a child `SipSession` and overrides the behavior of a `SipSession` marked with `invalidateWithReady(true)`. If the parent `SipApplicationSession` times out when it contains a child `SipSession` that is not yet ready to be invalidated, an application could detect it in the `sessionExpired` callback of the `SipApplicationSessionListener` and extend the lifetime of the `SipApplicationSession` using `setExpires`. Failing to do so will cause the `SipApplicationSession` expiry to forcefully invalidate the child `SipSession` even if it is not yet ready to be invalidated.

For example, to extend the lifetime of a `SipApplicationSession` when one of its child sessions is not yet in the ready-to-invalidate state, one could handle the callback in the `SipApplicationSessionListener` as follows:

```
@SipListener
public class MySipApplicationSessionListener
    implements SipApplicationSessionListener {
    ...
    sessionExpired(SipApplicationSessionEvent sase) {
        SipApplicationSession sas = sase.getApplicationSession();
        Iterator sessions = sas.getSessions("SIP");
        while (sessions.hasNext()) {
            SipSession ss = (SipSession) sessions.next();
            if (! ss.isReadyToInvalidate()) {
                sas.setExpires(extensionTime);
                return;
            }
        }
    }
}
```

```
...
}
```

If an application invokes any method on an invalidated `SipSession` object, the container **SHOULD** throw an `IllegalStateException`. Note that the methods introduced on the `SipSession` and `SipApplicationSession` classes in this specification throw `IllegalStateException` when invoked against invalid session objects. However, some of the existing methods (from v1.0) which did not throw this exception will continue to not throw it in this specification for backwards-compatibility reasons. A new `isValid()` method is provided on both `SipApplicationSession` and `SipSession` interfaces to check whether the session object has been invalidated.

6.2.4.2 Important Semantics

This specification introduces the Invalidate When Ready mechanism for session invalidation. The `invalidateWhenReady` flag is true by default for both `SipApplicationSessions` and `SipSessions` for applications written compliant to v1.1 (or later) of this specification. This means that by default, all sessions are observed by the container and are invalidated automatically after they transition to the ready-to-invalidate state. This automatic cleanup helps in releasing resources as soon as possible and helps containers in managing their resources efficiently.

The `invalidateWhenReady` flag is false by default for both `SipApplicationSessions` and `SipSessions` for applications written compliant to v1.0 of this specification. Hence, sessions belonging to v1.0 applications will not be subject to the automatic session cleanup. By default, such v1.0 applications should see no change in session lifetime behavior on containers compliant with this specification. v1.0 applications that want the automatic session cleanup must explicitly call `setInvalidateWhenReady(true)` on each session. These applications may also implement the `sessionReadyToInvalidate()` callback on the session listeners.

If applications use the explicit invalidation mechanism for session invalidation, it is recommended that applications implemented a SIP dialog cleanup mechanism using the application composition function ([15 Application Selection And Composition Model](#)). The deployer may choose to always deploy a B2BUA application in the application path which implements the desired deployer policy as regards the duration of SIP dialogs. For example, in the simplest case, such an application could set a timer when a SIP call is first established through it. On expiration of this timer, this B2BUA application would tear down the call by appropriately sending termination SIP messages (BYE) on both of its dialogs.

6.2.5 The RequestDispatcher Interface

There is a need for a structuring mechanism, that allows applications to consist of multiple servlets that handle various parts of the application. In the case of HTTP, servlets may use a `RequestDispatcher` to forward a request to another servlet or to include the output of another servlet in its own response. SIP servlets can use the `RequestDispatcher` interface the same way—as a mechanism for passing a request or response to another servlet in the same application using the `forward` method and can be thought of as an “indirect method call”. (In the case of SIP, the name “`RequestDispatcher`” is a bit of a misnomer as it applies equally well to responses and requests). The `include` method has no meaning for SIP servlets.

6.2.6 The SipSession Handler

The `RequestDispatcher` interface provides one structuring mechanism for SIP servlets. Additionally, the `SipSession` interface has a notion of a handler—a reference to a servlet that is effectively the callback interface through which the container dispatches all incoming events related to that `SipSession` to the application.

When a servlet application creates a new initial request, a `SipSession` springs into existence. At this point, the container assigns the applications one servlet to be the handler for the newly created `SipSession`. This servlet selection is described in the [16 Mapping Requests To Servlets](#). The `SipSession` handler will be invoked to handle responses for all requests as well as new incoming requests. When a `SipSession` is created as part of processing an initial incoming request, the servlet object invoked becomes the initial handler for that `SipSession`.

Regardless of how a `SipSession` was created, the application may subsequently specify that a different servlet within the same application should become the handler for a particular `SipSession`. This is done by calling the `SipSession` method:

```
void setHandler(String name);
```

The argument is the same as in the call to `ServletContext.getNamedDispatcher`, i.e., it is a servlet name. Servlet names are typically associated with servlets through the deployment descriptor or using the `@SipServlet(name)` annotation.

6.2.7 Binding Attributes into a SipSession

A servlet can bind an object attribute into a `SipSession` by name. Any object bound into a session is available to any other servlet that belongs to the same `ServletContext` and that handles a request identified as being a part of the same session.

Some objects may require notification when they are placed into, or removed from, a session. This information can be obtained by having the object implement the `SipSessionBindingListener` interface. This interface defines the following methods that will signal an object being bound into, or being unbound from, a session.

- `valueBound`
- `valueUnbound`

The `valueBound` method must be called before the object is made available via the `getAttribute` method of the `SipSession` interface. The `valueUnbound` method must be called after the object is no longer available via the `getAttribute` method of the `SipSession` interface.

6.3 Last Accessed Times

The `getLastAccessedTime` method of the `SipApplicationSession` and `SipSession` interfaces allows a servlet to determine the last time a session was accessed before the current message was handled. A session is considered to be accessed when a message that is part of the session is handled by the servlet container. The last accessed time of a `SipApplicationSession` will thus always be the most recent last accessed time of any of its contained protocol sessions.

Whenever the last accessed time for a `SipApplicationSession` is updated, it is considered refreshed i.e., the expiry timer for that `SipApplicationSession` starts anew.

6.4 Important Session Semantics

6.4.1 Message Context

When passed to applications, SIP servlet request and response objects are always associated with a `SipSession` object which, in turn, belong to a `SipApplicationSession`. This specification refers to the combination of application, `SipApplicationSession`, and `SipSession` as the message context. When routing subsequent requests, containers must determine which applications to invoke and also the context in which to invoke them. When multiple applications execute on the same request, they will do so in different message contexts—each application has its own session objects which are independent of those of other applications. Application data placed into a session object by application A is not accessible to application B.

When applications are invoked to process subsequent requests, the message context will be identical to that of the initial request. Suppose, for example, that an application is invoked to handle an INVITE and proxies it with record-routing enabled. When a BYE for the same dialog

is received, the same application is invoked and in the same context, that is, `getApplicationSession` and `getSession` are required to return the same session objects when invoked on the INVITE and BYE requests as well as on corresponding response objects.

Note that while logically all SIP messages passed to applications have `SipApplicationSession` and `SipSession` objects associated with them, for performance reasons, containers may choose to defer creation of session objects when applications cannot observe the difference.

6.4.2 Threading Issues

This specification does not mandate any particular threading model.

Containers may invoke a single SIP servlet application in more than one concurrent thread, and these threads may have active access to the same `SipSession` and `SipApplicationSession` objects at the same time. In such container implementations, the application developer has the responsibility for synchronizing access (possibly with container-specific means) to these session objects and any attributes and objects held by these session objects, as appropriate.

Alternatively, container implementations may provide a specific threading model (pertaining to access to session objects), an explicit locking API or implicit serialized access semantics. In such cases, the container behavior should be documented in sufficient detail to facilitate application thread safety and porting of applications between different vendors.

6.4.3 Distributed Environments

For an application marked as distributable the container must be able to handle all objects placed into `SipSession` and `SipApplicationSession` (generically called sessions in this section) using the `setAttribute` method according to the following rules.

- The container **MUST** accept objects that implement the `Serializable` interface.
- The container **MUST** accept object of the type `SipServletMessage`.
- The container **MAY** choose to support storage of other objects in the `SipSession`, such as references to Enterprise Java Beans and transactions.
- Migration of sessions will be handled by container-specific facilities.

The servlet container **MAY** throw an `IllegalArgumentException` if an object is placed into the sessions that is not `Serializable` or for which specific support has not been made available. The `IllegalArgumentException` **MUST** be thrown for objects where the container cannot support the mechanism necessary for migration of a session storing them.

These restrictions mean that the developer is ensured that there are no additional concurrency issues beyond those encountered in a non-distributed container.

The container provider can ensure scalability and quality of service features like load-balancing and failover by having the ability to move a session object and its contents from any active node of the distributed system to a different node of the system.

If distributed containers persist or migrate sessions to provide quality of service features, they are not restricted to using the native JVM serialization mechanism for serializing session and their attributes, however the containers **MUST** provide consistent environment after de-serialization as one would expect Java serialization mechanism to provide. Developers are not guaranteed that containers will call `readObject()` and `writeObject()` methods on session attributes if they implement them, but are guaranteed that the `Serializable` closure of each attribute will be preserved. Defining closure around individual attributes mean that each individual attribute can be saved independently without having to save the entire `SipSession` or `SipApplicationSession` on an attribute addition or update, of course the shared objects accross the attributes shall be referenced as in normal Java serialization mechanism.

Containers **MUST** notify any session attributes implementing the `SipSessionActivationListener` and `SipApplicationSessionActivationListener` during migration of a session. They **MUST** notify listeners of passivation prior to serialization of a session, and of activation after deserialization of a session when such a migration takes place.

Developers writing distributed applications should be aware that since the container may run in more than one Java VM, the developer cannot depend on static or instance variables for storing application states. They should store such states using an EJB or a database.

In addition to this the containers **MUST** also ensure that the `Serializable` "info" object associated with the `ServletTimer` is also migrated or persisted alongside the sessions if the `ServletTimer` is chosen to be persistent.

7 SIP Servlet Applications

A SIP servlet application is a collection of servlets, class files, and other resources that comprise a complete application on a SIP server. The SIP application can be bundled and run on multiple containers from multiple vendors.

By default, an instance of a SIP application must run on one VM at any one time. This behavior can be overridden if the application is marked as “distributable” via its deployment descriptor or via the `@SipApplication(distributable)` annotation. An application marked as distributable must obey a more restrictive set of rules than is required of a normal servlet application. These rules are set out throughout this specification.

7.1 Relationship with HTTP Servlet Applications

The directory structure used for SIP servlet applications is very similar to the one defined for HTTP servlets, in particular, the deployment descriptor, class files, libraries etc. of SIP servlet applications exists underneath the `WEB-INF/` directory. This allows converged HTTP and SIP servlet applications to be packaged in a single archive file and deployed as a unit without having any part of the SIP application be interpreted as content to be published by the Web server.

7.2 Relationship to ServletContext

The servlet container must enforce a one-to-one correspondence between a servlet application and a `ServletContext`. A `ServletContext` object provides a servlet with its view of the application. This is true for converged SIP and HTTP applications, also: all servlets within an application see the same `ServletContext` instance.

7.3 Elements of a SIP Application

A SIP application may consist of the following items:

- Servlets
- Utility classes
- Static resources and content (text and speech announcements, configuration files, etc.)
- Descriptive meta information which ties all of the above elements together.

7.4 Deployment Hierarchies

This specification defines a hierarchical structure used for deployment and packaging purposes that can exist in an open file system, in an archive file, or in some other form. It is recommended, but not required, that servlet containers support this structure as a runtime representation.

7.5 Directory Structure

A SIP application exists as a structured hierarchy of directories. As mentioned above, for compatibility with the HTTP servlet archive file format, the root of this hierarchy serves as the document root for files intended to be published from a Web server. This feature is only actually used for combined HTTP and SIP servlet applications.

A special directory exists within the application hierarchy named “WEB-INF”. This directory contains all things related to the application that are not in the document root of the application. For SIP-only applications, everything typically resides under WEB-INF. The WEB-INF node is not part of the public document tree of the application. No file contained in the WEB-INF directory may be served directly to a client by the container. However, the contents of the WEB-INF directory are visible to servlet code using the `getResource()` and `getResourceAsStream()` method calls on the `ServletContext`. Any application specific configuration information that the developer needs access to from servlet code yet does not wish to be exposed to a web client may be placed under this directory. Since requests are matched to resource mappings case-sensitively, client requests for `‘/WEB-INF/foo’`, `‘/WEB-INF/foo’`, for example, should not result in contents of the web application located under `/WEB-INF` being returned, nor any form of directory listing thereof.

The contents of the WEB-INF directory are:

- The `/WEB-INF/sip.xml` deployment descriptor.

- The /WEB-INF/classes/* directory for servlet and utility classes. The classes in this directory are available to the application classloader.
- The /WEB-INF/lib/* .jar area for Java Archive files. These files contain servlets, beans, and other utility classes useful to the web application. The web application classloader can load class from any of these archive files.

The application classloader must load classes from the WEB-INF/ classes directory first, and then from library JARs in the WEB-INF/ lib directory.

7.6 Application Names

In order to identify SIP servlet applications, a new mandatory element, `<app-name>`, is defined under the `<sip-app>` element in the sip.xml deployment descriptor file of each application. The names of Sip Servlet applications must be unique within a container instance or across clustered containers under common administrative control for application selection to function properly. It is recommended that application developers follow the Java class naming convention when naming applications, e.g. "org.sipservlet.app.voicemail", to avoid naming conflicts with other developers. The container is responsible for detecting and handling naming conflict, when a new application is deployed. Though how the application is identified within the container is something internal to it and is like a symbol, it is strongly recommended that containers make use of the deployment context while identifying the applications. Specifically the containers **SHOULD** create symbols identifying the application as

"myj2eeapp/mysipapp/org_sipservlet_app_voicemail" where the voice mail application with the following deployment descriptor is packaged in a war archive by the name of "mysipapp.war" which is packaged in an ear archive by the name of "myj2eeapp.ear".

Listing 7-1 Example of sip.xml file illustrating app-name and servlet-name

```

<sip-app>
  <app-name>org.sipservlet.app.voicemail</app-name>
  ...
  <servlet>
    <servlet-name>depositServlet</servlet-name>
    <servlet-
class>org.sipservlet.app.voicemail.DepositServlet</servlet-class>
    ...
  </servlet>
  <servlet>
    <servlet-name>retrievalServlet</servlet-name>
    <servlet-

```

```
class>org.sipservlet.app.voicemail.RetrievalServlet</servlet-  
class>  
    ...  
    </servlet>  
    ...  
</sip-app>
```

If the aforementioned application was a v1.0 style application with similar packaging and were to be deployed on a container compliant with this specification, the application name **SHOULD** be "myj2eeapp/mysipapp".

The containers **MAY** further distinguish these names by adding their own versioning tokens.

If the application is specified by annotation, chapter [18 Java Enterprise Edition 5 Container](#) provides the procedures for determining the name of such applications.

7.6.1 Example Application Directory Structure

The following is a listing of all the files in a sample SIP servlet application:

```
/WEB-INF/sip.xml  
/WEB-INF/wakeup.wav  
/WEB-INF/lib/foo.jar  
/WEB-INF/classes/WakeupServlet.class
```

The following is the same example but with an HTTP servlet component included:

```
/wakeup.html  
/register.html  
/WEB-INF/sip.xml  
/WEB-INF/web.xml  
/WEB-INF/wakeup.wav  
/WEB-INF/lib/foo.jar  
/WEB-INF/classes/RegisterWakeupCall.class  
/WEB-INF/classes/WakeupServlet.class
```

In this latter case there are separate SIP and HTTP deployment descriptors: sip.xml and web.xml, respectively. Resources not under WEB-INF/ will be part of the content served from the HTTP part of the converged application.

7.7 Servlet Application Archive File

Servlet applications can be packaged and signed into a Servlet Archive format (sar) file using the standard Java Archive tools. For example, a click-to-dial application might be distributed in an archive file called click2dial.sar.

When packaged into such a form, a META-INF directory will be present, which contains information useful to Java Archive tools. This directory cannot be directly served as content by the container in response to a web client's request, though its contents are visible to servlet code via the `getResource()` and `getResourceAsStream()` calls on the `ServletContext`.

Such an archive deployed to a Java EE 5 compliant servlet container may reference a .jar file or directory by naming the referenced .jar file or directory in a Class-Path header in the referencing JAR file's Manifest file as per chapter 8 of the Java EE 5 Specification. The referenced .jar file or directory is named using a URL relative to the URL of the referencing JAR file.

Note that the .sar archive format is modeled on .war archive format and since the .sar archives can contain the web application components, a .war archive should also be able to contain the sip application components. It should not matter to the converged container as the contents of the archive determine what components are in them. This specification establishes the equivalence of .sar and .war archive formats in the context of SIP Servlet containers. For the purpose of discussion in this specification the archive file is called Servlet Archive or SAR for short.

7.8 SIP Application Configuration Descriptor

The following are types of configuration and deployment information in the SIP application deployment descriptor (the Deployment Descriptor schema is presented in [19 Deployment Descriptor](#)):

- ServletContext init parameters
- Session configuration
- Servlet definitions
- Security

7.9 Dependencies On Extensions

When a number of applications make use of code or resources, they will typically be installed as library files in the container in current implementations of servlet containers. These files are often common or standard API that can be used without portability being sacrificed. Files used only by

one or a few applications will be made part of the servlet application to be available for access. Application developers need to know what extensions are installed on a servlet container, and containers need to know what dependencies on such libraries servlets in a SAR/WAR may have, in order to preserve portability.

Servlet containers are recommended to have a mechanism by which servlet applications can learn what JAR files containing resources and code are available, and for making them available to the application. Containers should provide a convenient procedure for editing and configuring library files or extensions.

It is recommended that application developers provide a META-INF/MANIFEST.MF entry in the SAR/WAR file listing extensions, if any, needed by the SAR/WAR. The format of the manifest entry should follow standard JAR manifest format. In expressing dependencies on extensions installed on the servlet container, the manifest entry should follow the specification for standard extensions defined at <http://java.sun.com/j2se/1.5.0/docs/guide/extensions/versioning.html>.

Servlet containers should be able to recognize declared dependencies expressed in the optional manifest entry in a SAR/WAR file, or in the manifest entry of any of the library JARs under the WEB-INF/lib entry in a SAR/WAR. If a servlet container is not able to satisfy the dependencies declared in this manner, it should reject the application with an informative error message.

7.10 Servlet Application Classloader

The classloader that a container uses to load a servlet in an application archive must not allow the archive to override J2SE or Java servlet API classes. It is further recommended that the loader not allow servlets in the archive access to the servlet containers implementation classes.

It is recommended, however, that the application classloader be implemented so that classes and resources packaged within the SAR/WAR are loaded in preference to classes and resources residing in container-wide library JARs.

7.11 Replacing a Servlet Application

A server should be able to replace an application with a new version without restarting the container. When an application is replaced, the container should provide a robust method for preserving session data within that application.

7.12 Servlet Application Environment

Java EE defines a naming environment that allows applications to easily access resources and external information without the explicit knowledge of how the external information is named or organized.

As servlets are an integral component type of Java EE, provision has been made in the SIP servlet application deployment descriptor for specifying information allowing a servlet to obtain references to resources and enterprise beans. The deployment elements that contain this information are:

- `env-entry`
- `ejb-ref`
- `resource-ref`

The `env-entry` element contains information to set up basic environment entry names relative to the `java:comp/env` context, the expected Java type of the environment entry value (the type of object returned from the JNDI lookup method), and an optional environment entry value. The `ejb-ref` element contains the information needed to allow a servlet to locate the home interfaces of an enterprise bean. The `resource-ref` element contains the information needed to set up a resource factory.

The requirements of the Java EE environment with regards to setting up the environment are described in chapter 5 of the Java™ Platform, Enterprise Edition 5 (Java EE 5) Specification. Servlet containers that are not part of a Java EE compliant implementation are encouraged, but not required, to implement the application environment functionality described in the Java EE specification. If they do not implement the facilities required to support this environment, upon deploying an application that relies on them, the container should provide a warning.

Since Servlet specification v2.5 introduced the usage of annotations based resource injection and this specification is based on it, the annotations defined in Servlet specification v2.5 are supported, like `@Resource`, `@EJB` etc. see [18.2.1 Servlet 2.5 alignment](#) for details.

8 Application Listeners and Events

Version 2.3 of the Java Servlet Specification introduced the notion of application listeners [Servlet API, chapter 10]. The SIP Servlet API requires full support for application listeners and events as defined in that document, except for the HTTP specific events defined in package `javax.servlet.http`. The servlet context events defined in `javax.servlet` must be supported. Additionally, this specification defines some SIP specific events for which support is also required.

For full details on application listeners and events, see [Servlet API, chapter 10].

8.1 SIP Servlet Event Types and Listener Interfaces

The SIP specific events types and the listener interfaces used to monitor them are as follows.

- Listeners on `SipApplicationSession`:
 - `javax.servlet.sip.SipApplicationSessionListener`: Implementations of this interface will receive notifications when a `SipApplicationSession` has been created, destroyed, has timed out or is in the ready-to-invalidate state.
 - `javax.servlet.sip.SipApplicationSessionAttributeListener`: Implementations of this interface will receive notifications when attributes are added, removed or replaced from a `SipApplicationSession`.
 - `javax.servlet.sip.SipApplicationSessionBindingListener`: Attributes implementing this interface will receive notifications when they are bound or unbound from a `SipApplicationSession`.

- `javax.servlet.sip.SipApplicationSessionActivationListener`:
Implementations of this interface will receive notifications when a `SipApplicationSession` has been activated or passivated.
- Listeners on `SipSession`:
 - `javax.servlet.sip.SipSessionListener` : Implementations of this interface will receive notifications when a `SipSession` has been created, destroyed or is in the ready-to-invalidate state.
 - `javax.servlet.sip.SipSessionActivationListener` : Implementations of this interface will receive notifications when a `SipSession` has been activated or passivated.
 - `javax.servlet.sip.SipSessionAttributeListener`: Implementations of this interface will receive notifications when attributes are added, removed or replaced from a `SipSession`.
 - `javax.servlet.sip.SipSessionBindingListener`: Attributes implementing this interface will receive notifications when they are bound or unbound from a `SipSession`.
- Error listener:
 - `javax.servlet.sip.SipErrorListener`: Implementations of this interface will receive notification when an expected ACK or PRACK is not received.
- Timer listener:
 - `javax.servlet.sip.TimerListener` : Implementations of this interface will receive notification when a `ServletTimer` has fired.
- `SipServlet` listener:
 - `javax.servlet.sip.SipServletListener` : Implementations of this interface will receive notification on `SipServlet` initialization.

9 Timer Service

The timer service is a container-provided service that allows applications to schedule timers and to receive notifications when timers expire. Timers are managed by the container like other resources, and may optionally be persistent in which case they are stored along with session data. Timers can be scheduled to expire once after a specified time, or repeatedly at specified intervals.

The timer support consists of three interfaces: `TimerService` which is used when creating timers, `ServletTimer` which represents timers and is passed to callbacks, and `TimerListener` which is the callback interface implemented by the application and invoked by the container on timer expiration.

9.1 TimerService

The `TimerService` interface is implemented by containers and is made available to applications as a `ServletContext` parameter with the name `javax.servlet.sip.TimerService`. The `TimerService` provides the following methods:

```
ServletTimer createTime(SipApplicationSession appSession,
                      long delay,
                      boolean isPersistent,
                      java.io.Serializable info);

ServletTimer createTime(SipApplicationSession appSession,
                      long delay,
                      long period,
                      boolean fixedDelay,
                      boolean isPersistent,
                      java.io.Serializable info);
```

`ServletTimer` objects are associated with an application session. The application may store data in the application session and retrieve it later when the timer fires. The `getTimers` method of the `SipApplicationSession` interface returns a `java.util.Collection` containing all `ServletTimer` objects currently scheduled and associated with that session. Similarly, the `getTimer(String id)` method on the `SipApplicationSession` returns the specific `ServletTimer` object that is currently scheduled and is associated with that session.

The meaning of the arguments to the `createTimer` method is:

- `appSession` – the application session that the new `ServletTimer` will be associated with
- `delay` – delay in milliseconds before the `ServletTimer` is to expire the first time
- `period` – time in milliseconds between successive timer expirations
- `fixedDelay` – specifies whether a repeating timer is fixed-delay or fixed-rate. The semantics are similar to those of `java.util.Timer`: in both cases the repeating timer expires at approximately regular intervals but with fixed-delay execution rescheduling of the repeating timer ignores any “lateness” in previous expirations whereas fixed-rate timers are rescheduled based on absolute time.
- `isPersistent` – if true the `ServletTimer` should be reinstantiated if the server is shut down and subsequently restarted. During the restart, the container will call the `TimerInterface` for a timer that has expired during the shutdown. The `SipApplicationSession` associated with the `ServletTimer` should be persistent.
- `info` – application information to be delivered along with the timer expiration. This is useful for determining the significance of a timer expiration in applications that sets multiple timers per application session.

9.2 ServletTimer

The `ServletTimer` interface represents scheduled timers. The application session, timer id and the serializable information object can be retrieved from the `ServletTimer` in the expiration callback. The `ServletTimer` interface also lets applications cancel timers and obtain information regarding last and next scheduled expiration time.

The `ServletTimer` interface provides the following methods:

```
SipApplicationSession getApplicationSession();
java.io.Serializable getInfo();
String getId();
```

```
long scheduledExecutionTime();  
long getTimeRemaining();  
void cancel();
```

9.3 TimerListener

Applications are notified of expiration of timers through the `TimerListener` interface. This interface has a single method:

```
void timeout(ServletTimer timer);
```

Applications using timers must implement the `TimerListener` interface and must declare the implementation class in a listener element of the SIP deployment descriptor (as described below) or use the `@SipListener` annotation (as described in [18.2.4 @SipListener Annotation](#)). For example, an application might include a class `com.example.MyTimerListener` that implements `javax.servlet.sip.TimerListener`. It would then declare this listener in the SIP deployment descriptor as follows:

```
<listener>  
  <listener-class>com.example.MyTimerListener</listener-class>  
</listener>
```

There may be at most one `TimerListener` implementation declared in the deployment descriptor.

Timer Service

10 Proxying

One important function of SIP is the ability to route requests, that is, for incoming requests to decide which destination or destinations should receive the request. The ability to proxy requests is essential to many SIP services. In some cases the service may have a choice between proxying and redirecting but many services require proxying because they need to see responses and/or stay on the signaling path.

One of the most important differences between the HTTP and SIP Servlet APIs is that whereas HTTP servlets execute on origin servers only and are concerned only with responding to incoming requests, SIP servlets are typically located on proxy servers and must be able to proxy incoming requests as well as respond to them directly.

Proxying is initiated and controlled via a `Proxy` object obtained from an incoming `SipServletRequest` and its associated `ProxyBranch` object. There is at most one `Proxy` object per SIP transaction, meaning that `SipServletRequest.getProxy()` will return the same `Proxy` instance whenever invoked on the original request object or on any other message belonging to that transaction.

10.1 Parameters

A number of `Proxy` parameters control various aspects of the proxying operation:

- **recurse**: flag specifying whether the servlet engine will automatically recurse or not. If recursion is enabled the servlet engine will automatically attempt to proxy to contact addresses received in redirect (3xx) responses. The default value is true.

- **recordRoute**: flag controlling whether the application stays on the signaling path for this dialog or not. This should be set to true if the application wishes to see subsequent requests belonging to the dialog. The default value is false.
- **parallel**: flag specifying whether to proxy to multiple destinations in parallel (true) or sequentially (false). In the case of parallel search, the server may proxy the request to multiple destinations without waiting for final responses to previous requests. The default value is true.
- **supervised**: whether the servlet is invoked to handle responses. Note that setting the supervised flag to false affects only the transaction to which the `Proxy` object relates. The default value is true.
- **proxyTimeout**: The timeout for the proxy in general. In case the proxy is a sequential proxy then this value behaves like the `sequential-search-timeout` which is deprecated since v1.1. In case the proxy is a parallel proxy then this timeout acts as the timeout for the entire proxy i.e each of its parallel branches before it starts to send out CANCELs waiting for final responses on all INVITE branches and sends the best final response upstream. The default value is the value of the `proxy-timeout` element of the deployment descriptor or a container specific value, if this is not specified.
- **addToPath**: flag controlling whether the application adds a Path header to the request, thereby adding itself into the Path thus created by the REGISTER requests as per RFC3327.

The `recordRoute` flag may only be set before the first call to `Proxy.proxyTo` or `Proxy.startProxy()`. Any attempt to set it afterwards results in an `IllegalStateException` being thrown. The `recurse`, `parallel`, and `supervised` flags as well as the `proxyTimeout` parameter may be modified by applications while a proxying operation is in progress.

Note: Deprecation of Stateless Proxy in SIP Servlet Specification:

The `Proxy` implemented on containers compliant with this specification MUST always be transactional stateful and they MUST NOT allow an application to change the transactional behaviour of `Proxy`. Therefore both `Proxy.setStateful()` and `Proxy.getStateful()` are deprecated in this specification.

For backward compatibility purposes:

- `Proxy.setStateful(boolean stateful)` now does nothing.
- `Proxy.getStateful` now always returns true.

Though stateless proxying by applications is deprecated in this release, a proxy implemented on containers compliant with this specification **MUST** forward messages statelessly under certain conditions as described in [RFC 3261, Section 16.7, point 10 and 16.10]

10.2 Operation

There are two ways in which an application can perform proxying operation using the `Proxy` object obtained from an incoming `SipServletRequest`:

1. By calling `proxyTo()` method on the `Proxy` object and passing the `URI(s)` to proxy the request to.
2. By creating `ProxyBranch` object(s) from the `Proxy` and then invoking `startProxy()` method on the `Proxy`.

When creating the `Proxy` object for an incoming INVITE request, the server also sends a 100 provisional response upstream. This is done to stop retransmissions from the client. The server **MAY** choose to send the 100 provisional response irrespective of the application type for the same reason.

Before proxying the request the application may modify the `Proxy` parameters as appropriate and may also modify the request object itself. This includes adding or removing headers but the proxy **MUST NOT** add to, modify or remove the message body as per [RFC 3261 16.6.1].

The `proxyTo()` method is overloaded to take either a single `URI` or a `List` of `URIs` specifying downstream destination(s). If the list of `URIs` passed to the `proxyTo()` method contains any duplicates (based on the definition of equality for the `URI` type), the proxy **MUST** ignore proxying to the duplicate `URIs` as per [RFC 3261 16.5]. Compliant servlet engines are required to be able to handle SIP and SIPS `URIs` and may know how to handle other `URI` schemes, for example, `Tel` `URIs`.

For each `URI` passed to it in one of the `proxyTo` methods, the container creates a new branch on which the request is proxied. The proxied request will have the request `URI` replaced with that of the specified destination `URI` and is routed either based on that modified request `URI`, or based on the top Route header field value, if one is specified.

Until a final response is forwarded upstream, the application may invoke the `proxyTo` method any number of times to add additional addresses to the set of destinations the container should proxy to.

Applications that wish to stay on the signaling path but which do not perform any routing decisions or otherwise influence the call setup may proxy with record-routing enabled without changing the request `URI`, specifically it will do the following:

```
public void doInvite(SipServletRequest req) {  
    ...  
    Proxy p = req.getProxy();  
    p.setRecordRoute(true);  
    p.proxyTo(req.getRequestURI());  
    ...  
}
```

The request will be proxied either to other applications according to the application composition mechanism or towards the external destination specified by the request `URI`.

Note that the SIP servlet programming model is asynchronous, therefore the application code may not process a request in the context of the service method upcall, but save the request instead and act on it asynchronously. For example, it is perfectly fine for an application to obtain the `Proxy` object from a request in the context of a service method, but defer the act of starting the `Proxy` until later, for instance trigger it by an application timer expiration. In other words, there is no expectation that the `Proxy` will be started by the application code in the context of the service method.

10.2.1 Proxy Branches

A `ProxyBranch` object represents a proxy branch. Explicitly creating proxy branches, modifying them individually and then starting the proxy process is another mechanism of proxying. The `ProxyBranch` object can be created from the `Proxy` object using the following method -

```
java.util.List<ProxyBranch> createProxyBranches(java.util.List<? extends  
URI> targets)
```

As an example -

```
..  
public void doInvite(SipServletRequest req) {  
    ...  
    Proxy p = req.getProxy();  
    p.setRecordRoute(true);  
    p.createProxyBranches(targetList);  
    p.startProxy();  
}
```

```
...
}
```

The effect of the above operation is same as `p.proxyTo(targetList)`; The proxy branch will not be used directly by the proxy application in most cases except when the application requires having some branch specific changes to the request.

Only the following operations are allowed on the `ProxyBranch` and the associated `SipServletRequest` object.

1. push different route headers to the associated `SipServletRequest` with the proxy branch.
2. cancel a certain proxy branch, by issuing a CANCEL request. The mechanism that the container uses to resolve the race condition of receiving a final response of the branch and mechanism of not sending CANCEL until the provisional response is received are same as defined for `Proxy.cancel()` method which cancels all the proxy branches and all branches created recursively.
3. set different record-route, path parameter, recurse flag or proxy timeout for the branch.
4. add/remove non system headers from the `SipServletRequest` before the request is proxied.
5. any other method invoked on `SipServletRequest` object (like `send()`, `createCancel()` etc), not relevant for this context MUST throw `IllegalStateException`.

Until a final response is forwarded upstream, the application may invoke the `createProxyBranches()` and `startProxy()` methods any number of times to add additional addresses to the set of destinations the container should proxy to.

The `createProxyBranches()` method explicitly and the `proxyTo()` method implicitly create the `ProxyBranch(es)`. These proxy branches can be retrieved from the `Proxy` object using the `ProxyBranch Proxy.getProxyBranch(URI)` on the `Proxy` interface. These are the top level branches created by user code providing a single URI or a list of URIs. Further `java.util.List<ProxyBranch> Proxy.getProxyBranches()` on `Proxy` interface returns all the top level branches associated with the `Proxy`.

Additionally any `ProxyBranch` might recurse if the recurse flag is true for the proxy/branch and the branch received the 3xx class response with alternate Contact headers. It is possible to retrieve the recursed branches by using the `ProxyBranch` method `java.util.List getRecursedProxyBranches()`.

10.2.2 Pushing Route headers

As mentioned in the SIP specification, a proxy may have a local policy that mandates that a request visit a specific set of proxies before being delivered to the destination [RFC 3261, section 16.6]. This is done by pushing a set of SIP/SIPS URIs on to the Route header field of the request.

This is the purpose of the `pushRoute` method of the `SipServletRequest` interface. The argument identifies a proxy that should be visited before the request reaches its final destination. If only one Route header field value is added, the container may choose to not actually push a Route header field value but rather to use the alternative of bypassing the usual forwarding logic, and instead just sending the request to the address, port, and transport for the proxy specified in the single `pushRoute` call. This is subject to the constraints specified in [RFC 3261].

10.2.3 Sending Responses

Applications may generate informational responses of their own before or during a proxying operation. This is done the same way it's done when acting as a UAS that is downstream to a proxy.

Once a request has been proxied, the application will not usually generate a final response of its own. However, there are cases where being able to do so is useful and so it's allowed but with some additional constraints. As long as no final response has been sent upstream, the application may create and send its own final response. The container behaves as if the final response was received from a (virtual) branch of the proxy transaction with the following qualifications:

- A virtual branch applies only in the context of a proxy i.e., `request.getProxy()` MUST be called before sending any response from the virtual branch upstream. Once the application calls `request.getProxy()`, the proxy could send the response upstream either before or after calling `proxyTo()`. If the response is sent upstream before calling `request.getProxy()`, the subsequent proxying operation with `Proxy.proxyTo()` MUST throw an `IllegalStateException`.
- If it's a 2xx response, it is sent upstream immediately without notifying the application of its own response.
- A non-2xx final response generated while the proxy transaction has outstanding branches contributes to the response context as any response received from a real branch. If it's eventually selected as the best response, the container will perform the usual best-response callback.

- If the best response received was a non-2xx and the application generated its own final response in the `doResponse` callback (be it a 2xx or non-2xx), then that response is sent immediately without invoking the application again for its own generated response.

This last item allows applications to create, for example, a different error response from the one chosen by the container as the best response. Note that if an application does generate its own final response when passed the best response received, it cannot also proxy to more destinations.

These rules are designed to guarantee that SIP servlet containers, when observed from the outside, do not violate the SIP specification.

From the container perspective when a response from the virtual branch is sent upstream, a derived `SipSession` is created for the virtual branch as per [6.2.3.2 Derived SipSessions](#). The original `SipSession` associated with the incoming request corresponds to the `Proxy` while the derived one represents the virtual UAS. Containers **MUST** restrict to creating only one virtual branch for each `Proxy`.

If the final response sent upstream was generated by the application itself, the container must update `SipSession` state as if it was a UAS (which in fact it is).

Alternative to sending a response itself, a proxy could use application composition to delegate the responsibility of sending the response to another application. The composition approach is recommended in cases where more than one virtual branch may be required. However, the virtual branch approach is simpler in that it makes sure that no other application is invoked between the proxy and the UAS.

10.2.4 Receiving Responses

The servlet container is responsible for automatically forwarding the following responses received for a proxying operation upstream:

- all informational responses other than 100
- the best response received when final responses have been received from all destinations.
- all 2xx responses for INVITE
- exactly one 2xx response for non-INVITE as per [RFC 3261 16.7, para - Check responses for forwarding]

Additionally, if the supervised flag is true, the servlet engine invokes the application for these responses before forwarding them upstream. The application may modify responses in the notification callback. In this case the modified responses are forwarded upstream. This is useful,

for example, for application level gateways that need to modify addresses in the body of responses.

As in the UAC case, applications are not invoked for incoming 100 responses.

When a 6xx response is received, the server CANCELS all outstanding branches and will not create new branches. Similarly, when a 2xx response is received, the server CANCELS all outstanding branches and will not create new branches unless the proxy is configured to not cancel branches with `setNoCancel(true)`.

When an application is invoked to handle the best final response received and it is not a 2xx or 6xx, the application may add addresses for further destinations through the `Proxy.proxyTo` method or by creating additional proxy branches and re-starting the proxy. The effect is that a new branch is created for each additional destination, as if the method had been invoked before a (tentative) best answer had been passed to the application. If, in the upcall informing a servlet of the best response received, the servlet proxies to one or more additional destinations, the container does not immediately forward the best response received so far as the new branch may result in a “better” response. The ability to call `proxyTo` in the callback for best response received is useful, for example, for writing call-forward-no-answer type services.

A consequence of the described behavior is that an application may be notified of more than one final response for a transaction. This can happen either because

1. the application proxied to more destinations in the notification for a final response
2. one or more 2xx retransmissions were received, see [10.2.4.1 Handling 2xx Responses to INVITE](#) below
3. multiple destinations returned 2xx responses.

In the first two cases, the application may end up being notified of the same final response more than once.

10.2.4.1 Handling 2xx Responses to INVITE

The 2xx responses to INVITEs are special in that they cause both client and server transactions to terminate immediately with retransmissions bypassing the transaction layer [RFC 3261, sections 13.3.1.4 and 17.1.1].

Ordinarily, SIP proxies handle 2xx retransmissions by forwarding them statelessly upstream based on the Via header field. In the case of an application server, we have the additional complication that one or more of the proxying applications (there may be more than one due to application composition) may wish to modify the 2xx response before it’s forwarded upstream.

In this case it's necessary to ensure that 2xx retransmissions are modified in the exact same way as the 2xx first received for the INVITE.

This can be achieved in several different ways and it is left to implementations to choose one. One solution is to not terminate INVITE client and server transactions when first seeing a 2xx for an INVITE. This enables the application server to handle 2xx retransmissions in a manner similar to the original 2xx, that is, invoke applications as necessary and then forward it upstream. For this reason, applications proxying INVITEs must be prepared to receive retransmissions of 2xx responses and must modify such retransmissions the same way they did the 2xx originally received.

10.2.4.2 Correlating responses to proxy branches

When notified of an incoming response, it is sometimes useful for applications to be able to determine which of several branches the response was received for when there was more than one branch outstanding. Applications identify branches by request URI, and so to test whether a response was for a particular branch it is sufficient to compare the request URI of that branch with the URI object previously passed to the `proxyTo` or `createProxyBranches` method by the application.

For incoming responses to proxied requests, `SipServletResponse.getRequest()` returns a `SipServletRequest` object representing the request that was sent on the branch on which the response was received. The request URI can then be obtained from the branch request object and be compared against destination URIs previously passed to the `proxyTo` or `createProxyBranches` method.

Containers are required to ensure that the request URI of the branch's request object is equal to the URI object that the application passed to the `proxyTo` method. (Recursion can, of course, cause branches to get created that the application didn't explicitly request and for which it has no URI.)

Servlets may be interested in being notified of intermediate final responses before the best final response is selected as per [RFC 3261, Section 16.6]. Version 1.0 of this specification did not offer any mechanism to peek at these intermediate responses.

To handle final responses received on each proxy branch, this specification introduces a new method on the `Servlet` class – `doBranchResponse(SipServletResponse resp)`. A servlet implementation that is interested in being notified of intermediate final responses received on individual branches must override this method. This method will be invoked by the container only if the supervised flag is true for the `Proxy`. A typical use-case for using the

`doBranchResponse` method would be to create and proxy to additional branches on the receipt of a non-2xx final response.

Any attempt to add additional branches in `doBranchResponse()` for a cancelled branch or a cancelled proxy will result in `IllegalStateException`. Only after the last branch has received a final response, the container must determine the best final response and pass it to the `doResponse()` method. Note that if the `doBranchResponse()` is not overridden then `doResponse()` method will be invoked only for the best final response as before.

10.2.4.3 Sequential Search Timeout

As of v1.1 of this specification the sequential search timeout may only be set using the general `proxyTimeout` parameter rather than the `sequentialSearchTimeout` parameter which has now been deprecated.

This section clarifies the case where a SIP servlet is acting as a sequential search proxy and the sequential search timeout has expired before a final response has been received for the current branch.

The sequential search timeout defined by the SIP servlet API is semantically similar to the Timer C with some differences explained below.

The following procedure describes the handling in case of sequential searches, sequential search timer is referenced as SST:

1. Forward request and start branch.
2. Start the SST on receipt of a provisional response. SST does not start before a provisional response is received.
3. If SST expired but final response not received then send a CANCEL to the INVITE transaction and move on to the next branch.
4. If SST value is greater than Timer-C then Timer-C SHOULD assume the value of SST.
5. If final response received within SST, then process response normally.

10.2.4.4 Handling 3xx responses

If the recurse flag of the proxy is set to false then it is up to the application developer to keep track of the contact addresses received in the redirect (3xx) responses. If the contact addresses received in multiple 3xx responses contain the same URI, the container SHOULD throw an `IllegalStateException` if the application attempts to proxy to the same URI again. This is to prevent creation of duplicate branches.

If the recurse flag of the proxy is set to true then the container SHOULD proxy only to unique contact address ignoring any duplicates. This may result in a case where a branch does not recurse on receiving a 3xx response with an alternate contact as a branch for that contact already exists. Calling `getRecurseProxyBranches()` on such a branch would result in an empty list. Every recursed contact results in a new branch and duplicate contacts are ignored as per [RFC 3261, Section 16.5]. The container MUST remove the redirect contacts from the 3xx response since they have produced new branches. A 3xx response without any contacts can never be a best response of a proxy according to [RFC 3261 16.7, point 4] and should never be propagated up to the application.

10.2.5 Sending CANCEL

An INVITE proxy operation can be aborted by invoking `Proxy.cancel()` or `ProxyBranch.cancel()`. An overloaded version of the cancel method is available so that application developers and containers MAY specify the reason for cancelling the `Proxy` or `ProxyBranch` in accordance with [RFC 3326]. Invoking cancel will cause the container to terminate all outstanding branches by sending CANCEL requests to the corresponding proxied INVITEs or cancel the specified branch respectively. The proxy operation then proceeds as in the normal case, as per the SIP specification. As far as response handling is concerned, the act of cancelling a branch can be thought of as a way of speeding up the generation of a final response. The application will still be invoked to handle responses (assuming the supervised flag is true). Calls to `Proxy.cancel()` have the effect of cancelling all branches currently in progress and clearing the proxy transaction's current target set. If the application subsequently calls `proxyTo()` with additional targets, the proxy transaction will create branches for those targets as usual. Alternatively applications can create new Proxy branches by invoking `Proxy.createProxyBranches()`.

The invocation of both variants of `cancel` method results in calling the cancel recursively on the `ProxyBranch` objects if the proxy recursed and there are recursed proxy branches.

As in the UAC case, applications are not invoked when CANCEL responses are received.

The comments made in [11.1.9 Sending CANCEL](#) regarding the container having to delay the sending of a CANCEL until a provisional response has been received apply to proxy transaction branches, also.

10.2.6 Receiving CANCEL

When receiving a CANCEL for a transaction for which a `Proxy` object exists the server responds to the CANCEL with a 200 and

- if the original request has not been proxied yet the container responds to it with a 487 final response
- otherwise, all branches are cancelled, and response processing continues as usual

In either case, the application is subsequently invoked with the CANCEL request. This is a notification only, as the server has already responded to the CANCEL and cancelled outstanding branches as appropriate. The race condition between the server sending the 487 response and the application proxying the request is handled as in the UAS case as discussed in [11.2.3 Receiving CANCEL](#).

10.2.7 Sending ACK

The SIP specification requires stateful proxies to generate ACKs for non-2xx final responses to INVITE requests. The purpose of these ACKs is to stop response retransmissions at the downstream server and have no semantic significance. Therefore, in the SIP Servlet API, the servlet container is responsible for generating ACKs for non-2xx final responses, and so SIP servlets should never generate ACKs for proxied requests.

10.2.8 Receiving ACK

ACKs for non-2xx final responses are just dropped. ACKs for 2xx's are treated as other subsequent requests, and are proxied by the container. Applications should not attempt to proxy ACKs explicitly.

If the application proxied the corresponding INVITE with record-routing enabled, the application is invoked before the ACK is proxied so as to give it an opportunity to modify the ACK. In this case, as when proxying the original request, the ACK is proxied downstream in its modified form.

10.2.9 Handling Subsequent Requests

As discussed in [15.6 Responses, Subsequent Requests and Application Path](#), a “subsequent” request is one that is dispatched to applications based on a previously established application path as opposed to initial requests that are dispatched to applications based on application selection process.

When proxying with the recordRoute flag being true, the server may receive subsequent requests in the same dialog. In these cases processing takes place as described above except that the server is responsible for proxying the request to the destination specified in the top Route header as specified in the SIP specification [RFC 3261, section 16.6].

The application is still passed the request object, though, and may modify it in the upcall before the server proxies it downstream. Applications should not attempt to explicitly proxy subsequent requests, any attempt to do so results in an `IllegalStateException`. The `isInitial` method on the `SipServletRequest` interface returns true if the request is initial as defined above, and can be used as an indication of whether the application should explicitly proxy an incoming request or not.

Applications are allowed to reject subsequent requests, but only when doing so in the upcall notifying the application of the subsequent request. After having performed the upcall the container will proxy the request unless the application generated a final response for it. The ability of proxy applications to respond to subsequent requests is needed, for example, for applications wishing to perform proxy authentication programmatically.

10.2.10 Max-Forwards Check

The Max-Forwards header serves to limit the number of elements a SIP request can traverse on the way to its destination. It consists of an integer that is decremented by one at each hop. RFC 3261 specifies that for a request to be proxied it must have a Max-Forwards value greater than 0 [RFC 3261, section 16.3].

Since a SIP Servlet container cannot know a priori whether an application will proxy a request or not, it cannot perform the Max-Forwards check before invoking the application. Instead, the container performs the check when the application invokes `getProxy()` on a `SipServletRequest`, and if Max-Forwards is 0 a `TooManyHopsException` is thrown. If unhandled by the application, this will be caught by the container which must then generate a 483 (Too many hops) error response. However, if the Max-Forwards is 0 for an OPTIONS request, the Proxy may not throw a `TooManyHopsException` but rather respond to the request.

Max-Forwards header is used within container for inter container loop detection also see [15.10 Loop Detection](#).

10.3 Proxying and Sessions

When an application proxies a request it may subsequently be invoked to handle requests and responses related to that transaction or to subsequent transactions, if the application record-routed.

Whenever an application is invoked because an event occurred on an existing transaction or for a subsequent request and the application has previously obtained a `SipApplicationSession`

and/or `SipSession`, the container must ensure that those same session objects are associated with subsequent requests and responses that the application is invoked to handle.

Also, as a result of forking proxies an initial request may result in multiple dialogs being established. [6.2.3 Creation of SipSessions](#) discusses under which circumstances a container must associate an incoming response or even a subsequent request with a `SipSession` cloned from that of an existing dialog.

10.4 Record-Route Parameters

When record-routing, it is often convenient for a SIP proxy to push state to the dialog endpoints, the user agents, and have it returned in subsequent requests of the dialog. This can be achieved by adding parameters to the URI of the Record-Route header field inserted by the proxy. The Record-Route header field value, including URI parameters, will contribute to the route set of the user agents and will be returned to the proxy in a Route header in subsequent requests. The loose routing mechanism of RFC 3261 ensures that a SIP proxy gets back the URI it record-routed with in a Route header field value of the subsequent request.

The `Proxy.getRecordRouteURI` method allows a record-routing proxy application (be it dialog stateful or stateless) to set parameters on the Record-Route header that will be inserted by the container into a proxied request:

```
SipURI getRecordRouteURI();
```

The URI returned by the `getRecordRouteURI` method should be used only for pushing application state in the form of URI parameters to the user agents. Applications must not set SIP URI parameters defined in RFC 3261. This includes `transport`, `user`, `method`, `ttl`, `maddr`, and `lr`. Other components of the URI, e.g. `host`, `port`, and URI scheme must also not be modified by the application. These Record-Route URI components will be populated by the container and may or may not have valid values at the time an application proxies a request.

Parameters added to Record-Route header field values in this manner can be retrieved from subsequent requests using the `getParameter` method of `SipServletRequest` or through access to the popped Route header, discussed in [5.6.1 Parameters](#).

Note that this mechanism guarantees that when a proxy application adds parameters to a URI obtained through the `getRecordRouteURI` method, the value of those parameters can be retrieved by the same application using `SipServletRequest.getParameter` on subsequent requests in that dialog (assuming the user agents involved are well-behaved). However, there is no guarantee that the parameters go unchanged into an actual Record-Route header. For example,

due to application composition, the parameters may in some cases go into a Contact header, and also, the container may choose to insert only one Record-Route header when there are multiple record-routing proxies on the application path. In this case the container would have to encode parameters so as to avoid name clashes between applications. Implementations may even choose to store the parameter set locally, possibly along with other dialog related data, and retrieve it based on a dialog ID computed for subsequent incoming requests. There is no hard bound defined for the size of data that can be pushed to endpoints in Record-Route header fields but it is recommended that application writers keep in mind that some implementations may not function correctly with large size data.

This version of the SIP Servlet API does not allow proxy applications to push different state to the two endpoints of a dialog. The UAS will copy the Record-Route header field of requests unchanged into 2xx responses and the proxy application is not given the opportunity to modify its own previously inserted Record-Route parameters when processing 2xx responses. This would effectively rewrite the Record-Route header field value and cause state received in subsequent requests from the caller to differ from that received in subsequent requests from the callee. Such a feature would potentially be useful but may have an adverse effect on performance and security as the Record-Route header in responses cannot be protected end-to-end if rewritten by proxies.

10.5 Path Header and Path Parameters

The Path introduced in [RFC 3327] header provides the semantics of a Path created by a REGISTER request, which shall be followed by subsequently created dialogs. The concept is similar to Record-Route but defines such a path outside of dialogs. The Path is added by intermediate proxies and maintained by the Registrar.

The Home Proxy shall then look up the Path alongside the AOR binding and add pre-loaded routes to the request destined for the registered UA. The requirements on SIP Servlet API can be summarized as -

1. Proxy application may want to add the `Proxy` (itself) to the Path in REGISTER
2. A proxy application besides adding its own Path may also push arbitrary Path [RFC 3327 Section 5.2]

Similar to `getRecordRouteURI()` (see [10.4 Record-Route Parameters](#)) the `getPathURI()` API is additionally defined on `Proxy` interface, this gives applications a chance to set some parameters on the Path URI. These parameters can subsequently be retrieved as the Request URI parameters or from the popped route mechanism as defined in [5.6.3 Popped Route Header](#).

The SIP option tag "path" should be present in immutable `List` of extensions supported as described in section [3.2 Extensions Supported](#) in addition to the `100rel` option tag.

Path is no different from Record-Route in semantics, except that it creates a path for subsequent dialogs, therefore Path header is a "System Header" (see [5.4.2 System Headers](#)), managed by SIP Servlet Container.

A `pushPath()` method is defined on `SipServletRequest` similar to `pushRoute()` method. However, `pushPath()` shall throw an `IllegalStateException` on non-REGISTER Requests.

11 Acting as a User Agent

This chapter describes how SIP servlets perform various UA operations such as initiating and responding to requests. The SIP specification lists client and server transaction state machines that define at which points in processing various actions are allowed. As mentioned in [5.2 Implicit Transaction State](#), when an application attempts to perform an action that would violate the SIP state machine, an `IllegalStateException` is thrown by the container.

The notions of SIP client and server are meaningful only in relation to a particular transaction. An application acting as a UA will often have to be able to act as both a client and a server in a dialog.

11.1 Client Functions

11.1.1 Creating Initial Requests

When initiating a request a UAC first creates a `SipServletRequest` object and then sends it, possibly after having modified it. Creation of the request object is done by invoking the overloaded `createRequest` method on a `SipFactory` if the request is an initial request, that is, if it doesn't belong to an existing SIP dialog.

More precisely, a SIP servlet carries out the following steps in order to send a new request which does not belong to an existing dialog:

- look up a `SipFactory` as an attribute on the `ServletContext`, see [3.1 The SipFactory](#)
- invoke one of the overloaded `SipFactory.createRequest` methods

- modify the resulting `SipServletRequest` as appropriate, for example, set the content
- invoke `send` on the request object

The `SipFactory` interface defines the following methods for creating new requests:

```
SipServletRequest createRequest(  
    SipApplicationSession appSession,  
    String method,  
    Address from,  
    Address to);  
  
SipServletRequest createRequest(  
    SipApplicationSession appSession,  
    String method,  
    URI from,  
    URI to);  
  
SipServletRequest createRequest(  
    SipApplicationSession appSession,  
    String method,  
    String from,  
    String to) throws ServletException;
```

The returned `SipServletRequest` exists in a new `SipSession` which belongs to the specified `SipApplicationSession`. The handler for the newly created `SipSession` is the application's default servlet for v1.0 applications, see [6.2.6 The SipSession Handler](#) and the main servlet for v1.1 applications, see [16.2 Servlet Selection](#). This can be changed by the application by invoking `SipSession.setHandler`. The container is required to assign the returned request a fresh, globally unique Call-ID as defined by the SIP specification. The From and To headers are as specified by the from and to parameters, respectively. The container is responsible for adding a CSeq header. The request method is given by the method parameter. The default value of the request URI is the URI of the To header, with the additional requirement that for REGISTER requests the user part of a SIP request URI is empty. The application can change this by invoking `setRequestURI` on the newly created `SipServletRequest`.

Note: ACK and CANCEL requests have special status in SIP. [11.1.7 Sending ACK](#) and [11.1.9 Sending CANCEL](#) discuss how and when applications generate those requests. The `SipFactory` methods listed above MUST throw an `IllegalArgumentException` when invoked to create ACK or CANCEL requests.

Once constructed, the application may modify the request object, for example, by setting the request URI, adding headers, or setting the content. The request is then sent by invoking the

`send()` method on `SipServletRequest`. The request is routed by the container based on the request URI and any Route headers present in the request, as per the SIP specification. A fourth version of `SipFactory.createRequest` is intended specifically for writing back-to-back user agents, and is discussed in [12.1 B2BUA Helper Class](#).

11.1.1.1 Copying From and To Addresses

Ordinarily when `Address` objects are set as header field values of a message, they are not copied. However, the From and To header fields are special in that they are container managed and have certain requirements regarding tag parameters imposed from the SIP specification. For these reasons, the `SipFactory.createRequest` method makes a deep copy of the *from* and *to* arguments before making them the values of the From and To header fields of the newly created request. If the copied *to* `Address` has a *tag* parameter it's removed as the To header field of outgoing initial requests must be tag-free. The copied *from* `Address` is given a fresh *tag* parameter according to the SIP specification. Any component of the *from* and *to* URIs not allowed in the context of SIP From and To headers are removed from the copies. This includes, headers and various parameters as referenced in [RFC 3261, Section 20].

The copied *from* and *to* `Address` objects are associated with the new `SipSession`. If a dialog is established, the container must update the *to* *tag* to the value chosen by the peer SIP element. Subsequent requests belonging to the same `SipSession` will have the same From and To headers. Applications can retrieve the copied *from* and *to* `Address` objects (reflecting new tags) from either the `SipSession` or the newly created request, but are not allowed to modify them.

11.1.2 Creating Subsequent Requests

For subsequent requests, that is, requests made in an already established dialog, applications use the following method on the `SipSession` representing the dialog:

```
SipServletRequest createRequest(String method);
```

This method constructs a request with container provided values for request URI and Call-ID, From, To, CSeq, and Route headers. Again, the application may modify the request before invoking it by invoking the `send` method on the request object.

11.1.3 Pushing Route Header Field Values

A UAC may push Route header field values, basically SIP or SIPS URIs, onto the Route header field of the request. The effect is that the request will visit the set of proxies identified in those

Route values before being delivered to the destination. This is what is known as a pre-loaded Route in the SIP specification.

The `SipServletRequest` interface defines the `pushRoute` method for adding entries to the Route header field:

```
void pushRoute(SipURI uri);
```

The `pushRoute` method is applicable to be used by UAs only till the time when the dialog has not yet been established. This means that `pushRoute` can only be done on the initial requests. Subsequent requests within a dialog follow the route set. Any attempt to do a `pushRoute` on a subsequent request in a dialog **MUST** throw an `IllegalStateException`.

The Route header behaves like a stack – `pushRoute` adds items to the head of the list and the SIP routing algorithm pops items from the top. Therefore, if `pushRoute` is called more than once, the request will visit the most recently added proxy first. The same mechanism is available for proxying applications and is discussed in [10.2.2 Pushing Route headers](#).

11.1.4 Sending a Request as a UAC

Once created, a request is sent by invoking `send()` on the `SipServletRequest` object:

```
void send() throws IOException;
```

If the `send` method call throws an exception it means the request was not successfully sent and the application will not receive any callbacks pertaining to that transaction, that is, the application will not receive any responses for this request.

If the `send` method does not throw an exception, the container is obliged to present the application with a final response except for an ACK request. Containers may send the request asynchronously in which case sending may fail after the `send` method has returned successfully. In this type of situation, the container will generate its own final response except for an ACK request. In this particular case, a 503 (Service Unavailable) response would be appropriate as per [RFC 3261 8.1.3.1 and 16.9].

11.1.5 Receiving Responses

The UAC application is invoked for all incoming responses except 100 responses (and retransmissions). The servlet invoked is the current handler for the `SipSession` to which the request belongs (see [6.2.6 The SipSession Handler](#)). The container invokes `Servlet.service`

with `SipServletResponse` and a null value for the request argument. If the servlet extends the `SipServlet` abstract class, the response will be further dispatched based on the value of the status code.

11.1.5.1 Handling Multiple Dialogs

Due to forking at downstream proxies it is possible that multiple 2xx responses will be received for a single INVITE request. In this (and similar) cases the container clones the original `SipSession` for the second and subsequent dialogs, as detailed in [6.2.3.2 Derived SipSessions](#). The cloned `SipSession` object will belong to the same `SipApplicationSession` as the original `SipSession` and contain the same application data but invoking its `createRequest` method will create requests belonging to the second or subsequent dialog, that is, with a `To` tag specific to that dialog.

11.1.6 Transaction Timeout

If timeout occurs in the transaction layer as specified by RFC3261 (i.e. In INVITE client transaction, if Timer B fires; and in non-INVITE client transaction, Timer F fires) , the container generates its own 408 Request Timeout final response and passes it to the application.

11.1.7 Sending ACK

In SIP, final responses to INVITE requests trigger sending of an ACK for that response from the UAC to the UAS. ACKs to non-2xx final responses are needed for reliability purposes only and are sent hop-by-hop, that is, they are generated by proxies as well as the UAC. ACKs to 2xx responses, on the other hand, signal completed session setup and may carry semantically useful information in the body, for example IP addresses and codecs for the media streams of the session. These ACKs are sent end-to-end from the UAC all the way to the UAS. Generally speaking, only ACKs for 2xx responses are of interest to services, and so SIP servlet containers are responsible for generating ACKs for non-2xx responses, while applications are responsible for generating ACKs for 2xx responses. A request object representing the ACK is created by calling the `SipServletResponse` method:

```
SipServletRequest createAck()
```

The application may modify the returned ACK object before invoking `send` on it. It is the containers responsibility to retransmit application generated ACKs for 2xx's when a 2xx retransmission is received and the container must not deliver the 2xx retransmission to the UAC

application. It is recommended that containers generate ACKs for non-2xx final responses prior to invoking the application, so as to stop response retransmission as soon as possible.

11.1.8 Sending PRACK

RFC 3262 introduced Reliable provisional responses and the new PRACK method. This RFC also added two new system headers RSeq and RAck. In order to allow for generation of PRACKs the method

```
SipServletRequest createPrack()
```

is to be used on the `SipServletResponse`. The RAck header will be populated by the container in such PRACKs according to procedures specified in RFC 3262.

11.1.9 Sending CANCEL

A UAC can cancel an INVITE request in progress by sending a CANCEL request for the INVITE. A SIP servlet acting as a UAC can invoke the following `SipServletRequest` method on the original INVITE request object:

```
SipServletRequest createCancel()
```

The application sends the returned CANCEL request by invoking `send` on it. Note that responses to CANCEL requests are not passed to the application. UACs and proxies are not allowed to cancel an INVITE request before a 1xx response has been received [RFC 3261, section 9.1]. SIP Servlet applications may do so, though. It is the containers responsibility to delay sending of the CANCEL until a provisional response has been received.

11.2 Server Functions

By definition, a SIP servlet that responds to an incoming request with a final response becomes a UAS for the corresponding transaction.

11.2.1 Sending Responses

A servlet may generate a number of provisional responses as well as a single final response for an incoming request. It does so by invoking `createResponse` on the request object. The resulting `SipServletResponse` is then subsequently sent, possibly after having been modified, by invocation of the `send` method. 2xx responses to INVITEs are special in that they are

retransmitted end-to-end. The SIP specification is defined in terms of a layered software model in which the transaction user (a UAS in this case) accepting an INVITE is responsible for periodically retransmitting the 2xx [RFC 3261, section 13.3.1.4]. In the SIP Servlet API this task must be handled by the container. This is fully within the scope of the logical model used for purposes of presentation in RFC 3261.

11.2.2 Receiving ACK

Applications are notified of incoming ACKs for 2xx responses to INVITEs which the application sent upstream. Applications are not notified of incoming ACKs for non-2xx final responses to INVITE. These ACKs are needed for reliability of final responses but are not usually of interest to applications.

It is possible that a UAC fails to send an ACK for an accepted INVITE request before the transaction times out. This is communicated to the application through the `SipErrorListener` mechanism discussed in [8.1 SIP Servlet Event Types and Listener Interfaces](#).

11.2.3 Receiving CANCEL

When a CANCEL is received for a request which has been passed to an application, and the application has not responded yet or proxied the original request, the container responds to the original request with a 487 (Request Terminated) and to the CANCEL with a 200 OK final response, and it notifies the application by passing it a `SipServletRequest` object representing the CANCEL request. The application should not attempt to respond to a request after receiving a CANCEL for it. Neither should it respond to the CANCEL notification.

Clearly, there is a race condition between the container generating the 487 response and the SIP servlet generating its own response. This should be handled using standard Java mechanisms for resolving race conditions. If the application wins, it will not be notified that a CANCEL request was received. If the container wins and the servlet tries to send a response before (or for that matter after) being notified of the CANCEL, the container throws an `IllegalStateException`.

11.2.4 Handling Error Conditions

11.2.4.1 Receiving unimplemented subsequent requests

When a UAS application does not handle subsequent requests (for example, `doMessage()` is not over-ridden to handle an in-dialog MESSAGE request) the container SHOULD send back a 501 Not Implemented response. However, this does not apply to ACKs for 2xx responses for INVITE.

11.2.4.2 Handling Pending Invites

As per [RFC 3261 14.2], if the container receives a second INVITE before it sends the final response to a first INVITE with a lower CSeq sequence number on the same dialog, it **MUST** return a 500 (Server Internal Error) response to the second INVITE and **MUST** include a Retry-After header field with a randomly chosen value of between 0 and 10 seconds.

If the container receives an INVITE on a dialog while an INVITE it had sent on that dialog is still in progress, it **MUST** send a 491 (Request Pending) response to the received INVITE.

12 Back To Back User Agents

A back-to-back user agent (B2BUA) is a SIP element which acts as an endpoint for two or more dialogs and forwards requests and responses between those two dialogs in some fashion. Experience has shown that they are one of the most frequently used application types.

B2BUA's are sometimes considered undesirable because of their potential to break services. This potential stems from the fact that they sit between two endpoints and in some way mediate the signaling between the endpoints. If the B2BUA doesn't know about an "end-to-end" service being used between those two endpoints it may inadvertently break it.

B2BUA's are, however, an important tool for SIP application developers and as such are supported by the SIP Servlet API. This specification adds a new helper class and some new methods to the SIP Servlet API making the B2BUA pattern very easy to implement.

12.1 B2BUA Helper Class

A B2BUA helper class contains all the useful methods for a B2BUA operation. The B2BUA helper class instance can be retrieved from the `SipServletRequest` by invoking `getB2buaHelper()` on it.

```
B2buaHelper getB2buaHelper() throws IllegalStateException;
```

This also indicates to the container that this application wishes to be a B2BUA.

Any UA operation would be permitted by the application but the application cannot act as Proxy after that, so any invocation to `getProxy()` must then throw `IllegalStateException`.

Similarly the `getB2buaHelper()` method would throw an `IllegalStateException` if the application has already retrieved a proxy handle by an earlier invocation of `getProxy()`.

12.2 Creating new B2BUA Request

The behavior specified here aims at minimizing the risk of breaking end-to-end services by suggesting that all unknown headers be copied from the incoming request to the outgoing request.

When an application receives an initial request for which it wishes to act as a b2bua, it may invoke the `createRequest` method available on `B2buaHelper`:

```
SipServletRequest B2buaHelper.createRequest(  
    SipServletRequest origRequest,  
    boolean linked,  
    java.util.Map<java.lang.String, java.util.Set> headerMap)  
    throws IllegalArgumentException
```

This method creates a request identical to the one provided as the first argument according to the following rules:

- All unknown headers and Route, From and To headers are copied from original request to the new one. A new From tag is assigned by the container to this newly created request.
- Record-Route and Via header fields are not copied. As usual, the container will add its own Via header field to the request when it's actually sent outside the application server.
- The headers in the new request can be taken for the optional `headerMap` which is a `Map` of headers that will be used in place of the ones from `origRequest`. eg.

```
    {"From" => {sip:myname@myhost.com},  
     "To"   => {sip:yourname@yourhost.com} }
```

where "From" and "To" are keys in the map. The only headers that can be set using this `headerMap` are non-system headers and From, To and Route headers. For Contact header if present only the user part and some parameters are to be used as defined in [4.1.3 The Contact Header Field](#). The values in the map is a `java.util.Set` to account for multi-valued headers.

The values in `headerMap` MUST override the values in the request derived from the `origRequest`. Specifically they do not append header values. Attempt to set any other system header results in an `IllegalArgumentException`.

- The "linked" boolean flag indicates whether the ensuing `SipSession` and `SipServletRequest` are to be linked to the original ones. The concept of linking is discussed in [12.3 Linked SipSessions And Linked Requests](#).
- For non-REGISTER requests, the Contact header field is not copied but is populated by the container as usual.

These methods are included for convenience and performance. Like other `createRequest` methods, the returned request belongs to a new `SipSession`.

Note: The `SipFactory.createRequest(SipServletRequest origRequest, boolean sameCallId)` method has been deprecated in this specification as the usage of this method with `sameCallId` flag as "true" actually breaks the provisions of [RFC 3261] where the Call-ID value is to be unique accross dialogs. Instead, the use of `B2buaHelper.createRequest(SipServletRequest origRequest)` is recommended.

12.3 Linked SipSessions And Linked Requests

12.3.1 Explicit Session Linkage

In a B2BUA there usually are two `SipSessions` (though there can be more than two). The most common function of a B2BUA is to forward requests and responses from one `SipSession` to the other, after performing some transformation [application of business logic]. The B2BUA helper class simplifies the usage of this pattern. It optionally links the two `SipSessions` and `SipServletRequest`s when you use the `B2buaHelper` to create the new request.

```
SipServletRequest B2buaHelper.createRequest (
    SipServletRequest origRequest,
    boolean linked,
    java.util.Map<java.lang.String, java.util.Set> headerMap)
```

The effect of this method when the `linked` parameter is true, is to create a new `SipServletRequest` using the original request, such that the two `SipSessions` and the two `SipServletRequest`s are linked together. When the two `SipSessions` and requests are linked then you may be able to navigate from one to the other.

The sessions linked can then later be accessed as -

```
doSuccessResponse(SipServletResponse response) {
    ....
    otherSession = B2buaHelper.getLinkedSession(response.getSession());
```

```
// do something on otherSession
....
}
```

The `getLinkedSession` is a method defined on the `B2buaHelper` class. The helper works like a `Visitor` to the `SipSession` (and other classes) and encapsulates the functionality useful to a B2BUA implementation.

Note that 1-1 linking is a convenience function and in no way mandatory for B2BUA functionality, in other words, the old style B2BUAs using just v1.0 API, where the peer session/request was stored as attributes, are also supported without any change.

Similar to `SipSessions`, the linked `SipServletRequest` can be obtained from the method: `B2buaHelper.getLinkedSipServletRequest(SipServletRequest)`.

Besides the `B2buaHelper.createRequest()` method, the linking can also be explicitly achieved by calling:

```
B2buaHelper.linkSipSessions(session1, session2) throws
IllegalArgumentException;
```

The `IllegalArgumentException` is thrown when the sessions cannot be linked together, such as one or both sessions have terminated or belong to different `SipApplicationSessions` or one or both have been linked to some different `SipSessions`.

Also for unlinking -

```
B2buaHelper.unlinkSipSessions(session) throws IllegalArgumentException;
```

unlinks other session linked with this session.

One `SipSession` at any given time can be linked to only one other `SipSession` belonging to the same `SipApplicationSession`.

The linkage at the `SipServletRequest` level is implicit whenever a new request is created based on the original with link argument as true. There is no explicit linking/unlinking of `SipServletRequests`.

12.3.2 Implicit Session Linkage

Another useful method on `B2buaHelper` for subsequent requests is:

```
B2buaHelper.createRequest(SipSession session, SipServletRequest
origRequest, java.util.Map<java.lang.String, java.util.Set>
headerMap) throws IllegalArgumentException
```

Here the session is the `SipSession` on which this subsequent request is to be sent, the `origRequest` is the request received on another `SipSession` using which this request is to be created and the `headerMap` can contain any non system header which needs to be overridden in the resulting request, any attempt to set any system header results in an `IllegalArgumentException`. The semantics of this method is similar to the `SipSession.createRequest`. This also results in automatically linking the two `SipSessions` (if they are not already linked) and the two `SipServletRequest`s

12.4 Access to Un-Committed Messages

A method on `SipServletMessage` defines the committed semantics for a message -

```
public boolean isCommitted();
```

The conditions under which a message is considered committed is detailed under [5.2 Implicit Transaction State](#).

A method defined on the new B2BUA Helper class would give the application a list of un-committed messages in the order of increasing CSeq based on the UA mode. This is because there may be more than one request/response uncommitted on a `SipSession`.

```
List<SipServletMessage> B2buaHelper.getPendingMessages(SipSession,
UAMode);
```

The `UAMode` is an enum with values of `UAC` or `UAS`. The same session can act as `UAC` or `UAS`, the `UAMode` indicates messages pertaining to which mode is queried.

For example, consider a B2BUA involved in a typical INV-200-ACK scenario that receives an ACK on one leg and wishes to forward it to the other. The B2BUA could call

```
B2buaHelper.getPendingMessages(leg2Session, UAMode.UAC)
```

to retrieve the pending messages which (as per [5.2 Implicit Transaction State](#), point 7) would contain the original 200 response received on the second leg. The B2BUA could then create the ACK using the

`SipServletResponse.createAck()` API. A PRACK request could be created in a similar way from a reliable 1xx response.

12.5 Original Request and Session Cloning

The incoming request that results in creation of a `SipSession` is termed as the original request, a response to this original request can be created by the application even if the request was committed and application does not have a reference to this request. This is necessary because the B2BUA application may require to send more than one successful response to a request. For example, when a downstream proxy forked and more than one success responses are to be forwarded upstream. This can only be required on initial requests, as only original requests shall need such multiple responses.

```
SipServletResponse
    B2buaHelper.createResponseToOriginalRequest(SipSession session, int
        status, String reasonPhrase) throws IllegalStateException;
```

The response thus generated MUST have a different "To" tag from the other responses generated to the request and must result in a different `SipSession`. In this (and similar) cases the container clones the original `SipSession` for the second and subsequent dialogs, as detailed in [6.2.3.2 Derived SipSessions](#). The cloned session object will contain the same application data but its `createRequest` method will create requests belonging to that second or subsequent dialog, that is, with a "To" tag specific to that dialog.

12.5.1 Cloning and Linking

In the case above when more than one response is received on the UAC side, it results in the cloned UAC `SipSession`, when the response is sent on the UAS side using the original request it is in context of the cloned UAS `SipSession`. These `SipSessions` are thus pair-wise linked for easy navigation.

If UAS-1 is the `SipSession` on which the incoming request was received and UAC-1 is the `SipSession` on which the outgoing request for relayed, then in case of multiple 2xx responses one response will be processed by the UAC-1 `SipSession` but when another 2xx response is received the container MUST clone the UAC-1 `SipSession` to create UAC-2 `SipSession` to process this new response. The `B2buaHelper` has the convenient `getLinkedSession` method to navigate from UAS-1 to UAC-1 and vice-versa. As for the cloned `SipSession` UAC-2 the `B2buaHelper` MUST be able to furnish the UAS-2 (a clone of UAS-1) when the `getLinkedSession` is invoked with UAC-2. Containers may chose to clone the UAS-2

`SipSession` lazily. The applications then can create the response to the original request but now in the context of UAS-2 by invoking the method -

```
SipServletResponse  
    B2buaHelper.createResponseToOriginalRequest(SipSession session-uas-2,  
        int status, String reasonPhrase) throws IllegalStateException;
```

Back To Back User Agents

13 Converged Container and Applications

SIP Servlet API is designed to facilitate the creation of SIP based applications in an easy servlet programming model. The SIP Servlet API was based on the Servlet API and thus was a peer to the HTTP Servlet API. Many of the interesting applications require the use of both SIP and HTTP protocols in a single application use-case, as an example a conference call application providing a web portal for management and monitoring, a classic click-to-call application, a CBSNA application etc. Therefore the usage of SIP Servlet API and HTTP Servlet API together in a converged application is a natural fit. Based on the fact that HTTP Servlet API is closely associated with the Java EE standards and also based on the real world requirements to have a comprehensive convergence model, a number of features are required to be supported by a converged container to enable not only SIP and HTTP convergence but also SIP Servlets and rest of Java EE. This section details the converged container features in detail.

13.1 Converged Application

A converged application is an application that not only has SIP Servlet component but also has at least one of HTTP Servlet or Java EE component like EJBs.

A converged application having SIP and other non HTTP Java EE components **MUST** be packaged as a single ear (Enterprise Archive) file. SIP and Web components of a converged application **SHOULD** be packaged as a single sar or war archive, either standalone or part of an ear archive. The SIP and Web applications **MAY** however be packaged into a single ear archive file with different individual sar/war files.

13.1.1 SIP and Java EE Converged Application

An application ear file with its optional application.xml deployment descriptor is the precinct of a converged application. A converged application having a SIP and a Java EE component other than Web MUST be packaged and deployed as a single ear archive file. EAR or Enterprise archive file is defined as application delivery format for Java EE applications and is specified in Java EE (Java EE 5) specifications.

If the application.xml is used for specifying the modules included in the application archive, the deployment descriptor's <web> element MUST be used for SIP Servlet applications. e.g.

```
<module>
  <web>
    <web-uri>mysipapp.war</web-uri>
  </web>
</module>
```

If the application.xml deployment descriptor is not present, then all files in the application package with a filename extension of .war or .sar and which contain the sip.xml deployment descriptor are considered converged SIP components.

This application scope is used to determine the extent of availability of convergence features in Java classes as we will see in the following sections on `SipFactory` and `SipSessionsUtil` injections.

13.1.2 SIP and HTTP Converged Application

The reason why the SIP and HTTP convergence is special cased is because SIP Servlet and HTTP Servlet APIs have a shared base API and by virtue of that some Servlet API specific constructs are additionally available to them for better convergence features. The converged SIP and Web application may be packaged as a single standalone sar or war archive as described in [7 SIP Servlet Applications](#).

For converged SIP and Web application packaged like this -

1. The archive MUST have the packaging directory structure as in [7.4 Deployment Hierarchies](#)
2. SIP and HTTP Servlets MUST have the same view of the application, which includes:
 - Context parameters
 - Context attributes
 - Context-listeners and context-attribute listeners

- Application class-loader
- Application scoped JNDI, which can be looked up against "java:comp/env" on the `InitialContext`.

A converged SIP and Web applications MAY also be packaged as independent sar/war and war files respectively, within a single application ear file. If such a packaging structure is used then the convergence features are similar to SIP and Java EE components, specifically the benefits of a single sar/war file of shared context are not available.

13.2 Accessing SIP Factory

SIP servlet containers MUST make a `SipFactory` instance available to applications through a `ServletContext` attribute with name "javax.servlet.sip.SipFactory" (See [3.1 The SipFactory](#)). Access to `SipFactory` instance can be made by the servlet applications, in order to make `SipFactory` accessible to other Java EE components in the application, the container MUST make `SipFactory` instance available by means of an Annotations based Dependency Injection.

The `@Resource` annotation defined in the Common Annotations for Java Platform (JSR250) is to be used for `SipFactory` injection.

When this annotation is applied to a field of type `SipFactory` in a converged application on a converged container, the container MUST inject an instance of `SipFactory` into the field when the application is initialized. This annotation can be used in place of the existing `ServletContext` lookup for the `SipFactory` from within a `Servlet`.

```
SipFactory sf =
    (SipFactory) getServletContext().getAttribute(SIP_FACTORY);
```

Usage above and

```
@Resource
private SipFactory sf;
```

are equivalent. For converged containers, the injected `SipFactory` would appear as `sip/<appname>/SipFactory` in the application scoped JNDI tree, where the appname is the name of the application as identified by the container (see [7.6 Application Names](#)).

In this example, this annotation can be used in any SIP Servlet or a Java EE application component deployed on the converged container in the same EAR archive file. The container must inject an instance of `SipFactory` into the field ("sf" in this example) at the time of

application initialization. For the purposes of associating a `Servlet`, with, say the responses received to the request created using the `SipFactory`, the container must use the application's main servlet as defined in [16.2 Servlet Selection](#). This can be changed to another servlet through the `setHandler` method of the `SipSession`.

Note that containers MAY allow the `@Resource` annotation to be present outside of SIP applications. In such cases, the name element of the `@Resource` annotation MUST identify the JNDI name of the desired factory to inject. Otherwise, an invocation of `SipFactory.createApplicationSession()` would be unable to determine the intended SIP application.

13.3 Accessing SipApplicationSession By ID

`SipApplicationSession` instance is identified by a certain container specific ID, it is frequently required to access the `SipApplicationSession` given its ID. Converged containers MUST provide a `SipSessionsUtil` lookup object using which the converged applications can access the `SipApplicationSession` instance by its ID. The reference to this lookup utility can be obtained from the `ServletContext` attribute `javax.servlet.sip.SipSessionsUtil` or by use of the `@Resource` annotation as described in [18.2.7 Annotation for SipSessionsUtil Injection](#).

Similar to the constraints on injection on `SipFactory`, the `SipSessionsUtil` lookup MUST NOT return the `SipApplicationSession` instance across application boundaries. Specifically if

```
getApplicationSessionById(java.lang.String applicationSessionId)
```

is invoked for an `applicationSessionId` not belonging to a certain application, the invocation MUST return null. The application boundary or scope as used in this application is the packaging structure of an ear archive file.

Similarly, a `SipApplicationSession` instance created with a specific session key (using `SipFactory.createApplicationSessionByKey()`) can be looked up by :

```
getApplicationSessionByKey(java.lang.String applicationSessionKey, boolean create)
```

13.4 Encoding URLs

The `SipApplicationSession.encodeURL()` encodes a URL (for example, HTTP URL) with the application session ID using the `sipappsessionid` parameter. The encoded URL SHOULD have the application session ID encoded in a way such that the parameter value which encodes the application session ID is unique across implementations. The recommended way is to use the java package name of the implementation, like `"com.acme.appsession"`.

This mechanism can be used by applications to encode the HTTP URL with the `SipApplicationSession` ID. The encoded URL can then be sent out through some out of band mechanism to an external UA. When the HTTP Request comes back to the converged container with this URL, the container MUST associate the new `HttpSession` with the encoded `SipApplicationSession`.

In case the HTTP request is not a new request but a follow-on request already associated with an existing HTTP session then converged containers MUST use the HTTP session association mechanism described in the HTTP Servlet specification, to route the request to the right HTTP session. If that HTTP session was not associated with the encoded `SipApplicationSession` in the request then that association MUST occur. This mechanism is similar to how the (deprecated) `encodeURI()` operates for SIP.

13.5 Association of HTTP Session With SipApplicationSession

[6.4.1 Message Context](#) defines the concept of a message context and how the `SipSessions` are associated with the `SipApplicationSessions`. In a similar way the `HttpSessions` MUST be associated with `SipApplicationSessions` as described below.

- Even though the `SipFactory` instance is available in the HTTP scope through the context attribute or the annotation and using the `SipFactory` a new `SipApplicationSession` can be created, the recommended way of creating a `SipApplicationSession` is to use the

```
ConvergedHttpSession.getApplicationSession()
```

The converged containers MUST make available the `javax.servlet.sip.ConvergedHttpSession`, an extension of `javax.servlet.http.HttpSession` to the applications. The effect of invoking this method is that the application session if currently associated is returned, however if no application session is associated it is created, associated with the `HttpSession` and returned.

- Another way a new `ConvergedHttpSession` gets associated with a `SipApplicationSession` is when the HTTP request has an URL encoded with an existing `SipApplicationSession` as described in [13.4 Encoding URLs](#).

13.6 Finding Parent `SipApplicationsSession` From A Protocol Session

It is often required to access the parent `SipApplicationSession` from the protocol session (SIP or HTTP). `SipSession` has a method defined `getApplicationSession()` to access the parent `SipApplicationSession`. For access to parent `SipApplicationSession` from `HttpSession` the applications can use `javax.servlet.sip.ConvergedHttpSession` interface which has a convenient method to access the parent `SipApplicationSession`.

13.7 Encoding HTTP URLs With HTTP Sessions In Non HTTP Scope

Given an application with multiple sessions (SIP and HTTP) it is often required to create an encoded HTTP URL such that the subsequent HTTP requests are routed to the correct HTTP session. This encoding is different from the one defined in [13.4 Encoding URLs](#) because in this case the `HttpSession` ID is encoded with parameter "`jsessionId`" as defined in the HTTP Servlet specification. The methods for this are defined on the `ConvergedHttpSession` interface.

14 Container Functions

14.1 Division of Responsibility

One of the goals of the SIP Servlet API is to keep the application programming model simple (see [1.1 Goals of the SIP Servlet API](#)). This is accomplished by delegating a number of tasks to the container and freeing up the applications to focus on their business logic. Tasks such as management of network listen points, SIP message and transaction processing, handling of headers (see [5.4.2 System Headers](#)) like Via, Contact, CSeq, RSeq, Path etc. are handled by the containers.

As a thumb rule, functionality that is "behind" the SIP Servlet API is managed by the container. As an example the API exposes the `SipSession` interface and allows applications to perform operations like creation of an in-dialog request, getting and setting attributes, invalidation etc on the `SipSession` but it does not allow the application to create a new `SipSession` explicitly. The `SipSession` is an object managed by the container and is created as a result of incoming request processing or creation of a new request by the application.

The API strives to strike a balance between the ability to write simple SIP Servlet applications but at the same time allowing for powerful constructs using the base API.

Consider the example of SIP specific Event handling [RFC 3265]. The SIP Servlet API provides for handling of SUBSCRIBE and NOTIFY requests. As these requests are dialog creating the container manages these SIP dialogs transparently for applications. However, the "Subscription" which is the application specific property is to be maintained by applications. The RFC 3265 specifies at least two mechanisms for installation of subscription, namely SIP SUBSCRIBE and administrative installation. Besides this the Subscription duration can be specified either through the SUBSCRIBE request's "Expires" header or it can be conveyed in the specific event package,

also it is possible for the application to accept the Subscription for a shorter duration and convey the same in a 200 OK Response's "Expires" header. Container however provides primitives like a Timer API, access to Session, Request etc to facilitate the applications carrying out such functions.

The above two examples illustrate that due care has been taken to provide the right level of abstraction to the application developer. Having said that, it is expected that some repetitive use cases will be found, some patterns which shall have to be re-used within SIP Servlet Applications. With this specification is provided a powerful mechanism for application composition [chapter [15 Application Selection And Composition Model](#)], we are hopeful that future applications shall use this feature to re-use application features or patterns.

14.1.1 Protocol Compliance

The responsibility of protocol compliance is again divided between the container and the applications. The protocol behavior of container managed objects and state like sessions, transactions, dialog state, address/URI formats etc. is all enforced by the container. The applications influence the state changes by exercising the API methods provided for such purpose. In some cases it is possible that the underlying state machine is in disagreement with the action that the application wishes to perform. In all such cases the container **MUST** throw an `java.lang.IllegalStateException`, these actions include not only the base RFC 3261 assertions and state machine but also the API constraints, expressed also as `java.lang.IllegalArgumentException`. As much as possible the more obvious of these violations are described in the API documentation, however it is not possible to list all such situations where these exceptions may be thrown. Implementors **MUST** make sure that the protocol behavior of all the container managed objects is consistent with the SIP RFCs they comply with.

Further the API declares the `Exceptions` that get thrown when either the state, context or parameters are not correct. Implementors **MUST** throw these exceptions when the specified conditions occur. As a general guideline even for the conditions which are not mentioned specifically in the API the implementations **SHOULD** -

- Throw an `IllegalArgumentException` when any argument in the method call is not suitable in the context, or the argument is an object the state of which is not suitable for the context
- Throw a `NullPointerException` where the method tries to set an attribute of the object and the null attribute is not allowed by the domain object.

- Throw an `IllegalStateException` where the method receiver (object on which the method is called) is not in a state where the invocation is suitable.
- Throw a `ConcurrentModificationException` if the base collection gets modified when the `Iterator` is being used unless the collection is otherwise concurrent.

The idea is to provide an environment of least surprise to the application developer.

14.2 Multihomed Host Support

Containers are required to manage the network listen points. An application might use the information about the used inbound interface to decide about invocation of specific features or an application may decide to mandate the use of a certain outbound interface in order to send the message to a specific domain or host.

The following are four use cases for a deterministic selection of network interface on a multihomed host:

- The container executes on a multi-homed host, with an interface used for SIP signalling and a different interface for management. The container should only receive and send SIP messages on the former interface.
- The container executes on a multi-homed host, which has a network interface into a trusted and a second network interface into an untrusted domain, these domains may represent an internal and an external network. The application logic could differ based on whether a request comes from a trusted or untrusted element. The application must be able to identify the inbound address on which the message was received. The use of the remote address or its type is not sufficient, as the address will most probably be assigned dynamically and both networks may use either public or private address ranges.
- The container executing on a multi-homed host, connected to a trusted and an untrusted network domain (same setup as in use case 2), is sending out a request. For security requirements internal and external traffic, e.g. traffic to own servers and external customers has to be separated on physical level. Thus the application needs the ability to mandate the use of a particular outbound interface.
- The container may execute on a multi-homed host which has network interfaces into a number of networks which are separate on the physical layer. Each network interface will be on separate IP addresses and may be able to each route to the same addresses. Some applications may wish to choose the network by selecting an interface instead of allowing the operating system routing tables to automatically handle the selection. In order to send originating requests to the correct network, the application needs the ability to mandate the use of a particular outbound interface.

The applications can access the list of SIP URIs which can be used by the application to send the outbound requests. The applications can access this through a `ServletContext` attribute `javax.servlet.sip.outboundInterfaces`.

The container **MUST** implement a method to set the outbound interface for a certain session: `setOutboundInterface(InetAddress address)`, this interface must then be used by the container for all messages associated with that `SipSession`. This method is provided on `Proxy`, `ProxyBranch` and `SipSession` object. Invocation of this method also impacts the system headers generated by the container for this `SipSession`, such as the `Via` and the `Contact` headers. The supplied IP address is used to construct these system headers.

This setting may be overwritten by subsequent calls to the method. In case there is no specific outbound interface set, the container shall apply default behaviour of interface selection.

Once the interface is selected on a multihomed host, the SIP servlet container must populate the Systems Headers (see [5.4.2 System Headers](#)) accordingly.

14.2.1 Application Composition on Multihomed Container

In case there are multiple applications involved in a call then the following must be the behavior of the `setOutboundInterface()` method (see [15 Application Selection And Composition Model](#) for details) -

- For last application in the chain acting as a proxy if it does not call `setOutboundInterface()`, then setting by previous applications (if any) is used.
- For the last application acting as a proxy if it invokes `setOutboundInterface()`, then it overrides the setting by the previous applications (if any).
- If any application in the chain acts as UAS (or a B2BUA), then the previous applications setting (if any) is not relevant anymore.

It is expected that the interface IP addresses returned by the `ServletContext` parameter `javax.servlet.sip.outboundInterfaces` shall include the IP addresses of virtual interfaces or publicly observable IP addresses of the load balancers in case of clustered setup. However, it shall be up to the applications to decide which interface to use.

Note: This mechanism of explicit outbound interface selection is to be used by the applications that are aware of the network topology by virtue of some configuration or other means and have strong reasons to use this mechanism. It is expected that most of the applications shall not need to use this advanced feature.

15 Application Selection And Composition Model

SIP servlet application servers are typically provisioned with many different applications. Each application provides specific functionality, but, by invoking multiple applications to service a call, the deployer can build a complex and complete service. This modular and compositional approach makes it easier for application developers to develop new applications and for the deployer to combine applications from different sources and manage feature interaction. A typical example from traditional telephony is a call-screening application and a call-forwarding application. If the application server receives an incoming INVITE destined to a callee who subscribes to both services, both applications should be invoked.


One key requirement for application composition is that the correct set of applications be invoked in the correct order to service a call. Therefore, an important function of SIP servlet application servers is to select applications for invocation and route SIP messages to them. When determining which applications to invoke and in which order, containers treat initial requests, subsequent requests and responses differently. Generally speaking, an initial request is a request for which the container has no a priori knowledge. It may or may not be a request capable of establishing a dialog [RFC 3261, Section 12]. Appendix B contains a precise definition of initial requests. Initial requests are routed based on an application selection process. The routing of an initial request establishes a path of applications. Subsequent requests and responses are then routed along the path taken by the initial request.

The principles of application independence and composition described in this section are adapted from the Distributed Feature Composition (DFC) architecture [DFC1998], [CN2004], [SE2005]. The original definition of DFC [DFC1998] has been improved as a result of extensive experience with it.

15.1 Application Selection Process

A key component of the SIP application selection procedure is a logical entity called the Application Router. The Application Router plays a central role in the SIP application selection process. A SIP servlet application server is thus made up of the container, the Application Router, and a number of applications, as shown in [Figure 15-1](#). The concerns of these three entities are cleanly separated, and this section discusses their separate roles.

Figure 15-1 The SIP Servlet Container, Application Router, and Applications



15.1.1 The Role of the Application Router

The Application Router is called by the container to select a SIP servlet application to service an initial request. It embodies the logic used to choose which applications to invoke. An Application Router is required for a container to function, but it is a separate logical entity from the container. The Application Router is solely responsible for application selection and must not implement application logic. For example, the Application Router cannot modify a request or send a response.

The Application Router implements the `SipApplicationRouter` interface, which defines the API between the container and the Application Router. There is no direct interaction between the Application Router and applications. It is also important to note that, besides the information passed by the container, the Application Router is free to make use of any other information or data stores. How it accesses that information and what information it makes use of is a matter of its implementation and is not defined in this specification.

The role of the deployer is defined in the Servlet API. The deployer in a SIP servlet environment controls application composition by defining and deploying the Application Router

implementation. Giving the deployer control over application composition is desirable because it is the deployer who is most aware of and responsible for the totality of services provided to his or her subscribers. Furthermore, this specification intentionally allows the Application Router implementation to consult arbitrary information or data stores. This is because the deployer maintains subscriber information and this information is often private and valuable.

15.1.2 The Role of Applications

In contrast to the deployer, application developers are solely concerned with the application logic of individual applications. An application executes independently of other applications and should not make any assumptions about the presence or location of other applications. This promotes modularity and re-use and facilitates application development.

As part of its application logic, an application may choose to proxy initial requests, relay initial requests as a B2BUA, terminate initial requests as a UAS, or send initial requests as a UAC. Although these actions influence application composition, the decision of which application to invoke for a request issued by an application rests solely with the Application Router. It follows that the Application Router must be aware of the intention of the application when it sends a request. In some cases, it can be implicitly inferred. However, in some cases the application MAY use a new API call to indicate its intention explicitly. This is crucial for correct application composition, and will be discussed in [15.2.2 Sending an Initial Request](#).

15.1.3 The Role of the Container

The container receives initial requests from external entities or from an internal application, calls the Application Router to obtain the application to invoke, and then invokes the selected application.

Since an Application Router is required for the functioning of a container, all compliant container implementations MUST provide at least a default Application Router as defined in Appendix C. The default Application Router enables the support of simple application composition based on a configuration file. Container implementations MAY also provide alternative implementations of Application Routers. In fact it is expected that implementations will each have their Application Routers suited to the needs of a deployment situation and means to configure them.


The container informs the Application Router when new SIP applications are deployed or undeployed using the `SipApplicationRouter.applicationDeployed()` and `SipApplicationRouter.applicationUndeployed()` callback methods respectively. Each SIP application packaged within an EAR is registered or unregistered with the Application Router individually.

15.1.4 Application Independence

When the container receives an initial SIP request from an external SIP entity, the container invokes the application selection process to select the first application to service this request. If this first application subsequently proxies the request, or if it acts as a B2BUA and on the UAC side sends a new initial request, the container again invokes the application selection process to select and invoke the next application often within the same container. In this way, as applications proxy or send requests, a chain of applications is created. The chain ends when either an application acts as a UAS or when no more applications are selected and the request is sent to an external SIP entity. It is important to note that the application composition happens as a side effect of applications sending the requests and no special APIs are used for composition.

There are two very important considerations that guide the way application composition happens on SIP Servlet containers.

- Each application in a chain shall see a consistent view of the context in which it is operating regardless of how many other applications are involved in the same communications session. Because the context is consistent regardless of the presence of other applications, application developers are guaranteed that they can write an application as if it were the only one involved in the communication session. Applications should not know nor need to know whether a request they send or proxy is going to be sent directly outside the container or to another application.
- Application composition chain shall be composed independent of the location of the application. It is possible that the constituent applications in the chain are spread over multiple container instances. In other words the invocation sequence should be independent of deployment.

Figure 15-2 Maintaining SIP Context with Multiple Applications Active on Same Container

Note: The SIP Servlet Applications are modeled as isolated, individually deployable component of larger composite applications or services. It is foreseeable that in large systems which maximally benefit from the innovations of unrelated developers and development projects, the relationship between deployed SIP Servlet application implementations will be defined solely by the deployer and/or the implementor of the Application Router. Composition of those SIP Servlet Applications into application usages of the SIP network itself is facilitated in this way, and the developer of a given SIP Servlet application should make no assumptions about the model for deployment and method of composition of these composite applications. In such cases it is possible that SIP Servlet applications may be deployed within a single SIP Servlet container or may be deployed in multiple SIP Servlet containers, each representing a different but compliant implementation of this specification.

15.1.5 Subscriber Identity and Routing Regions

SIP Servlet applications typically act on behalf of subscribers. Some applications are only invoked when their subscriber is originating a call, i.e. the subscriber is the caller. Examples include Speed Dial and Outbound Call Screening. Other applications are only invoked when their subscriber is receiving a call, i.e. the subscriber is the callee. Examples include Call Blocking and Location Service. Finally, there are applications that are invoked for both calling and called subscribers. For example Call Waiting and 3-Way Calling are invoked when their subscribers are either placing or receiving a call. There are also applications that provide services independent of any particular subscriber. For example, a call logging application may be invoked for all calls, regardless of whether the subscriber is a caller or callee.

The concepts of Subscriber Identity and Routing Region are introduced here so that an application may learn in which region it is invoked and on behalf of which subscriber. When the Application Router selects an application to service a request, it **MUST** specify the routing region in the `SipApplicationRouterInfo` object that it returns to the container. The routing region can be one of `ORIGINATING_REGION`, `TERMINATING_REGION`, `NEUTRAL_REGION` or one of their sub-regions. The Application Router **MUST** also specify the subscriber identity in the `SipApplicationRouterInfo` object when the routing region is set to either `ORIGINATING_REGION` or `TERMINATING_REGION`. The Application Router **MAY** specify the subscriber identity when the routing region is set to `NEUTRAL_REGION`. The subscriber identity is a `javax.servlet.sip.URI` object. The container **MUST** make the subscriber identity and routing region values as set by the Application Router available to the selected application using the `SipSession.getSubscriberURI()` and `SipSession.getRegion()` methods.

In most cases, an application selected to serve the calling subscriber is invoked in the originating region while an application selected to serve the called subscriber is invoked in the terminating region. Often, an application not designed to serve any specific subscriber (such as the call logging application mentioned above) will be invoked in the neutral region and the subscriber identity will not be set by the Application Router. (Note that the Application Router sets the routing region and subscriber identity along with the application name on the `SipApplicationRouterInfo` object that is returned to the container as part of the `SipApplicationRouter.getNextApplication()` call.)

The routing regions are extensible and an implementation of the Application Router may further subdivide these regions into smaller regions if necessary. For example, one may define new sub-regions - `TERMINATING_REGISTERED` and `TERMINATING_UNREGISTERED` - within the terminating region by extending `SipApplicationRoutingRegion` as shown below:

```
TerminatingRegisteredRegion extends SipApplicationRoutingRegion {
    public TerminatingRegisteredRegion () {
        super("TERMINATING_REGISTERED", TERMINATING_REGION);
    }
}

TerminatingUnRegisteredRegion extends SipApplicationRoutingRegion {
    public TerminatingUnRegisteredRegion () {
        super("TERMINATING_UNREGISTERED", TERMINATING_REGION);
    }
}
```

Deployers could thus define applications to be invoked in the above sub-regions. An application like the Location Service application defined in the `TERMINATING_REGISTERED` sub-region would be invoked for subscribers whose terminals have registered with the registrar; while applications like the Voicemail application defined in the `TERMINATING_UNREGISTERED` region would be invoked only for those subscribers who have not yet registered.

15.1.6 Routing Directives

Routing directives are provided by applications when they send or proxy an initial request. Routing directives provide a means for applications to indicate their intention clearly to the Application Router which can then perform application selection accurately. Sections 15.2 through 15.4 discuss routing directives in more detail.

15.2 Application Environment and Behaviour

This section is intended for application developers. The concepts relevant to application developers are discussed, including how an application indicates its routing intention, and what additional information related to application composition an application can obtain from the container.

15.2.1 Receiving an Initial Request

When an application is invoked to service an initial request, the container calls the application's `doXXX()` method with the `SipServletRequest` in a new `SipSession`. The application can obtain the identity of the subscriber that the application is serving, and the routing region in which it is invoked from the `SipSession` object. These attributes do not change throughout the lifetime of the `SipSession` object.

15.2.2 Sending an Initial Request

Whenever an application proxies or sends an initial request, the container invokes the Application Router to obtain the name of the next application that should service the request. The Application Router needs to know whether the sent request is related to a request that the application has received earlier. Though related, the new request could be a modified version of the originally received request or a brand new request. If the sent request is not related to any received request, the Application Router starts a new application selection process. On the other hand, if the sent request is related to a received request, the Application Router makes use of state information associated with the received request, and resumes the selection process from the previous state.

This is crucial to the selection of applications. As an example, consider a deployment where all initial INVITE requests received by the container are serviced by two applications, A and B, in this order. When an INVITE request is received externally, or if an application Z acts as a UAC and sends a INVITE request, the container invokes the Application Router. The Application Router starts the selection process from the beginning and selects A. When A is invoked and proxies the INVITE request, again the container invokes the Application Router. Now, the Application Router must be made aware that this is related to the previous request so that it can determine that A has already been invoked and that B must be invoked as the next application.

When an application acts as a UAC and sends a request, the request is not based on any previously received request. That is, the application intends the request to be regarded as a new request unrelated to any other request. Although the application intention is implicit in this case, it is clear that a new application selection process should take place.

Listing 15-1 Application acts as UAC and sends initial request

```
+-----+
|
| req = factory.createRequest
|      (appSession,          | req
|      "INVITE", from, to); |----->
| req.send();               |
|                           |
+-----+
```

req starts a new selection process

In the case of a proxy application, there is no ambiguity. It is clear that the application's intention is to send the request along the way, or in other words to continue the call.

Listing 15-2 Application proxies an initial request

```
+-----+
req1 |                               | req2
----->| proxy = req1.getProxy(); |----->
|      proxy.proxyTo(...);      |
|                               |
+-----+
```

req2 is a continuation of req1

If the application acts a B2BUA, and uses the `SipFactory.createRequest(SipServletRequest origRequest, boolean sameCallId)` or `B2buaHelper.createRequest(SipServletRequest origRequest, Map headers)` methods, then again it is clear the the application's intention is to continue the call.

Listing 15-3 Application acts as B2BUA to relay initial request

```

+-----+
req1 |                                     | req2
----->| req2 = factory.createRequest(req1, ...); |----->
      | req2.send();                       |
      |                                     |
+-----+

req2 is a continuation of req1

```

In the cases encountered so far, the application intentions are clear from context. Such intentions are referred to as routing directives. In the first case the application implicitly signaled a NEW directive, and in the second and third cases the application implicitly signaled a CONTINUE directive.

Unlike the cases described above, there are times where an explicit directive is required.

Listing 15-4 Application acts as B2BUA to relay initial request and signals the CONTINUE directive explicitly

```

+-----+
req1 |                                     | req2
----->| req2 = factory.createRequest(appSession, |----->
      | "INVITE", newFrom, req1.getTo());      |
      | // copy headers and content from req1  |
      | req2.setRoutingDirective(CONTINUE, req1); |
      | req2.send();                           |
+-----+

req is a CONTINUE

```

In case, an application wishes to send a new request based on a received request but intends the request to be regarded as a new request, thereby initiating a new selection process. The application could create the new request separately and copy all the headers and content, but it is far easier to create the request based on an existing request and then indicate a routing directive explicitly by making use of a method in the `SipServletRequest` class introduced in version 1.1,

```
setRoutingDirective(SipApplicationRoutingDirective dir, SipServletRequest
req)
```

Listing 15-5 Application acts as B2BUA to send NEW request

```

      +-----+
req1  |                                     | req2
- - - ->| req2 = factory.createRequest(req1, ...); |----->
      | req2.setRoutingDirective(NEW, null);    |
      | req2.send();                            |
      |                                         |
      +-----+

req2 is a NEW
```

In the above case, a new application selection process will begin for the req2.

Finally, a third directive, REVERSE, is necessary but is used less commonly. REVERSE is used when an application reverses the direction of the call. There are two cases. In the first case, an application invoked to service the caller for an initial request now wishes to place a call back to the caller. A practical example is a Call Waiting (CW) application. When a subscriber to CW places a call, CW is invoked to serve the caller. When this call is still up, the subscriber receives an incoming call. CW alerts the subscriber, who then activates CW to put the first callee on hold, and switch to the new caller. If the subscriber finishes the conversation with the new caller and hangs up the phone, forgetting that there is another call on hold, the CW feature places a call to the subscriber. In this call, the subscriber is now the callee.

In the second case, an application which is invoked to service the callee in a request, now wishes to place a call on behalf of that same subscriber as the caller. A practical example is a three-way calling (3WC) application. When the subscriber receives a call, the 3WC application is invoked to serve the subscriber as the callee. In the middle of the call, the subscriber activates the 3WC application to call a third person, with the subscriber as the caller.

In both cases, the direction of the call is reversed, and the applications must specify the REVERSE directive so the Application Router is able to select the next application correctly.

Listing 15-6 Application reversing the call direction by placing a call back to the caller

```

      +-----+
      |                                     |
req1  | req2 = factory.createRequest(appSession, |
      |                                     |
      +-----+
```

```

----->|  "INVITE", newFrom, req1.getFrom(),...);|
Caller   |  // copy headers and content from req1  |
         |  req2.setRoutingDirective(REVERSE, req1); |
<-----|  req2.send();                          |
         +-----+

req2 is a REVERSE

```

It is important to note that in the cases where an application is sending a request based on a previously received request, the transaction of the received request may have completed, or indeed the SIP dialog may have completed and terminated. Indeed, in 3WC the transaction is already completed, and in busy retry the SIP dialog has already terminated. The application may store the request for later use. Note also that the second argument to `setRoutingDirective` must be an initial request received by the application, i.e. `SipServletRequest.isInitial()` must be true. The following table summarizes the routing directive under different scenarios:

Table 15-1 Routing Directives

Application Action	Directive
Req = factory.createRequest(appSession, "INVITE", from, to);	NEW
request.getProxy().proxyTo()	CONTINUE
req2 = factory.createRequest(req1, ...); req2.send();	CONTINUE
req2.setRoutingDirective(directive, req1);	Explicit directive

Note: If recursion is enabled, and the container automatically proxies to contact addresses received in redirect (3xx) responses, these requests are treated as if the application explicitly called `proxyTo()` to proxy to these contact addresses. Therefore they also have the CONTINUE directive.

15.3 Application Router Behavior

When the container receives an initial request from an external entity, or when an application acts as a UAC and sends an initial request with a NEW routing directive, the application selection

process is started fresh. In this case, the Application Router is called with the following information:

- The `SipServletRequest`
- The routing directive which is NEW

Based on the supplied information, and its configuration of which subscriber subscribes to which set of applications, and any other information that it may wish to use (e.g. time of day, network condition, external subscriber profile database), the Application Router returns the following information:

- the name of the selected application
- the subscriber identity that the selected application is to serve
- the routing region that the application serves in
- an optional route, which can be local or remote
- a route modifier which tells the container how to interpret the route
- an optional stateInfo Serializable object

The Application Router can return routes to the container via its

`SipApplicationRouterInfo.getRoutes()` method. The routes can be external or internal. External routes are used by the Application Router to instruct the container to send the request externally. Internal route is returned when the Application Router wishes to modify the popped route header as seen by the application code (through the `SipServletRequest.getPoppedRoute()` API call). When the request is received by the container the request may have had a Route header belonging to the container, which the container removes and makes available through the provisions of 5.6.3 Popped Route Header . The route modifier returned by the Application Router tells the container how to make use of the routes returned by it and also how the popped route needs to be presented. The route modifier can be one of the following enum values - ROUTE, ROUTE_BACK, NO_ROUTE.

- ROUTE modifier indicates that `SipApplicationRouterInfo.getRoutes()` returns valid routes. It is up to the container to decide whether they are external or if an internal route was returned. All of the routes returned MUST be of the same type, so the container can make the determination by examining the first route only.
- ROUTE_BACK directs the container to push its own route before pushing the external routes obtained from `SipApplicationRouterInfo.getRoutes()`.

- `NO_ROUTE` indicates that Application Router is not returning any routes and the `SipApplicationRouterInfo.getRoutes()` value, if any, should be disregarded.

The behavior of container with respect to the route modifiers is explained in [15.4.1 Procedure for Routing an Initial Request](#).

The `stateInfo` `Serializable` object is useful for the Application Router to store state information from one invocation to the next. An Application Router implementation may choose to put any information in the `stateInfo` object, and this object is opaque to the container, and not accessible to the applications. Typically, an Application Router implementation may store information such as subscriber identity, routing region, the name of the application last invoked, and a precomputed list of applications that are to be invoked next.

If the selected application subsequently proxies or sends a new initial request based on the first one with a `CONTINUE` or `REVERSE` routing directive, again the Application Router is called. This time, in addition to the `SipServletRequest` and routing directive, it is also supplied with the `stateInfo` object that it previously returned. In this way, the Application Router delegates the maintenance of the application selection state to the container, and thus it can be stateless with respect to each initial request it processes.

If the Application Router determines that no application is selected to service a request, it returns null as the name.

15.3.1 Order of Routing Regions

Because proximity of an application to its subscriber confers priority, it is beneficial for the management of feature interaction that originating applications are closest to the caller, and that terminating applications are closest to the callee. This can be satisfied if the following rules are followed: a) The originating region applications should be invoked first followed by terminating region applications. b) The applications that service a subscriber are contiguous, i.e. no insertion of applications that service other subscribers in between.

On the other hand, it is entirely possible that the Application Router progresses directly to the terminating region if the caller is not a subscriber, or the caller does not subscribe to any applications. It is also possible that this application server does not serve any originating subscribers or has determined through some means that the originating applications have already been invoked and it should only look for terminating applications.

15.3.2 Inter-Container Application Routing

This specification supports applications distributed across multiple containers. The Application Router may return external routes in its `SipApplicationRouterInfo.getRoutes()` method that point to other application servers that it wishes the request to be routed to. The container **MUST** then push the routes onto the Route header stack of the request and send the request externally. If this request arrives at a container compliant with this specification, it will invoke the Application Router residing in that container so that the application selection process may continue. The first Application Router may pass any state information to the second Application Router by embedding it in the Route header. This is in accordance with the cascaded services model [SERL] as the applications can reside on different hosts and still participate in the application composition process.

Note: Some architectures require that the originating and terminating applications be hosted on different servers. The deployer can easily accomplish this by configuring the Application Routers such that one server hosts only originating and the other only terminating applications. Either the subscriber data can be partitioned such that the first server only serves originating users and the second serves terminating users or the Application Routers can collaborate by passing some state information in the Route headers, indicating for example that first server has already completed the invocation of originating services.

15.4 Container Behavior

The container is responsible for instantiating and initializing the Application Router and providing to it the initial list of deployed applications. Further when new applications are deployed or when applications are undeployed, the container must also inform the Application Router.

The container is responsible for receiving an initial request from an external entity or from an application, invoking the Application Router to obtain the name of the application to service the initial request and dispatching the request to the main servlet within this application as described in [16 Mapping Requests To Servlets](#). The container is also responsible for maintaining application selection state including:

- the routing directive associated with this request
- routing region (originating, terminating, or neutral)
- acting on the route returned from the Application Router in conjunction with the route modifier

- arbitrary, opaque state information returned from the Application Router

The next section details the procedures for routing an initial request and when the above fields are set, followed by a description of the application router packaging and deployment model.

15.4.1 Procedure for Routing an Initial Request

When the container receives a new initial request, it first creates and initializes the various pieces of application selection state as follows:

Directive:

- If request is received from an external SIP entity, directive is set to NEW.
- If request is received from an application, directive is set either implicitly or explicitly by the application.

Application router stateInfo:

- If request is received from an application, and directive is CONTINUE or REVERSE, stateInfo is set to that of the original request that this request is associated with.
- Otherwise, stateInfo is not set initially.

Subscriber URI and Routing Region:

These are not set initially.

The following procedure is then executed:

1. Call the `SipApplicationRouter.getNextApplication()` method of the Application Router object. The Application Router returns a `SipApplicationRouterInfo` object, named 'result' for this discussion.
2. Check the `result.getRouteModifier()`
 - If `result.getRouteModifier()` is ROUTE, then get the routes using `result.getRoutes()`.
 - If the first returned route is external (does not belong to this container) then push all of the routes on the Route header stack of the request and send the request externally. Note that the first returned route becomes the top route header of the request.
 - If the first returned route is internal then the container MUST make it available to the applications via the `SipServletRequest.getPoppedRoute()` method and

ignore the remaining ones, if any. This allows the AR modify the popped route before passing it to the application.

- If `result.getRouteModifier()` is `ROUTE_BACK` then push a route back to the container followed by the external routes obtained from `result.getRoutes()` and send the request externally. Note that the container route SHOULD include AR state encoded as a route parameter in order for the AR to continue processing the application chain once the request returns back to the container.
- If `result.getRouteModifier()` is `NO_ROUTE` then disregard the `result.getRoutes()` and proceed.

3. Check the `result.getNextApplicationName()`

- If `result.getNextApplicationName()` is not null:
 - set the application selection state on the `SipSession`: `stateInfo` to `result.getStateInfo()`, region to `result.getRegion()`, and URI to `result.getSubscriberURI()`.
 - follow the procedures of Chapter 16 to select a servlet from the application.
- If `result.getNextApplicationName()` is null:
 - If the Request-URI points to a different domain, or if there are one or more Route headers, send the request externally according to standard SIP mechanism.
 - If the Request-URI does not point to another domain, and there is no Route header, the container should not send the request as it will cause a loop. Instead, the container must reject the request with 404 Not Found final response with no Retry-After header.

While routing requests within the container to the next application in the chain if the route modifier as returned by the Application Router is not `ROUTE` then the container MUST NOT push routes just to route the requests to the applications. SIP servlet applications, however, are free to push internal and external routes if they so require.

Note: As a guideline it is strongly recommended for applications do not rely on Via or Record-Route headers for their application logic, as within the container the container implementations may chose to optimize the handling of the system headers by either aggregating them while going out or just replacing them with one container level header and maintaining the state internally. The applications SHOULD instead use the `SipServletMessage` methods `getLocalXXX`, `getRemoteXXX` methods if they are interested in the upstream entity.

If `SipApplicationRouter.getNextApplication()` throws an exception, the container should send a 500 Server Internal Error final response to the initial request.

15.4.2 Application Router Packaging and Deployment

The container is responsible for loading and instantiating Application Router implementations. In order to be portable across containers, the Application Router implementation **MUST** be packaged in accordance with the rules specified by the Java SE [Service Provider framework](#). Specifically, the jar file containing the Application Router implementation must include `META-INF/services/javax.servlet.sip.ar.spi.SipApplicationRouterProvider` file. The contents of the file indicate the name of the concrete public subclass of the `javax.servlet.sip.ar.spi.SipApplicationRouterProvider` class. The concrete subclass must have a no-arg public constructor.

As specified by the Service Provider framework, the providers may be installed by :

1. Including the provider jar in the system classpath
2. Including the provider jar in the extension class path
3. Container-specific means

If the container uses classpath-based deployment, the first Application Router jar file found in the classpath is installed. In order to avoid ambiguity when multiple Application Router implementations are present in the classpath, this specification also defines a system property which instructs the container to load a given provider. The

`javax.servlet.sip.ar.spi.SipApplicationRouterProvider` system property can be used to override loading behavior and force a specific provider implementation to be used. For portability reasons, containers that provide their own deployment mechanism for the Application Router **SHOULD** observe the system property, if specified by the deployer.

15.5 Application Names

In order to identify SIP servlet applications, a new mandatory element, `<app-name>`, is defined under the `<sip-app>` element in the `sip.xml` deployment descriptor file of each application. The names of Sip Servlet applications must be unique within a container instance or across clustered containers under common administrative control for application selection to function properly. It is recommended that application developers follow the Java class naming convention when naming applications, e.g. "org.sipservlet.app.voicemail", to avoid naming conflicts with other developers. The container is responsible for detecting and handling naming conflict, when a new application is deployed. Though how the application is identified within the container is something internal to it and is like a symbol, it is strongly recommended that containers make use of the deployment context while identifying the applications. Specifically the containers **SHOULD** create symbols identifying the application as

"myj2eeapp/mysipapp/org_sipservlet_app_voicemail" where the voice mail application with the following deployment descriptor is packaged in a war archive by the name of "mysipapp.war" which is packaged in an ear archive by the name of "myj2eeapp.ear".

Listing 15-7 Example of sip.xml file illustrating app-name and servlet-name

```
<sip-app>
  <app-name>org.sipservlet.app.voicemail</app-name>
  ...
  <servlet>
    <servlet-name>depositServlet</servlet-name>
    <servlet-
class>org.sipservlet.app.voicemail.DepositServlet</servlet-class>
    ...
  </servlet>
  <servlet>
    <servlet-name>retrievalServlet</servlet-name>
    <servlet-
class>org.sipservlet.app.voicemail.RetrievalServlet</servlet-
class>
    ...
  </servlet>
  ...
</sip-app>
```

If the aforementioned application was a v1.0 style application with similar packaging and were to be deployed on a container compliant with this specification, the application name **SHOULD** be "myj2eeapp/mysipapp".

The containers **MAY** further distinguish these names by adding their own versioning tokens.

If the application is specified by annotation, Chapter 18 provides the procedures for determining the name of such applications.

15.6 Responses, Subsequent Requests and Application Path

If an initial request results in a SIP dialog being established and at least one application is on the signaling path, be it a UA or a record- routing proxy, the container sets up state so that other requests in the same dialog can be routed to the set of applications within the container which are on the signaling path. Requests that are routed internally in a container based on such state are

called subsequent requests in this specification. Subsequent requests are not dispatched to applications based on the application selection process.

Correct routing of subsequent requests and responses can be achieved in several ways, and it is up to implementations to choose one. It is recommended that when routing messages internally between the applications the container **SHOULD** make use of the SIP mechanism of adding the Via header and adding the Record-Route (if it is record-routing proxy or a B2BUA) to allow for a robust stateless composition chain. However it is entirely possible that the implementations **MAY** chose to hide the internal topology by maintaining the application path state outside of SIP message. The implementation **MUST** however make sure that the parameters added to the Record-Route headers by individual applications are made available through the Route header on the subsequent request. The goal is to provide SIP consistency both at the application interface level and external SIP entity level.

Subsequent requests originated from the caller follow the path, or more precisely a subset of the path, of the corresponding initial request. Subsequent requests originated from the callee as opposed to the caller follow the reverse path. Responses always follow the reverse of the path taken by the corresponding request. This is true for responses to both initial and subsequent requests. The application path is a logical concept and as such may or may not be explicitly represented within containers.

For illustration, suppose three applications, A, B, and C, are invoked one after another to process an initial INVITE request. All three applications proxy but only A and C record-route, and so the application path consists of A and C along with associated contexts (sessions). A subsequent BYE request received from the callee will then be routed based on this application path, and will be passed first to C and then to A.

The distinction between initial and subsequent requests also applies to dispatching of locally initiated requests. If, for example, application A initiates an INVITE, and this is passed to application B which proxies it with record-routing enabled towards a destination outside the application server, then a subsequent BYE from A for the same dialog will be passed to B and then further downstream. Proxying of subsequent requests is discussed in [10.2.9 Handling Subsequent Requests](#).

It is worth noting that a `SipSession` can belong to at most one application path. This is because initial requests are processed in the context of new `SipSessions` and because there is a one-to-one correspondence between application paths and SIP dialogs. If the initial request results in more than one dialog being set up, the container will create derived `SipSessions` for the second and subsequent paths being created, see [6.2.3.2 Derived SipSessions](#).

When a `SipSession` terminates, either because the `SipApplicationSession` it belongs to times out or is explicitly invalidated or because the `SipSession` itself is explicitly invalidated, it is removed from the application path it was on, if any. If application paths are represented explicitly within containers, they are removed when the path becomes empty.

15.7 Transport Information

The `SipServletMessage` class provides methods for an application to obtain information about the transport used to receive a message. The container **MUST** return the following values irrespective of whether the request is initial or subsequent.

- `getRemoteAddr()` - **MUST** return address of the remote SIP interface if the request was received from an external entity but if the request was internally routed (from one application to the next on the same container) then it **MUST** return the address of the container's SIP interface
- `getRemotePort()` - **MUST** return port of the remote SIP interface if the request was received from an external entity but for the internally routed message the container is free to choose any port value consistent with one that would be chosen were the container to have actually sent the message out and back to itself through the local TCP/IP stack
- `getLocalAddr()` - address of the SIP listening interface on which this message was originally received from the external SIP entity throughout the chain.
- `getLocalPort()` - port number of the SIP listening interface on which this message was originally received from the external SIP entity throughout the chain.
- `getTransport()` - actual transport on which this message was received from the external entity and for the internally routed request - TCP or TLS to indicate that this is a reliable transport.
- `getInitialRemoteAddr()` - the IP address of the upstream/downstream hop from which this message was initially received by the container. This method returns the same value regardless of which application invokes it in the same application composition chain.
- `getInitialRemotePort()` - the port number of the upstream/downstream hop from which this message was initially received by the container. This method returns the same value regardless of which application invokes it in the same application composition chain.
- `getInitialTransport()` - the name of the protocol with which this message was initially received by the container. This method returns the same value regardless of which application invokes it in the same application composition chain.

- for the incoming TLS connection the X.509 certificate MUST be made available only to the first application that receives the message from outside of the container through the `javax.servlet.request.X509Certificate` for requests and `javax.servlet.response.X509Certificate` for responses. However, if the container is able to authenticate the remote user based on the credentials in the certificate then that authentication information MUST be made available to all the applications as described in 17.6 Server Tracking of Authentication Information

15.8 Popping of Top Route Header

15.8.1 Container Behavior

With application composition the requirement of popping the top route header belonging to the container [as described in [5.6.3 Popped Route Header](#)] is to be done on receipt of request not just on the network interface from external SIP entity but also within the container when the request is received from one of the applications. This will happen when one of the applications pushes a route header that happens to belong to the container. This popping of container's route can happen irrespective of the request being initial or subsequent. Besides this the Application Router can also modify the route header originally popped on initial request processing as per [15.3 Application Router Behavior](#).

15.8.2 Top Route Processing Examples

This section contains three examples to help illustrate the behavior of route popping, Application Router, `SipServletRequest.getPoppedRoute()` as well as the `SipServletRequest.getInitialPoppedRoute()` API methods.

15.8.2.1 Request with two route headers arriving at the container

This example illustrates an application chain with two applications (App1, App2), where neither application pushes routes onto the incoming request.

(Alice) - AR - App1 - AR - App2 - AR - (Bob)

1. Request R from Alice arrives at the container.
2. The request contains two route headers, A_i and B_e. A_i represents a route pointing to the container, B_e means it is an external route. The AR receives the request R.

3. Inside of the AR, `R.getPoppedRoute()` returns `A_i` as `A_i` is popped by the container prior to passing the request onto the AR. Also, `R.getInitialPoppedRoute()` returns `A_i`. R still contains `B_e` route header. The AR returns `App1`, and `NO_ROUTE` modifier to the container.
4. `App1` executes.
5. Inside of `App1`, `R.getPoppedRoute()` returns `A_i` and R's top route is `B_e`.
`R.getInitialPoppedRoute()` returns `A_i`. `App1` proxies request to Bob, without calling `pushRoute()`.
6. The AR intercepts the request again.
7. Inside of the AR, `R.getPoppedRoute()` returns null and R still contains `B_e` as its top route.
`R.getInitialPoppedRoute()` still returns `A_i`. AR returns `App2` and `NO_ROUTE` modifier to the container.
8. `App2` executes.
9. In `App2`, `R.getPoppedRoute()` returns null, and R's top route is `B_e`.
`R.getInitialPoppedRoute()` returns `A_i`. `App2` proxies R to R's Request-URI (Bob). The AR executes.

AR returns null application to the container, and `NO_ROUTE`. This causes the container to send the request to `B_e`, as per [15.4.1 Procedure for Routing an Initial Request](#), point 3.

15.8.2.2 Application pushes route pointing back at the container

This example illustrates container behavior when the application pushes a route onto the incoming request pointing back to the container.

(Alice) - AR - `App1` - AR - (Bob)

1. Request R from Alice arrives at the container.
2. The request contains two route headers, `A_i` and `B_e`. `A_i` represents a route pointing to the container, `B_e` means it is an external route. The AR receives the request R.
3. Inside of the AR, `R.getPoppedRoute()` returns `A_i` as `A_i` is popped by the container prior to passing the request onto the AR. `R.getInitialPoppedRoute()` returns `A_i`. R still contains `B_e` route header. The AR returns `App1`, and `NO_ROUTE` modifier to the container.
4. `App1` executes.

5. Inside of App1, `R.getPoppedRoute()` returns `A_i` and R's top route is `B_e`.
`R.getInitialPoppedRoute()` returns `A_i`. App1 calls `R.pushRoute(C_i)`, where `C_i` represents a Route header pointing back at the container. Next, App1 proxies the request to Bob. The AR intercepts the request again.

Inside of the AR, `R.getPoppedRoute()` returns `C_i`, and R's top route header is `B_e`.
`R.getInitialPoppedRoute()` returns `A_i`. AR decides to return null application to the container, and `NO_ROUTE`. This causes the container to send the request to `B_e`, as per [15.4.1 Procedure for Routing an Initial Request](#), point 3.

15.8.2.3 Application pushes external route

This example illustrates container behavior when an application pushes an external route.

(Alice) - AR - App1 - AR - (Bob)

1. Request R from Alice arrives at the container.
2. The request contains two Route headers, `A_i` and `B_e`. `A_i` represents a Route pointing to the container, `B_e` means it is an external Route. The AR receives the request R.
3. Inside of the AR, `R.getPoppedRoute()` returns `A_i` as `A_i` is popped by the container prior to passing the request onto the AR. `R.getInitialPoppedRoute()` returns `A_i`. R still contains `B_e` Route header. The AR returns App1, and `NO_ROUTE` modifier to the container.
4. App1 executes.
5. Inside of App1, `R.getPoppedRoute()` returns `A_i` and R's top route is `B_e`.
`R.getInitialPoppedRoute()` returns `A_i`. App1 calls `R.pushRoute(D_e)`, where `D_e` represents a Route header pointing to an external address. Next, App1 proxies the request to Bob. The AR intercepts the request again.

Inside of the AR, `R.getPoppedRoute()` returns null, and R's top route header is `D_e`.
`R.getInitialPoppedRoute()` returns `A_i`. AR decides to return null application to the container, and `NO_ROUTE`. This causes the container to send the request to `D_e`, as per [15.4.1 Procedure for Routing an Initial Request](#), point 3.

15.9 Examples

The first example uses a simple two applications scenario to illustrate the interactions between container, Application Router and applications along with the concept of application path and routing of subsequent requests and responses. It also serves to illustrate the concept of message

context discussed in [6.4.1 Message Context](#). The examples that follow the first examine more complex application composition scenarios.

15.9.1 Example with Two Applications

Consider the Location Service (LS) application in [1.6.1 A Location Service](#), and a new Speed-Dial (SD) application. This latter application allows a calling party to specify a short speed-dial number (e.g. "1"), which it translates into a complete address based on prior configuration. LS allows the called party to control where calls are received. For the purpose of illustration, assume that both SD and LS applications record-route.

The following diagram shows the routing of an initial INVITE request. Assuming that the caller and callee of the INVITE request are both subscribers served by this server, SD is needed to serve the caller and LS is needed to serve the callee. The following step-by-step walkthrough of the call flow indicates message context as a triplet (app, as, ss) where app is the application name, as specifies an application session, and ss specifies a `SipSession`.

Example of Application Composition:

```
-----> SD -----> LS ----->
```

1. The container receives an INVITE request. The INVITE does not belong to an existing SIP dialog and so the container calls the Application Router with the request to obtain the name of the application to invoke.
2. The Application Router determines SD is the first application to serve the caller. The Application Router returns the name "SD", the caller's identity, and the originating region to the container together with some state information.
3. The container invokes the SD application in the context of (SD, as1, ss1).
4. SD, based on the caller's identity, performs database lookup to obtain the caller's speed-dial settings and proxies to the full address that corresponds to the speed-dial number.
5. The container receives the proxied request and again calls the Application Router with the request and the state information.
6. Based on the state information, the Application Router determines that SD has already been invoked and there are no other applications for the caller. Assuming that there are no applications in the neutral region, the Application Router proceeds to the terminating region and determines that LS is needed to service the callee. It returns the name "LS", the callee's identity, and the terminating region to the container, together with some state information.

7. The container invokes the LS application in the context of (LS, as2, ss2).
8. LS, based on the callee's identity, performs database lookup to obtain the callee's location settings and proxies the request to destination d1.
9. The container receives the proxied request and calls the Application Router with the request and state information.
10. The Application Router determines that no applications are required for d1, and returns null to the container indicating there is no further application to be invoked.
11. The container proxies the request towards d1, that is, outside the application server.
12. The container receives a 200 (OK) response for the INVITE. The 200 response is passed upstream along the reverse path of the request, i.e. the container passes it first to the LS application in context (LS, as2, ss2) and then to the SD application in context (SD, as1, ss1), and then sends it to the caller.
13. Since the 200 response establishes a dialog, an application path is created. This means that the container will route subsequent requests in this dialog to applications along the signaling path. In this case the application path consists of (SD, as1, as2) and (LS, as2, ss2).
14. An ACK for the 200 is received. This is recognized as being a subsequent request and is associated with the previously established application path. (Incoming ACKs for non-2xx final responses are needed for protocol reasons only and are simply dropped.)
15. The container passes the ACK to the SD application in the context (SD, as1, ss1) and when the upcall to SD returns, it passes the ACK to the LS application in the context (LS, as2, ss2). When the upcall to LS returns, the container proxies the ACK outside the application server according to standard SIP routing rules.
16. Assuming that the callee hangs up first, the BYE is passed along the application path in the reverse direction, i.e. it is passed to LS first, then SD.

Note that, logically, each application has its own set of SIP client and server transaction objects operating in accordance with the transaction state machines specified in the SIP specification [RFC 3261, chapter 17]. Likewise, logically, each proxy application executes its own instance of the proxy logic, for example, it has its own response context [RFC 3261, section 16]. This specification then "augments" the RFC 3261 defined state machines with additional rules, for example, that 100 responses, ACKs for non-2xx responses, and responses for CANCELs are not delivered to the applications.

15.9.2 Simple call with no modifications of requests

As an example, assume that Alice, Bob, and Carol all subscribe to the following applications:

- Originating region applications:
 - Originating call screening (OCS)
- Terminating region applications:
 - Call Forwarding (CF)
 - Incoming Call Logging (ICL)
- Applications in both originating and terminating regions:
 - Call Waiting (CW)
 - 3-Way Calling (3WC)

Assume that all these applications proxy or relay the initial requests they receive without modification. Also assume that the Application Router selects applications based on a simple ordering: the originating applications in the order CW, 3WC, OCS and the terminating applications in the order ICL, CF, 3WC, CW.

When the container receives a call from Alice to Bob, it calls the Application Router which selects CW to service Alice, because CW is the first application in Alice's originating region. When CW relays the INVITE as a B2BUA, the container again calls the Application Router. Because the application selection process is carried forward from the INVITE that CW received to the INVITE that CW is now sending, the Application Router determines that CW has already been invoked based on the current state of the selection process and that the next application is therefore 3WC. The same process is repeated when 3WC relays the INVITE request and the OCS application is invoked next. When OCS relays the INVITE request, the Application Router determines that there are no more originating applications for Alice. The application selection process then proceeds to the callee's i.e., Bob's terminating region, where ICL, CF, 3WC and CW are then invoked in turn. During this call, we assume that CF, although present, is not active. That is, it does not forward the call to another URI. When Bob's CW relays the INVITE, there are no more terminating applications for Bob and thus the container sends the request to Bob's UA.

When complete, the application selection process yields the following application path:

```
(Alice) - CW - 3WC - OCS ----- ICL - CF - 3WC - CW - (Bob)
```

```
Alice's originating      |      Bob's terminating
```

15.9.3 Modification of Headers

An application may change the originating address and/or the terminating address by modifying specific headers in the SIP request. As mentioned above, applications are usually subscribed to and are invoked to service an address either in the originating region or the terminating region. Therefore, if an application modifies one or both addresses, in general the application selection process will be affected.

As an example, assume Bob subscribes to Call Forwarding and configures it to forward all calls to Carol. When Alice calls Bob, the Call Forwarding (CF) application is invoked in the terminating region to service Bob. CF proxies the INVITE request to Carol's Request-URI. At this point, the Application Router typically would stop invoking other terminating region applications for Bob, and instead begin invoking Carol's terminating region applications.

Following on from the example above, the invocation sequence will be:

```
(Alice) - CW - 3WC - OCS ----- ICL - CF ----- ICL - CF - 3WC - CW -
(Carol)
```

```
Alice's Originating    Bob's Terminating    Carol's Terminating
```

The following is an example where the caller's address information is modified, resulting in a different set of originating applications being selected. Consider that Alice works at home as a customer care agent for her employer. She now subscribes to an additional originating application - Agent Identification (AI). If Alice is not working, this application does nothing. However, after Alice logs on to the employer's system, when Alice places a call this application modifies the caller URI (e.g. the From header) in the INVITE request to be sip:customer-care@employer.com and relays the request. The Application Router would then select originating applications subscribed to by the new caller URI, for example a Supervisor Monitor (SM) application may be selected so the supervisor may listen in and coach Alice.

```
(Alice) - AI ----- SM - ...
```

```
Alice's Originating    customer-care's originating
```

As another example, consider the case when Alice calls Carol and they are engaged in a conversation. The application path looks like:

```
(Alice) - CW - 3WC - OCS ----- ICL - CF - 3WC - CW - (Carol)
```

```
Alice's Originating    Carol's Terminating
```

When Alice activates 3WC to call Bob, 3WC continues the INVITE request it received previously, but modifies the callee to Bob. Because 3WC is in the originating region and the caller address is not modified, the application selection is not affected at this point and the remaining originating application, OCS, is selected. Then the terminating applications of Bob are selected. The resulting application path would then look like:

```
(Alice) - CW - 3WC - OCS ----- ICL - CF - 3WC - CW - (Carol)
              \
              \ - OCS ----- ICL- CF - 3WC - CW - (Bob)

      Alice's Originating           Bob's Terminating
```

Note that this usually happens some time after the initial Alice-Carol call was established. 3WC maintains a reference to the INVITE request it received initially so that it may use it to initiate a call to another party if it is ever activated. The application selection state information associated with that INVITE request has the same lifetime as the request.

15.9.4 Reversing the Direction of the Call

This example illustrates the use of the REVERSE directive. Consider the case when Carol calls Alice and they are engaged in a conversation. The application path is:

```
(Carol) - CW - 3WC - OCS ----- ICL - CF - 3WC - CW - (Alice)
```

In this state, Alice's 3WC application was invoked in a call where Alice is the callee. However, if Alice activates 3WC to call Bob, 3WC sends a new INVITE request where Alice is the caller and Bob is the callee. This distinction affects the application selection process. Normally, if 3WC is simply continuing the request in the same direction towards Alice as the callee, the next application invoked to serve Alice would be CW. However, in this case, the correct next application is OCS, followed by Bob's terminating region applications. The resulting application paths should look like:

```
(Carol) - CW - 3WC - OCS ----- ICL - CF - 3WC - CW - (Alice)
                                          \
                                          \ - OCS ----- ICL -
- CF - 3WC - CW - (Bob)
```

In other words, 3WC was previously operating in Alice's terminating region, but for the new call, 3WC is operating in Alice's originating region. Therefore, the Application Router selects Alice's originating applications, and in this case the next such application is OCS. In fact, one would observe that the application path from Alice's 3WC to Bob is identical to the corresponding portion.

15.9.5 Initiating a New Request

Consider a network-based alarm application that calls its subscriber at a preset time. This application acts as a UAC and sends a new INVITE request towards the subscriber, i.e. with an implicit NEW directive.

Such a request is treated in the same way as a request received from an external SIP entity. The application selection process is started afresh.

Using the alarm application (AL) as an example, it creates and sends an INVITE request to Alice with the From header set to alarm-service@example.com. The Application Router begins from the originating region of alarm-service@example.com, but in this case there are no originating applications to invoke. Alice's terminating applications are then selected and invoked as usual, resulting in the application path shown in diagram below:

```
AL ----- ICL - CF - 3WC - CW - (Alice)
```

15.10 Loop Detection

There is a possibility that loops may occur in invoked applications, for example, A proxies to B which proxies back to A, etc. It is important that such loops be detected and handled. This specification does not mandate a particular mechanism. One strategy that is consistent with the cascaded services model is to mimic the existing SIP "inter-host" mechanisms for loop detection in the "intra-host" case of SIP servlet containers. In particular, a simple and effective mechanism is to decrement the value of the Max-Forwards header whenever a request is proxied internally, or whenever a request is forwarded by a servlet acting as a B2BUA. [10.2.10 Max-Forwards Check](#) discusses the issue of when to generate 483 (Too Many Hops) error responses on basis of the Max-Forwards header.

15.11 Session Targeting

There are three session targeting mechanisms supported by this specification. All of them allow the container to associate seemingly unconnected initial requests with the same `SipApplicationSession`. The encode URI mechanism, defined in v1.0 of this specification,

has been deprecated in v1.1 of this specification as it is problematic and does not mesh well with the v1.1 application composition mechanism. The Session Key based targeting mechanism introduced in v1.1 of this specification is now the preferred mechanism to associate a request with a particular `SipApplicationSession`. It is also more powerful than the deprecated encode URI mechanism. Finally, the optional Join [RFC 3911] and Replaces [RFC 3891] mechanisms also allow an initial request (INVITE) to be targeted to a specific `SipSession` (SIP dialog) thereby targeting its parent `SipApplicationSession`. For this discussion, any request that invokes one of these three session targeting mechanisms is termed a targeted request. Unlike normal initial requests, targeted requests target a particular `SipApplicationSession` object and hence, implicitly, a particular application. Further, for two of these mechanisms, the encode URI and the Join/Replaces support, the session targeting includes the targeting to an application in addition to the `SipApplicationSession` object. That is, a targeted request to an encoded URI or a request with either a Join or a Replaces header is targeted to a particular application before the Application Router has been invoked to perform application selection. Thus, in addition to describing the mechanisms themselves, the sections below must address how targeted request mechanisms are harmonized with Application Router based application composition.

15.11.1 Session Targeting and Application Selection

Targeted requests employing the encode URI mechanism or the Join/Replaces mechanism are considered initial requests and therefore the container must invoke the Application Router to service such targeted requests just as with any other initial request. However, to simplify the Application Router implementation logic in recognizing that a request is targeted, a `SipTargetedRequestInfo` object is provided by the container to the Application Router. If present in a call to `getNextApplication()`, this object indicates to the Application Router that the request is targeted and further provides two pieces of information:

1. Targeted Request Type - This enumerated value indicates that the request contains a Join header, contains a Replaces header, or is targeted to an encoded URI.
2. Application Name - This string identifies the name of the application that owns the `SipApplicationSession` to which this request is targeted.

Note: There is no targeted request type corresponding to the Session Key based targeting mechanism defined in the v1.1 specification. That is because this mechanism is different from the other two: the Session Key based targeting mechanism operates only after the container has called the Application Router's `getNextApplication()` method to determine the next application to be invoked. For all of the three session targeting mechanisms the servlet to be invoked within the application **MUST** be the designated "main" servlet as discussed in [16.2 Servlet Selection](#). (unless of course the encode URI

targeted application is a v1.0 based application). The sections below describe each of these three session targeting mechanisms in more detail.

15.11.2 Session Key Based Targeting Mechanism

In general, application invocation results in creation of an `SipApplicationSession` object belonging to the application to be invoked. Thus, the processing of an initial request normally results in the creation of a new `SipApplicationSession` instance.

However, sometimes it is required to route all requests for a subscriber, application, region combination (or other factors) to a single `SipApplicationSession` instance.

As an example consider the CW (Call Waiting) application invoked on behalf of the subscriber Carol in the originating region, whose URI appeared in the From header. The application chain is subsequently formed with other applications and the dialog is established.

Consider the following case: An initial request is received and during the application selection process the Application Router concludes that the CW application needs to be triggered; further if the subscriber being served is Carol and the region is originating then the container must associate this request with an existing `SipApplicationSession`. The CW application should have a way to indicate its desire for such an association.

The applications indicate their desire to associate the request with existing `SipApplicationSessions` through a key generating mechanism described below. An application can use a method annotation `@SipApplicationKey` within one and only one of its classes, generally the `Servlet` class. This annotated static public method takes the (read-only) `SipServletRequest` as its argument and returns a `String`. The returned `String` is then used as a "key" to lookup the `SipApplicationSession`. Note that the applications themselves may generate this key based on the contents of the request, the routing region, the subscriber or any other relevant information. This application generated key is used by the container as an index for the application sessions. As the `application-session-id` must have the property of being unique across applications in a JVM, the containers may prefix app-names or other such application specific identifiers to such keys before assigning them as `application-session-ids`. This ID would then inform the container which `SipApplicationSession` instance this initial request be routed to.

The static method annotated by `@SipApplicationKey` is prohibited from modifying the request in any way, and any attempts to do so SHOULD result in an `IllegalStateException` thrown by the `SipServletRequest`.

While processing the initial request after selecting the application, the container MUST look for this annotated static method within the application. If found, the container MUST call the method

to get the key and generate a `application-session-id` by appending some unique identifier. If the resultant `application-session-id` identifies an existing `SipApplicationSession` within the container JVM, then the container **MUST** associate this request with that `SipApplicationSession` rather than creating a new one.

If the `application-session-id` thus generated does not identify an existing `SipApplicationSession` instance then a new instance **MUST** be created.

The absence of the `@SipApplicationKey` method annotation indicates that the application does not want to use the Session Key based targeting mechanism and so the container **MUST** (by default) create a new `SipApplicationSession` instance for every initial request.

It is strongly recommended that the same `application-session-id` be returned by the `getId()` method of the `SipApplicationSession`. If the annotation exists then the container should use the return value of the `@SipApplicationKey` annotated method in the return value of `getId()` after having added a unique prefix as required.

As an example, consider that there is a chat room application which desires that all requests with request URI "`sip:mychatroom1@example.com`" be handled using the same `SipApplicationSession`. The application defines a method with annotation `@SipApplicationKey` that takes the `SipServletRequest` as the argument.

1. An initial request comes into the container and the Application Router is consulted by the container.
2. The Application Router indicates to the container that the chat application needs to be invoked for this request.
3. The container calls the method annotated with `@SipApplicationKey` and gets an application session key. (Since the application wishes to send all requests with a certain Request-URI to the same `SipApplicationSession` it returns some key, based on the Request-URI, like a hash or perhaps just the Request-URI itself).
4. The container prefixes the key with the application name for uniqueness and then uses the resultant `application-session-id` to check if a `SipApplicationSession` already exists.
5. If found, that `SipApplicationSession` is used to associate the `SipServletRequest` and `SipSession` with, otherwise a new `SipApplicationSession` is created.

So in this example the first request with the Request-URI, "`sip:mychatroom1@example.com`", causes the creation of a `SipApplicationSession` object. Other requests with the same Request-URI, though still initial, are routed to the same `SipApplicationSession` instance through the use of Session Key targeting.

Note that since it is left up to applications to generate the keys for session association, they are free to make use of any information at their disposal to generate such a key. This accords tremendous flexibility to applications to associate initial requests or new `SipSessions` with existing `SipApplicationSessions`.

Similar to the other session targeting mechanisms described below, further processing of the request can result in an application chain different from the already existing chain of which the targeted `SipApplicationSession` is a part.

As noted above, this Session Key based targeting mechanism is different from the other two session targeting mechanisms described below. This is because this targeting mechanism does not operate until after the Application Router has selected an application to be invoked. That is, the session key based targeting mechanism does not constrain the Application Router in any way as regards to its selection of which application to invoke. Consequently, for targeted requests that make use of session keys, the `SipTargetedRequestInfo` argument passed to the Application Router's `getNextApplication()` method is set to null.

15.11.3 The Encode URI Mechanism

Note: The encode URI mechanism is deprecated.

The `encodeURI()` method is defined on the `SipApplicationSession`. When called with a URI it encodes an identifier of the `SipApplicationSession` in the URI, for example by adding the `sipappsessionid` parameter. In the case of SIP and SIPS URIs, the container may also rewrite the host, port, and transport protocol components of the URI based on its knowledge of local listen points.

Subsequently if this encoded URI appears in the top Route header or as the Request-URI in a request received by the container then the container must associate this request with the identified `SipApplicationSession`. If the container is not able to find the `SipApplicationSession` using the encoded URI, then it should treat the request as a normal, untargeted initial request.

This `encodeURI` method allows applications to correlate events which would otherwise be treated as being independent, that is, as belonging to different `SipApplicationSessions`. For example, an application might send an instant message with an HTML body to someone. The IM body may then contain a SIP URI pointing back to the SIP servlet container and to the `SipApplicationSession` in which the IM was generated, thus ensuring that an INVITE triggered by the IM recipient accessing that URI is associated with this `SipApplicationSession` when received by the container.

Upon receiving an initial request for processing, a container **MUST** check the topmost Route header and Request-URI (in that order) to see if it contains an encoded URI. If it does, the

container **MUST** use the encoded URI to locate the targeted `SipApplicationSession` object. If a valid `SipApplicationSession` is found, the container must determine the name of the application that owns the `SipApplicationSession` object. When calling the Application Router's `getNextApplication()` method, the `SipTargetedRequestInfo` object must be populated to indicate that this request is a targeted request of type, "ENCODED_URI", and must supply the application name that owns the identified `SipApplicationSession` object. If the container is unable to identify the `SipApplicationSession` object with the encoded URI, it **MUST** not identify this request as a targeted one and the `SipTargetedRequestInfo` argument in the `getNextApplication()` method is set to null.

If an application receives a request targeted to an encoded URI and subsequently either proxies it or relays it onwards as a B2BUA without having changed the encoded URI appearing in the Request-URI, the container will once again receive the targeted request for processing. In such a case, the container **MUST** handle the request in the same manner as described in the paragraph above.

Note: It is unexpected that an application receiving a request targeted to an encoded URI would proxy it or relay it onwards. Depending upon application and Application Router implementations, such behavior could result in a routing loop. By mandating that the container handle such requests in a consistent manner, the policy is established that it is not the container's responsibility to detect and/or remedy the possibility of routing loops due to unexpected treatment of targeted requests by applications.

15.11.4 Join and Replaces Targeting Mechanism

Container support for Join [RFC 3911] and Replaces [RFC 3891] is optional in this specification. Support for [RFC 3911] and [RFC 3891] is indicated by the presence of the "join" and "replaces" option tags respectively in the `javax.servlet.sip.supported` list. (see [3.2 Extensions Supported](#)) If a container supports Join and Replaces, the procedure followed by the container in handling a request with a Join or Replaces header is as follows:

1. If a container supporting [RFC 3911] or [RFC 3891] receives an initial INVITE request with Join or Replaces header, the container must first ensure that the request passes the RFC-defined validation rules.
2. For a request that passes the validation step, the container **MUST** attempt to locate the `SipSession` object matching the tuple from the Join or Replaces header. In locating this `SipSession`, the container **MUST** not match a `SipSession` owned by an application that has acted as a proxy with respect to the candidate `SipSession`. It must only match a `SipSession` for an application that has acted as a UA. If the matching `SipSession` is found, the container

must locate the corresponding `SipApplicationSession` object along with the application name of the application that owns the `SipApplicationSession`.

3. The container MUST populate a `SipTargetedRequestInfo` object with the type corresponding to whichever header appears in the request (e.g. `JOIN` or `REPLACES`) and with the application name that owns the identified `SipApplicationSession` and `SipSession`. The container MUST then pass the request to the Application Router's `getNextApplication()` method as a targeted request i.e., along with the populated `SipTargetedRequestInfo` object.
4. If no `SipSession` matching the tuple is found, the container MUST pass the request to the Application Router as an untargeted request, i.e., where the `SipTargetedRequestInfo` argument to `getNextApplication()` is null.
5. If the Application Router returns an application name that matches the application name found in step 2, then the container must create a `SipSession` object and associate it with the `SipApplicationSession` identified in step 2. The association of this newly created `SipSession` with the one found in step 2 is made available to the application through the `SipSessionsUtil.getCorrespondingSipSession(SipSession session, String headerName)` method.
6. If the Application Router returns an application name that does not match the application name found in step 2, then the container invokes the application using its normal procedure, creating a new `SipSession` and `SipApplicationSession`.
7. If the Application Router returns an application name of null, then the container MUST follow the procedure given in [15.4.1 Procedure for Routing an Initial Request](#) with the following modification: If the Request-URI does not point to another domain, and there is no Route header, the container should not send the request as it will cause a loop. Instead, the container must reject the request with a 404 Not Found final response with no Retry-After header.

If an application receives a request targeted to a particular `SipApplicationSession` because of the presence of a Join or Replaces header, and that application subsequently proxies it or relays it onwards as a B2BUA without having changed or removed the Join/Replaces header, the container will once again receive the targeted request for processing. In such a case, the container MUST handle the request in the same manner as described in the procedure defined above.

Note: As with applications handling encoded URI requests, it is unexpected that an application receiving a targeted request due to a Join/Replaces header would proxy it or relay it onwards leaving the Join/Replaces header intact in the request. Depending upon application and Application Router implementations, such behavior could result in a routing loop. By mandating that the container handle such requests in a consistent

manner, the policy is established that it is not the container's responsibility to detect and/or remedy the possibility of routing loops due to unexpected treatment of targeted requests by applications. Container implementations MAY allow configuration options to automatically reject all requests with Replaces and Join headers unconditionally or else when no matching `SipSession` is found.

15.11.5 Resolving Session Targeting Conflicts

A given initial request could employ one, two or all three session targeting mechanisms. For example, a request whose Request-URI is an encoded URI could also contain a Join header and the application selected to handle this request could use the Session Key based targeting mechanism to identify an existing `SipApplicationSession`. In such cases, it is possible that one, two, or even three distinct `SipApplicationSessions` could be simultaneously targeted by the three different mechanisms when processing a single initial request. In the presence of such conflicts, the following priority order is to be used by the container to determine which of the several possible `SipApplicationSession` objects is targeted by a particular initial request:

1. Join/Replaces
2. Session key based targeting
3. Encoded URI

In other words, if the container supports the Join/Replaces session targeting, then if a request has a Join or a Replaces header, then that header value is used by the container to identify an existing `SipApplicationSession`. If an application and `SipApplicationSession` are thereby identified, then that `SipApplicationSession` is the one targeted by the request even if the request has a Request-URI that is an encoded URI recognized by this container. Furthermore, in this case when invoking the application identified by the Join or Replaces header, the container MUST not consult the application's session key method even when that method is present. However, if an initial request has a Request-URI that is an encoded URI recognized by the container, the container MUST call the application's session key generating method even though an `SipApplicationSession` has already been identified by the encoded URI. In this case, according to the above priority list, if the session key based targeting mechanism identifies a different `SipApplicationSession` than the one identified by the encoded URI, the `SipApplicationSession` used when the application is invoked is the one identified by the session key targeting not the one identified by the encoded URI.

16 Mapping Requests To Servlets

[15 Application Selection And Composition Model](#) described the Application Router component and how it is instrumental in application selection process. The Servlet mapping is a mechanism for the selection of a servlet after the application has been determined by the container.

16.1 Multiple Servlets

One SIP Servlet application can comprise of many different servlets, potentially written by different developers. All these servlets, however, are still related such that together they provide the services of a complete application. Roughly, a servlet can be considered as means for separating responsibilities between these servlet classes or modularization of project such that many developers can work on the same application.

Having said that, the logic coded as different servlets could very well be coded into just one servlet class in the application (with business logic in non Servlet classes). Multiple servlets in a SIP Servlet application is not strictly a specification requirement but can be used to modularize code and distribute the responsibility of handling different SIP messages amongst them; this specification supports multiple servlets packaged in a single application archive.

16.2 Servlet Selection

Once an application is selected to be invoked on an initial request, one of the servlets within the application is invoked next. We have seen that the application can comprise of multiple servlets. The v1.1 version of the specification introduces the declaration for a main servlet. One and only one servlet amongst the servlets in the application can be marked as main servlet. When using this mechanism for servlet selection, if there is only one servlet in the application then this

declaration is optional and the lone servlet becomes the main servlet. However, if there are multiple servlets present in the application while using this mechanism then one of those servlets MUST be declared as the main servlet. This declaration can be done in the `@SipApplication` annotation's "mainServlet" element (see [18.2.3 @SipApplication Annotation](#)) or by the presence of the element `<main-servlet />` in the deployment descriptor. The main servlet thus declared MUST then be invoked for the initial request of the selected application.

The main servlet is tasked with the responsibility of processing the initial request. The main servlet can also forward the request to another servlet within the application. It can do that by using the `RequestDispatcher` interface, see [6.2.6 The SipSession Handler](#). If need be, the servlets can pass request attributes while passing requests back and forth amongst them to convey some information.

The servlet that handles the request becomes the handler for the `SipSession` that may come into existence and will be delivered the subsequent requests and responses directly by the container. The handling servlet or the main servlet can also designate any other servlet as the handler for subsequent request or responses by calling `SipSession.setHandler()` method.

Note: The servlet mapping mechanism defined in version 1.0 of this specification based on an XML rules language can be used in this version of the specification. Developers familiar with the v1.0 servlet mapping scheme who find it sufficient for servlet selection can make use of it even in applications compliant with v1.1 (or later) of this specification. Developers who find the v1.0 scheme insufficient for servlet selection could use the new main servlet mechanism for servlet selection. However, this specification makes a restriction that only one servlet selection mechanism, either the `<main-servlet />` OR `<servlet-mapping />` shall be employed for a given application. Once the Application Router chooses a particular application for an initial request, the container chooses the servlet to be invoked either using the `<main-servlet />` declaration OR the `<servlet-mapping />` rules. Applications employing both mechanisms via annotations or descriptor declarations MUST fail to deploy. If multiple servlets are present in applications, one of the two servlet selection mechanisms MUST be used.

16.2.1 Compatibility with v1.0 Specification

The applications compliant with v1.0 of this specification (JSR 116) can be deployed on containers compliant with this 1.1 specification without any change. When the Application Router selects the v1.0 application for invocation the container MUST select the servlet going through the servlet mapping rules. The main-servlet declaration can be used only by applications written compliant to v1.1 (or later) of this specification.

17 Security

Servlet applications are created by application developers who give, sell, or otherwise transfer the application to a deployer for installation into a runtime environment. Application developers need to communicate to deployers how the security is to be set up for the deployed application. This is accomplished declaratively by use of the deployment descriptor mechanism.

This chapter describes deployment representations for security requirements. Similar to servlet application directory layouts and deployment descriptors, this chapter does not describe requirements for runtime representations. It is recommended, however, that containers implement the elements set out here as part of their runtime representations.

17.1 Introduction

A servlet application represents resources that can be accessed by many users. These resources are accessed over unprotected, open networks such as the Internet. In such an environment, a substantial number of servlet applications will have security requirements.

Although the quality assurances and implementation details may vary, servlet containers have mechanisms and infrastructure for meeting these requirements that share some of the following characteristics:

- **Authentication:** The means by which communicating entities prove to one another that they are acting on behalf of specific identities that are authorized for access.
- **Access control for resources:** The means by which interactions with resources are limited to collections of users or programs for the purpose of enforcing integrity, confidentiality, or availability constraints.

- **Data integrity:** The means used to prove that information has not been modified by a third party while in transit.
- **Confidentiality or data privacy:** The means used to ensure that information is made available only to users who are authorized to access it.

17.2 Declarative Security

Declarative security refers to the means of expressing an application's security structure, including roles, access control, and authentication requirements in a form external to the application. The deployment descriptor is the primary vehicle for declarative security in servlet applications. The deployer maps the application's logical security requirements to a representation of the security policy that is specific to the runtime environment. At runtime, the servlet container uses the security policy representation to enforce authentication and authorization.

The security model applies to servlets invoked to handle requests on behalf of either caller or callee. The security model does not apply when a servlet uses the `RequestDispatcher` to invoke a static resource or servlet using a `forward`.

17.3 Programmatic Security

Programmatic security is used by security aware applications when declarative security alone is not sufficient to express the security model of the application. Programmatic security consists of the following methods of the `SipServletMessage` interface:

- `getRemoteUser`
- `isUserInRole`
- `getUserPrincipal`

The `getRemoteUser` method returns the user name the client used for authentication. The `isUserInRole` method determines if a remote user is in a specified security role. The `getUserPrincipal` method determines the principal name of the current user and returns a `java.security.Principal` object. These APIs allow servlets to make business logic decisions based on the information obtained.

If no user has been authenticated, the `getRemoteUser` method returns `null`, the `isUserInRole` method always returns `false`, and the `getUserPrincipal` method returns `null`.

Note: A response may contain credentials of the UAS. For this reason, the programmatic security methods apply to responses as well as to requests. However, since there is no mechanism for a

proxy to challenge a UAS upon unsuccessful response authentication, the SIP deployment descriptor cannot express a requirement that responses be authenticated.

The `isUserInRole` method expects a `String` `user` `role-name` parameter. A `security-role-ref` element should be declared in the deployment descriptor with a `role-name` sub-element containing the rolename to be passed to the method. A `security-role` element should contain a `role-link` sub-element whose value is the name of the security role that the user may be mapped into. The container uses the mapping of `security-role-ref` to `security-role` when determining the return value of the call.

For example, to map the security role reference "FOO" to the security role with `role-name` "manager" the syntax would be:

```
<security-role-ref>
  <role-name>FOO</role-name>
  <role-link>manager</role-link>
</security-role-ref>
```

In this case, if the servlet called by a user belonging to the "manager" security role made the API call `isUserInRole("FOO")`, the result would be true.

If no `security-role-ref` element matching a `security-role` element has been declared, the container must default to checking the `role-name` element argument against the list of `security-role` elements for the servlet application. The `isUserInRole` method references the list to determine whether the caller is mapped to a security role. The developer must be aware that the use of this default mechanism may limit the flexibility in changing rolenames in the application without having to recompile the servlet making the call.

17.4 Roles

A security role is a logical grouping of users defined by the application developer or assembler. When the application is deployed, roles are mapped by a deployer to principals or groups in the runtime environment. A servlet container enforces declarative or programmatic security for the principal associated with an incoming request based on the security attributes of the principal. This may happen in either of the following ways:

1. A deployer has mapped a security role to a user group in the operational environment. The user group to which the calling principal belongs is retrieved from its security attributes. The principal is in the security role only if the principal's user group matches the user group to which the security role has been mapped by the deployer.

2. A deployer has mapped a security role to a principal name in a security policy domain. In this case, the principal name of the calling principal is retrieved from its security attributes. The principal is in the security role only if the principal name is the same as a principal name to which the security role was mapped.

17.5 Authentication

A SIP user agent can authenticate a user to a SIP server using, for example, one of the following mechanisms:

- SIP digest authentication
- The P-Asserted-Identity header, thus trusting an upstream/downstream proxy to have authenticated the caller/callee, as specified in [privacy]
- The Identity and Identity-Info headers, as specified in [RFC 4474]

The HTTP Servlet specification also allows user authentication based on SSL. With SSL, TLS, and IPSec, it is actually the previous hop that is being authenticated and from which the user-data constraints are enforced. As proxies are more common in SIP, and may not be strongly associated with the UAC, the entity authenticating itself on an incoming TLS connection is not, generally speaking, the UAC itself.

For this reason SIP servlet containers will not typically perform authentication based on credentials received as part of a TLS handshake. However, it is possible that in some environments it is known that the TLS connection really does identify the UAC and in such cases it is reasonable for the container to reflect this knowledge in its implementation of the declarative and programmatic security features discussed here. These security features relate to end-users, not proxies, but it is up to individual containers to determine what constitutes an authenticated message.

17.6 Server Tracking of Authentication Information

As the underlying security identities (such as users and groups) to which roles are mapped in a runtime environment are environment specific rather than application specific, it is desirable to:

1. Make login mechanisms and policies a property of the environment the servlet application is deployed in.
2. Be able to use the same authentication information to represent a principal to all applications deployed in the same container, including converged applications, and

3. Require re-authentication of users only when a security policy domain boundary has been crossed.

Therefore, a servlet container is required to track authentication information at the container level (rather than at the servlet application level). This allows users authenticated for one servlet application to access other resources managed by the container permitted to the same security identity.

17.7 Propagation of Security Identity in EJBTM Calls

A security identity, or principal, must always be provided for use in a call to an enterprise bean. The default mode in calls to enterprise beans from servlet applications is for the security identity of a user to be propagated to the EJBTM container.

In other scenarios, containers are required to allow users that are not known to the servlet container or to the EJBTM container to make calls:

- SIP servlet containers are required to support access to applications by clients that have not authenticated themselves to the container.
- Application code may be the sole processor of signon and customization of data based on caller identity.

In these scenarios, a servlet application deployment descriptor may specify a `run-as` element. When it is specified, the container must propagate the security identity of the caller to the EJB layer in terms of the security role name defined in the `run-as` element. The security role name must have one of the security role names defined for the servlet application.

For servlet containers running as part of a Java EE platform, the use of `run-as` elements must be supported both for calls to EJB components within the same Java EE application, and for calls to EJB components deployed in other Java EE applications.

17.8 Specifying Security Constraints

Security constraints are a declarative way of annotating the intended protection of applications. A SIP servlet security constraint consists of the following elements:

- resource collection (a set of servlets)
- type of authentication
- authorization constraint

- user data constraint

A resource collection is a set of servlets and SIP methods. A servlet may have one or more security constraints associated with it. Before invoking a servlet to handle an incoming request, the container must ensure that all security constraints associated with that servlet are satisfied. If this is not the case, the request must be rejected with a 401 or 407 status code.

Note: SIP servlet resource collections are identified by names of servlets being invoked instead of URL patterns as in the HTTP Servlet API.

The authentication type is an indication of whether the container should return a 401 (Unauthorized) or 407 (Proxy Authentication Required) response status code when authenticating an incoming request.

An authorization constraint is a set of security roles at least one of which users must belong to for access to resources described by the resource collection. If the user does not belong to an allowed role, the user must be denied access to the resource. If the authorization constraint defines no roles, no user is allowed access to the portion of the servlet application defined by the security constraint.

A user data constraint describes requirements for the transport layer of the client server. The requirement may be for content integrity (preventing data tampering in the communication process) or for confidentiality (preventing reading while in transit). The container must at least use TLS to respond to requests to resources marked integral or confidential. If the original request was over TCP, the container must redirect the client to the TLS port.

The `login-config` element allows for configuration of the authentication method that should be used, the realm name that should be used for this application, and the configuration of a identity assertion scheme.

This specification introduces the `identity-assertion` element in `login-config` to specify in the deployment descriptor the mechanism to be used for identity assertion. This element can use one of the two identity assertion mechanisms:

- **P-Asserted-Identity:** Identity MUST be asserted using the P-Asserted-Identity header as described in [RFC 3325]. This mechanism is limited to trusted domains.
- **Identity:** Identity MUST be asserted using the Identity and Identity-Info headers as described in [RFC 4474]. This mechanism provides a cryptographic approach to assure the identity of the end users that originate SIP requests, especially in an interdomain context.

An application can also choose if the identity assertion scheme specified is REQUIRED by the application or just SUPPORTED using the element.

When P-Asserted-Identity scheme is REQUIRED by the application, the P-Asserted-Identity header MUST be present in the request. If the P-Asserted-Identity header is not present the container MUST reject the request with a 403 response. If authorization of the Identity specified by P-Asserted-Identity header fails, the container MUST return a 403 response.

When P-Asserted-Identity scheme is SUPPORTED by the application, the container checks the presence of the P-Asserted-Identity header in the message. If the header is not present then any other authentication mechanism configured by the user as part of the login configuration is used, for example, Digest Authentication.

When using the Identity mechanism, the Identity and Identity-Info headers are used to assert the identity of the user as defined in RFC 4474. The REQUIRED or SUPPORTED behavior is same as described above for the P-Asserted-Identity mechanism except that the container MUST return a 428 error response if the required headers are not present in the request.

For example, here is a sample configuration for the `login-config` element:

```
<login-config>
  <auth-method>DIGEST</auth-method>
  <realm-name>example.com</realm-name>
  <identity-assertion>

  <identity-assertion-scheme>P-Asserted-Identity</identity-assertion-scheme>
    <identity-assertion-support>SUPPORTED</identity-assertion-support>
  </identity-assertion>
</login-config>
```

In the above example, for every servlet that has a security constraint configured in the application, the container would use the P-Asserted-Identity header for authorization, if present. If P-Asserted-Identity header is absent then the request will be subjected to the Digest authentication mechanism configured in the element.

17.9 Default Policies

By default, authentication is not needed to access resources (servlets). Authentication is needed for requests for a resource collection only when specified by the deployment descriptor.

17.10 Authentication of Servlet Initiated Requests

This section defines the mechanism to provide authentication information on Servlet initiated requests which are challenged by the remote Proxy or UAS.

A new API introduced in this specification enables:

- the applications to set the authentication parameters while letting the container do the actual processing of the authentication headers that should be added to the request.
- the application and the container to store authentication information into an object that can be used in other requests or when additional challenges are received either from a proxy/proxies (407), server/servers (401). Note that when forking is done a 401 challenge may be received from multiple servers (but for the same realm).
- the container to add the required authentication headers to the request prior to sending it again.

The container **MUST** implement the logic to manipulate the authentication headers.

`AuthInfo` is an object that may be saved and passed to the container by the application. Only the container has knowledge of its content and format. Therefore, different implementations of the container can put different content in the object according to the implementation and according to the actual authentication mechanisms that are implemented by the container.

```
AuthInfo SipFactory.createAuthInfo()
    Returns an empty instance of AuthInfo.
```

Following new APIs are added to support this mechanism -

```
Iterator SipServletResponse.getChallengeRealms()
    Returns an iterator over all the realm(s) associated with the current
    challenge response.
```

```
AuthInfo.addAuthInfo(int statusCode, java.lang.String realm,
    java.lang.String userName,
    java.lang.String passWord)
```

- `statusCode` is the 401/407 that was returned in the challenge that is being provided for in the API call.
- `realm` is the realm that was returned in the challenge the is being provided for in this API call

This adds authentication info in the `AuthInfo` object for the challenge of Type and Realm.

```
SipServletRequest.addAuthHeader(SipServletResponse challengeResponse,  
                                AuthInfo authInfo)
```

This adds the appropriate authentication header to the request. After doing this the application can resend the request again. Note that if the request contains multiple challenges the container can add multiple authentication headers.

```
SipServletRequest.addAuthHeader(SipServletResponse challengeResponse,  
                                java.lang.String userName,  
                                java.lang.String passWord)
```

This adds the appropriate authentication header to the request. After doing this the application can resend the request again. This is a shortcut API to be used in cases where the application does not want or needs to use the `AuthInfo` object.

17.10.1 Description of the Procedure

When receiving a challenge response (401/407) on a prior request, the application can get the challenge type (status code) and the realm from the challenge response. The application can then provide authentication parameters that will be added to the `AuthInfo` object. After adding the parameters to the `AuthInfo` object, it is possible to ask the container to add the appropriate header to the request and after that the application can send the request again. If the 2xx response to a request that was successfully authenticated included the Authentication-Info header, the container MAY make use of it while sending subsequent requests.

Security

18 Java Enterprise Edition 5 Container

18.1 Java 5

J2SE 5.0 is the minimum version of the underlying Java platform with which SIP Servlet containers compliant with this specification must be built.

Even though the API may be revised to make use of new J2SE 5.0 constructs it must remain fully backwards compatible for the applications written for the 1.0 version of this specification.

18.2 Annotations and Resource Injection

The Java Metadata specification (JSR-175) provides a means for an application to provide configuration and dependency information in the Java code. This metadata, or annotations, are used to provide a faster and easier way for application development.

This section describes annotations and resource injection in a SIP Servlet 1.1 compliant container, reflecting the work done in the Java EE through the Common Annotations for the Java Platform (JSR-250) and the Servlet 2.5 specification (JSR-154 Version 2.5).

The goal of this annotations work is to eliminate the need for the deployment descriptor. Here is a mapping from the deployment descriptor to the function in the annotations.

Table 18-1 Mapping from deployment descriptors to annotations

Item	Deployment descriptor	Annotation element
Icons	small-icon or large-icon	@SipApplication smallIcon or largeIcon

Table 18-1 Mapping from deployment descriptors to annotations

Display Name	display-name	@SipApplication displayName
Description	description	@SipApplication description
Distributable	distributable	@SipApplication distributable
Context parameters	context-param , param-name , param-value	Not applicable.
Listener	listener-class	@SipListener
Servlet name	servlet-name	@SipServlet name
Application name	app-name	@SipApplication name
Servlet class	servlet-class	@SipServlet
Initialization parameters	init-param , param-name , param-value	Not applicable, suggested to use constants.
Startup order	load-on-startup	@SipServlet loadOnStartup
Proxy timeouts	proxy-timeout	@SipApplication proxyTimeout
Session timeouts	session-timeout	@SipApplication sessionTimeout
Resources	resource-*	@Resource , @Resources
Security Roles	security-role *	@DeclaresRole
EJBs	ejb-ref *	@EJB
Run as	run-as *	@RunAs
Web Services	Not applicable	@WebServiceRef

18.2.1 Servlet 2.5 alignment

Reflecting on the work done in the Servlet 2.5 (JSR-154 Version 2.5) specification, this specification will require the components defined in that specification be supported. At the time

of this writing, the annotations defined in the Servlet 2.5 specification included `@Resource`, `@Resources`, `@PostConstruct`, `@EJB`, `@WebServiceRef`, `@DeclaresRole`, and `@RunAs`. Those definitions in the Servlet 2.5 specification remain the same with the following additional classes required to be injected.

Table 18-2 Interfaces and Classes which Require Annotation Support

Servlet	<code>javax.servlet.sip.SipServlet</code>
Listeners	<code>javax.servlet.sip.SipApplicationSessionListener</code>
	<code>javax.servlet.sip.SipApplicationSessionActivationListener</code>
	<code>javax.servlet.sip.SipSessionAttributeListener</code>
	<code>javax.servlet.sip.SipSessionListener</code>
	<code>javax.servlet.sip.SipSessionActivationListener</code>
	<code>javax.servlet.sip.SipErrorListener</code>
	<code>javax.servlet.sip.TimerListener</code>

With this specification, annotations defined as part of Servlet specification 2.5 and some common Java EE annotations as defined in JSR 250 are also included. For a description of annotations listed but not described here please refer to Servlet specification 2.5 and JSR 250. The complete list of annotations required to be supported is -

- `@RunAs`
- `@DeclaresRole`
- `@Resource`
- `@Resources`
- `@EJB`
- `@WebServiceRef`
- `@PostConstruct`
- `@PreDestroy`
- `@SipServlet`
- `@SipApplication`

- `@SipListener`
- `@SipApplicationKey`

Resource Injection is supported for classes whose lifecycle is controlled by the container. A non Java EE compliant implementation of this specification MUST support SIP specific annotations but it is not required to support the Java EE specific annotations.

18.2.2 @SipServlet Annotation

The `SipServlet` annotation allows for the `SipServlet` metadata to be declared without having to create the deployment descriptor. Certain values from the deployment descriptor are not needed since these are declared in the annotation in the source file of the servlet itself. First from the deployment descriptor, `servlet-class` is not needed since the annotation is declared in the class. Also, `init-param` is not useful since it could just as easily be a static constant in the source file where the annotation exists.

The first element is the `name`. The `name` element is equivalent to the `servlet-name` element in the deployment descriptor. If the name is not provided, the servlet name used will be the short name of the class annotated.

The second element is the `applicationName` that this servlet is associated with. The application will contain the proxy settings, session settings, listeners, and some basic configuration. The `applicationName` element is an optional element. If the application name is not set, and it is not available from the deployment descriptor, the container must check if the package this servlet class is within is a package which has a `@SipApplication` annotation. If so, the container must bind this servlet to that application without the application name being set. If not, the container SHOULD treat this as a deployment error, as specified in [18.2.3 @SipApplication Annotation](#).

The `description` element in this annotation contains the declarative data of a servlet and can help the consumer of this application understand better what the `SipServlet` is and what it does. This element is optional and an empty string will be used in its place if one is not provided.

The `loadOnStartup` element of the annotation is similar to the `load-on-startup` element of the deployment descriptor. This element specifies the starting order of the servlet application within the system. If the value is a negative number, the container may choose when to start up this servlet. If the value is zero or a positive number, the container must start the servlets beginning with the lowest number. If servlets have the same `loadOnStartup` value, the container may choose which one to start first. This element defaults to a negative number, allowing the container to choose when to startup this servlet if no other value is given.

Here is the definition of the `@SipServlet` annotation:

```

package javax.servlet.sip.annotation;
import java.lang.annotation.Inherited;
import java.lang.annotation.Retention;
import java.lang.annotation.Target;
import static java.lang.annotation.ElementType.TYPE;
import static java.lang.annotation.RetentionPolicy.RUNTIME;
@Target({TYPE})
@Retention(RUNTIME)
@Inherited
public @interface SipServlet {
    String name() default "";
    String applicationName() default "";
    String description() default "";
    int loadOnStartup() default -1;
}

```

The first example is a basic annotated class to become a `SipServlet`. Here, because no elements are specified, the name and display name of the servlet will be "Weather"

```

package com.example;
import javax.servlet.sip.SipServlet;
@SipServlet
public class Weather extends SipServlet {

```

Here in the second example, the servlet name is "WeatherService" and in the absence of application name, the display name defaults to "WeatherService".

```

package com.example;
import javax.servlet.sip.SipServlet;
@SipServlet (name = "WeatherService")
public class Weather extends SipServlet {

```

In this next example, the servlet name is "WeatherService" and since the application name is specified, the display name is "WeatherApplication".

```

package com.example;
import javax.servlet.sip.SipServlet;
@SipServlet (name = "WeatherService",
    applicationName = "WeatherApplication")
public class Weather extends SipServlet {

```

The class annotated by the `@SipServlet` annotation must override all methods it wishes to implement functionality for. For example, if the class wanted to perform some action on INVITE methods, it would do the following:

```
package com.example;
import javax.servlet.sip.annotation.SipServlet;
import javax.servlet.sip.SipServletRequest;
import javax.servlet.sip.SipServlet;
@SipServlet (name = "WeatherService",
             applicationName = "WeatherApplication",
             description = "Provides weather information",
             loadOnStartup=1)
public class Weather extends SipServlet {
    public void doInvite(SipServletRequest req){
        //perform action
    }
}
```

18.2.3 @SipApplication Annotation

The `@SipApplication` annotation is used to create an application level annotation for a collection of `SipServlets` and to maintain the application level configuration which was provided by the DD. The SIP application is a logical entity which contains a set of servlets and listeners with some common configuration. This annotation provides that configuration. Since this annotation is at the package level, all servlets within this package will belong to the same application unless they specify another application name using the `@SipServlet (applicationName)` annotation.

Note that regardless of the method used (DD or annotations), there can only be one SIP application registered with the container per .war or .sar archive. Violations of this rule SHOULD result in a deployment error.

The first element is the `name`. The `name` element is equivalent to the `app-name` in the deployment descriptor, and like in the deployment descriptor, it is mandatory. This is the name the listeners and servlets reference when adding themselves to this logical application. Note that v1.1 applications require that the application name be specified by either the deployment descriptor or this element. Failure to do so SHOULD result in a deployment error.

The second element is the display name of the application. The `displayName` element defaults to the application name unless specified through its own element in the annotation. This element is equivalent to the `display-name` element in the deployment descriptor.

The `largeIcon` and `smallIcon` elements in this annotation allow for specific icons to be used for this application. These are equivalent to the `large-icon` and `small-icon` elements in the

deployment descriptor. Each takes a simple string specifying the location of the image relative to the root path of the archive this class exists within. These elements are optional.

The `description` element in this annotation is equivalent to the `description` element in the deployment descriptor. This can help the consumer of this application understand better what the application is and what it does. Often, tools that consume applications will provide the descriptions to better understand what the application does. This element is optional and a empty string will be used in its place if one is not provided.

The `distributable` element indicates to the container whether this application is developed to properly function in a distributed environment. This element set equal to true is the equivalent to the `distributable` attribute in the deployment descriptor. Just as the descriptor without the `distributable` element defaults the application to not be distributable, the `distributable` element in the annotation defaults to false and must be set to true if the servlet developer wishes the application to function in a distributed environment.

The `proxyTimeout` element is equivalent to the `proxy-timeout` element of the deployment descriptor. The `proxyTimeout` should specify, in whole seconds, the default timeout for all proxy operations in this application. The container may override this value as a result of its own local policy.

The `sessionTimeout` element is equivalent to the `session-timeout` element in the deployment descriptor. The `sessionTimeout` should specify, in whole minutes, the default session timeout for all application sessions created in this servlet. `SipSessions` have no timeout independent of that of the containing application session. The lifetime of a `SipSession` is tied to that of the parent application session. If the timeout is zero or a negative number, the container must ensure the default behavior of the sessions is to never time out.

The `mainServlet` element of the annotation specifies which servlet is designated as the main servlet of the application. The concept of main servlet is described in [16 Mapping Requests To Servlets](#) of this specification. This corresponds to `main-servlet` element of the deployment descriptor.

```
package javax.servlet.sip.annotation;
import java.lang.annotation.Retention;
import java.lang.annotation.Target;
import static java.lang.annotation.ElementType.TYPE;
import static java.lang.annotation.RetentionPolicy.RUNTIME;
@Target ({PACKAGE})
@Retention (RUNTIME)
public @interface SipApplication {
    String name();
    String displayName() default "";
}
```

```

        String smallIcon() default "";
        String largeIcon() default "";
        String description() default "";
        boolean distributable() default false;
        int proxyTimeout() default 180; // seconds
        int sessionTimeout() default 3; // minutes
        String mainServlet() default "";
    }

```

To use the package level `@SipApplication` annotation, define it in a `package-info.java` file as follows:

```

@javax.servlet.sip.annotation.SipApplication(
    name="WeatherApplication",
    sessionTimeout=30,
    distributable=true)
package com.example;

```

The `WeatherService` servlet below does not define the `applicationName` element in the `@SipServlet` annotation:

```

package com.example;
import javax.servlet.sip.annotation.SipServlet;
import javax.servlet.sip.SipServletRequest;
import javax.servlet.sip.SipServletResponse;
    @SipServlet (name = "WeatherService")
public class Weather extends SipServlet {
    public void doInvite(SipServletRequest){
        //perform action
    }
}

```

The `HumidityChecker` servlet below defines the `applicationName` element in the `@SipServlet` annotation:

```

package com.example.humidity;
import javax.servlet.sip.annotation.SipServlet;
import javax.servlet.sip.SipServletRequest;
import javax.servlet.sip.SipServletResponse;
import javax.servlet.sip.SipServlet;
    @SipServlet (name = "HumidityChecker",
        applicationName="WeatherApplication")
public class Humidity extends SipServlet { . }
}

```

In the above example, the `WeatherService` servlet is automatically bound to the `WeatherApplication` application in its package, `com.example`. The `HumidityChecker` servlet

is in another package and has to specify the application name explicitly to be bound into the same application.

18.2.4 @SipListener Annotation

The @SipListener annotation allows the application developer to specify a listener without declaring it in the deployment descriptor of the application. The @SipListener annotation provides an alternative to the <listener> deployment descriptor element. The listener type is inferred from the interfaces implemented by the target class.

```
package javax.servlet.sip.annotation;

import static java.lang.annotation.ElementType.TYPE;
import java.lang.annotation.Inherited;
import java.lang.annotation.Retention;
import static java.lang.annotation.RetentionPolicy.RUNTIME;
import java.lang.annotation.Target;

@Target({TYPE})
@Retention(RUNTIME)
@Inherited
public @interface SipListener {
    String applicationName() default "";
    String description() default "";
}
```

The class annotated by the @SipListener must implement at least one of the listener interfaces described in this specification.

```
package com.example;
import javax.servlet.sip.annotation.SipListener;
import javax.servlet.sip.SipApplicationSessionListener;
import javax.servlet.sip.SipApplicationSessionEvent;

@SipListener (applicationName = "WeatherService")
public class MyApplicationSessionListener implements
SipApplicationSessionListener {
    public void sessionDestroyed(SipApplicationSessionEvent ev) {.}
    public void sessionCreated (SipApplicationSessionEvent ev){.}
    public void sessionExpired (SipApplicationSessionEvent ev){.}
}
```

18.2.5 @SipApplicationKey Annotation

The `@SipApplicationKey` annotation is used when the application wants to associate the incoming request (and `SipSession`) with a certain `SipApplicationSession`.

The method this annotation is attached to **MUST** be a public and static method, **MUST** return a `String` and **MUST** have a single argument of type `SipServletRequest`. If the method is not of this signature, the container **MUST** fail deployment of the application. The annotated method **MUST** not modify the `SipServletRequest` passed in. The `String` returned by the method is to be used as a key to associate the incoming request to an existing `SipApplicationSession` by the container as specified in [15.11.2 Session Key Based Targeting Mechanism](#). The container should treat a "null" return or an invalid session id as a failure to obtain a key from the application. It is recommended that the container create a new `SipApplicationSession` for the incoming request in such a case.

```
package javax.servlet.sip.annotation;
import java.lang.annotation.Inherited;
import java.lang.annotation.Retention;
import java.lang.annotation.Target;
import static java.lang.annotation.ElementType.TYPE;
import static java.lang.annotation.RetentionPolicy.RUNTIME;
@Target ({METHOD})
@Retention (RUNTIME)
@Inherited
public @interface SipApplicationKey {
    String applicationName() default "";
}
```

For example:

```
@javax.servlet.sip.annotation.SipApplication
package com.example;
import javax.servlet.sip.annotation.SipApplicationKey;
import javax.servlet.sip.SipServletRequest;
public class WeatherMapper {
    @SipApplicationKey
    public static String sessionKey(SipServletRequest req){
        return hash(req.getRequestURI +
                    getDomain(req.getFrom()));
    }
}
```

Note that only one `@SipApplicationKey` annotation can be present in a given SIP application (packaged as a SAR/WAR) and the container **MUST** enforce this restriction.

18.2.6 Annotation for SipFactory Injection

The `@Resource` annotation defined in the Common Annotations for Java Platform (JSR 250) is to be used for `SipFactory` injection.

When this annotation is applied to a field of type `SipFactory` in a converged application on a converged container, the container **MUST** inject an instance of `SipFactory` into the field when the application is initialized. This annotation can be used in place of the existing `ServletContext` lookup for the `SipFactory` from within a `Servlet`.

```
SipFactory sf =
    (SipFactory) getServletContext().getAttribute(SIP_FACTORY);
```

Usage above and

```
@Resource
private SipFactory sf;
```

are equivalent. For converged containers, the injected `SipFactory` would appear as `sip/<appname>/SipFactory` in the application scoped JNDI tree, where the `appname` is the name of the application as identified by the container (see [7.6 Application Names](#)).

In this example, this annotation can be used in any SIP `Servlet` or a Java EE application component deployed on the converged container in the same EAR archive file. The container must inject an instance of `SipFactory` into the field ("`sf`" in this example) at the time of application initialization. For the purposes of associating a `Servlet`, with, say the responses received to the request created using the `SipFactory`, the container must use the application's main `Servlet` as defined in [16.2 Servlet Selection](#). This can be changed to another `Servlet` through the `setHandler` mechanism in the `SipSession`.

Note that containers **MAY** allow the `@Resource` annotation to be present outside of SIP applications. In such cases, the `name` element of the `@Resource` annotation **MUST** identify the JNDI name of the desired factory to inject. Otherwise, an invocation of `SipFactory.createApplicationSession()` would be unable to determine the intended SIP application.

18.2.7 Annotation for SipSessionsUtil Injection

The `@Resource` annotation defined in the Common Annotations for Java Platform (JSR 250) is to be used to inject an instance of the `SipSessionsUtil` utility class for application session instance lookup.

This annotation can be used in place of the following `ServletContext` based lookup for the `SipSessionsUtil`.

```
SipSessionsUtil s =
    (SipSessionsUtil)
    getServletContext().getAttribute("javax.servlet.sip.SipSessionsUtil");
```

And the annotation based access -

```
@Resource
SipSessionsUtil s;
```

are equivalent. The injected `SipSessionsUtil` appears as `sip/<appname>/SipSessionsUtil` in the application scoped JNDI tree, where the `appname` is the name of the application as identified by the container (see [7.6 Application Names](#)). When multiple applications are packaged in the same EAR, access to `SipSessionsUtil` for a specific application is possible by specifying the JNDI `name` element of the `@Resource` annotation in a manner similar to [18.2.6 Annotation for SipFactory Injection](#). The `SipSessionsUtil` is described in [chapter 13 Converged Container and Applications](#).

18.2.8 Annotation for TimerService Injection

The `@Resource` annotation also can be used to inject an instance of the `TimerService` for scheduling timers.

This annotation can replace the following `ServletContext` based lookup of the `TimerService`.

```
TimerService t =
    (TimerService)
    getServletContext().getAttribute(TIMER_SERVICE);
```

And the annotation based access -

```
@Resource
TimerService t;
```

are equivalent. The injected `TimerService` appears as `sip/<appname>/TimerService` in the application scoped JNDI tree, where the `appname` is the name of the application as identified by the container (see [7.6 Application Names](#)). The `TimerService` is described in [9 Timer Service](#).

18.3 Annotation Parsing

The SIP servlet container parses Java class files in order to check for presence of the annotations described in the previous section. For .war and .sar files, the WEB-INF/classes as well as WEB-INF/lib directories are scanned. Containers may optionally process annotations for classes found elsewhere in the application's classpath.

19 Deployment Descriptor

This chapter specifies the SIP Servlet Specification, v1.1 requirements for SIP servlet container support of deployment descriptors. The deployment descriptor conveys the elements and configuration information of a servlet application between application developers, application assemblers, and deployers.

19.1 Differences from the HTTP Servlet Deployment Descriptor

The SIP servlet deployment descriptor is very similar to the HTTP servlet deployment descriptor. The main differences are:

- The root element is `sip-app` rather than `web-app`.
- Incoming requests are mapped to a defined main servlet rather than the HTTP specific URL mapping.
- The following elements (together with child elements) have no meaning for SIP and are not present in the deployment descriptor: `form-login-config`, `mime-mapping`, `welcome-file-list`, `error-page`, `taglib`, `jsp-file`.
- Resource collections are defined to be a named set of servlets rather than a set of URL patterns as in the case of HTTP.
- The notion of filters introduced in version 2.3 of the Servlet API is not supported. The application composition model of the SIP Servlet API allows multiple applications to

execute on a single request and, while this is not exactly the same as filters, it does meet many of the same requirements.

A SIP servlet application which does not use annotations ([18.2 Annotations and Resource Injection](#)) must have a deployment descriptor that conforms to the XML XSD. This deployment descriptor exists as a file `sip.xml` in the `/WEB-INF` directory of the SIP Servlet application.

19.2 Converged SIP and HTTP Applications

Applications may have both SIP and HTTP components (and potentially components related to protocols for other, as of yet unspecified, servlet APIs). When this is the case, there will be both a `web.xml` HTTP deployment descriptor conforming to the DTD defined by the HTTP Servlet API, and a `sip.xml` deployment descriptor.

A number of deployment descriptor elements apply to the application as a whole. The following rules specify how those elements are handled when they appear in both the SIP and HTTP descriptors (the rules apply equally to the case where there are deployment descriptors for other servlet APIs alongside SIP and/or HTTP):

- **distributable**: if this “tagging” element is present in one deployment descriptor, it must be present in the other.
- **context-param**: SIP and HTTP servlets belonging to the same application are presented with a single `ServletContext`. The set of initialization parameters available from the `ServletContext` is the union of parameters specified in `context-param` elements present in the SIP and HTTP deployment descriptors. A parameter may be configured in both descriptors but in that case must have the same value in both declarations.
- **listener**: the set of listeners of an application is the union of listeners defined in the SIP and HTTP deployment descriptors.

Also, the application `display-name` and icons should be identical and, if present, in one should be present in the other.

The `app-name` element is defined only in `sip.xml`. It is to be taken as the logical name for the SIP Servlet Application.

19.3 Deployment Descriptor Elements

The following types of configuration and deployment information exist in the SIP servlet application deployment descriptor and MUST be supported, with the exception of the security syntax, for all servlet containers:

- `ServletContext` `init` parameters
- Session configuration
- Servlet definitions
- Application lifecycle listener classes
- Error handler
- Security

Also, elements exist in the SIP servlet application deployment descriptor to support the additional requirements of servlet containers that are part of a Java EE application server. These elements allow for looking up JNDI objects (`env-entry`, `ejb-ref`, `ejb-local-ref`, `resource-ref`, `resource-env-ref`), and are not required to be supported by containers wishing to support only the servlet specification. See the schema comments for further description of these elements.

19.4 Rules for Processing the Deployment Descriptor

This chapter lists some general rules that servlet containers and developers must note concerning the processing of the deployment descriptor for a servlet application.

- Servlet containers should ignore all leading whitespace characters before the first non-whitespace character, and all trailing whitespace characters after the last non-whitespace character for PCDATA within text nodes of a deployment descriptor.
- Servlet containers and tools that manipulate servlet applications have a wide range of options for checking the validity of a SAR/WAR. This includes checking the validity of the deployment descriptor document held within. It is recommended, but not required, that servlet containers and tools validate deployment descriptors against the DTD document for structural correctness.

Additionally, it is recommended that they provide a level of semantic checking. For example, it should be checked that a role referenced in a security constraint has the same name as one of the security roles defined in the deployment descriptor.

In cases of non-conformant servlet applications, tools, and containers should inform the deployer with descriptive error messages. High-end application server vendors are encouraged to supply this kind of validity checking in the form of a tool separate from the container.

19.5 The SIP Servlet XSD

The XSD document [sip-app_1_1.xsd](#) describes the SIP Servlet deployment descriptor.

A Changes since v1.0

Changes in this specification since v1.0 in the order of appearance in the specification -

1. Servlet Life Cycle Listener - Added to remove any possibility of race conditions during initialization of the servlet. Also defined how servlet initialization relate to application deployment. [2.1.1 Servlet Life Cycle Listener](#)
2. Clear definition of Initial Request under various condition. Appendix [B Definition of Initial Request](#)
3. `Parameterable` interface - Support for SIP headers that can have parameters by providing easy accessors and mutators. [4.1.1 The Parameterable Interface](#)
4. Ability to add Contact header parameters and set user part. [4.1.3 The Contact Header Field](#)
5. Change to committed state of `SipServletMessage`. [5.2 Implicit Transaction State](#)
6. Because of evolution of servlet spec 2.4,
`SipServletResponse.setCharacterEncoding()` which extends `SipServletMessage` does not throw the `UnsupportedEncodingException`.
7. Specification for container to remove its own Route header and making the popped Route header available to applications. [5.6.3 Popped Route Header](#)
8. Changes in lifetime and accesibility of `SipServletMessage`. [5.8 Accessibility of SIP Servlet Messages](#)
9. The behavior of `SipServletMessage.getAcceptLanguage()` and `SipServletMessage.getAcceptLanguages()` has changed when no preferred locale is specified by the client. This is so that a caller of these methods can distinguish the case where

- no preferred locale is specified by the client from one where a preferred locale is specified. If no preferred locale is specified by the client, `getAcceptLanguage` returns null and `getAcceptLanguages` returns an empty `Iterator`. [5.9.1 Indicating Preferred Language](#)
10. New `getSession(Id)` method on `SipApplicationSession` to access the `SipSession` by its Id. [6.1.1 Protocol Sessions](#)
 11. Addition of a new mechanism for `SipApplicationSession` invalidation called Invalidation When Ready mechanism. [6.1.2.2 Invalidation When Ready Mechanism](#)
 12. Addition of a new mechanism for `SipSession` invalidation called Invalidation When Ready mechanism. [6.2.4.1.2 Invalidation When Ready Mechanism](#)
 13. Distributed containers to support not only `Serializable` objects but also instances of `SipServletMessage` as attributes. Also extent of serializable closure clarified. [6.4.3 Distributed Environments](#).
 14. Equivalence of .sar and .war archive formats for SIP Servlet Containers. [7.7 Servlet Application Archive File](#)
 15. New listeners for `SipApplicationSession` attributes. [8.1 SIP Servlet Event Types and Listener Interfaces](#)
 16. New listener for `SipApplicationSession` activation. [8.1 SIP Servlet Event Types and Listener Interfaces](#)
 17. New listener for `SipServlet` lifetime to receive initialization callback `servletInitialized()`. [8.1 SIP Servlet Event Types and Listener Interfaces](#)
 18. `SipSession` attribute listeners can be any class defined in the deployment descriptor. [8.1 SIP Servlet Event Types and Listener Interfaces](#)
 19. Support for RFC3327 (Path header) by extending the `Proxy` object. [10.5 Path Header and Path Parameters](#)
 20. Deprecation of transaction stateless proxy. [10.1 Parameters](#)
 21. Support for explicit `Proxy` branch creation and manipulation. [10.2.1 Proxy Branches](#)
 22. A general purpose `proxyTimeout` parameter applicable for both sequential and parallel proxy, deprecation of `sequential-search-timeout` parameter (not feature), optional override of `proxyTimeout` values for branches. [10.1 Parameters](#)
 23. Change in Sequential search timeout behavior. [10.2.4.3 Sequential Search Timeout](#)

24. Clarification that a 503 error response to be sent to the servlet if the sending of the message fails after the `send()` returns. [11.1.4 Sending a Request as a UAC](#)
25. Specification of Transaction timeout handling. [11.1.6 Transaction Timeout](#)
26. Removal of the stateless record routing section. Record routing applications by very nature are stateful.
27. Support for creating PRACK request by providing `createPrack()` method. [11.1.8 Sending PRACK](#) and [5.7.1 Reliable Provisional Responses](#)
28. Introduction of a B2BUA helper functionality in specification. The helper simplifies the writing of B2BUA applications by abstracting features like Session linkage, Request access on Sessions, Session cloning etc. [12 Back To Back User Agents](#)
29. SIP/HTTP and in general SIP/Java EE convergence. The features introduced in this specification bridge the gap between SIP Servlet applications and other Java EE components. [13 Converged Container and Applications](#)
30. A clear definition of a converged application scope and features available within that scope. [13.1 Converged Application](#)
31. Support for multihomed hosts. [14.2 Multihomed Host Support](#)
32. Detailed specification of Application selection and composition model. [15 Application Selection And Composition Model](#)
33. Specification of `encodeURI` mechanism in context of application composition. [15.11.1 Session Targeting and Application Selection](#)
34. The new Session Key based session targeting mechanism. [15.11.2 Session Key Based Targeting Mechanism](#)
35. Introduction of a designated "main" servlet in the application. [16.2 Servlet Selection](#)
36. Specification that only one servlet to be invoked in case of multiple servlets in mapping element. [16.2 Servlet Selection](#)
37. Modified Servlet triggering mechanism obviating the mapping rules. [16.2 Servlet Selection](#)
38. Mechanism for authentication of application initiated request. [17.10 Authentication of Servlet Initiated Requests](#)
39. Java 5 support. [18.1 Java 5](#)
40. Annotations for SIP Servlet container. [18.2 Annotations and Resource Injection](#)

41. New method for obtaining the application name from a `SipApplicationSession`. [7.6 Application Names](#)

A.1 Backward Compatibility considerations

While applications written to v1.0 of this specification (JSR 116) shall work without any change, there are a few changes that the container implementors must be aware of. In some extreme cases some changes may be required in the applications.

1. v1.0 style applications deployed on v1.1 or above container must still be invoked through the Application Router. Some v1.0 container implementations may have relied on servlet mapping rules to select applications, while servlet mapping rules shall still be consulted for servlet selection but the application selection solely resides with Application Router.
2. Contact header now allows user-part to be set and parameters to be added by the applications. Container implementations that used user-part or some named parameter for something should make appropriate adjustments.
3. Invalidate When Ready mechanism for session invalidation is on by default and will impact session lifetime as compared to v1.0. Please refer to [6.2.4.2 Important Semantics](#)
4. `SipApplicationSession.setExpires()` now does not throw an `IllegalArgumentException` on 0 or negative argument. Instead it is used to set the session to never expire, just as it was used to set through the deployment descriptor.
5. For v1.0 style applications with multiple servlets and their servlet-mapping rules one and only servlet selected [first matched going through the mapping] shall be invoked on initial request.
6. Because of evolution of servlet spec 2.4, `SipServletResponse.setCharacterEncoding()` which extends both `SipServletMessage` and `ServletResponse` now does not throw the `UnsupportedEncodingException` because it inherits a more generic method from `ServletResponse`. This is a binary compatible change, meaning that the compiled applications shall run on the v1.1 containers without change but it is not a source compatible change if the application is explicitly catching the `UnsupportedEncodingException` on `SipServletResponse.setCharacterEncoding()`
7. In absence of `SipServletRequest.getPoppedRoute()` API, some implementations of JSR 116 popped the Route header after the request visited the servlet. In these implementations, the applications may still have access to the Route header through existing methods such as `getHeader("Route")`. The applications using the old mechanism will have to change to use the new `SipServletRequest.getPoppedRoute()` and `SipServletRequest.getInitialPoppedRoute()` API.

A.2 Changes in the API since v1.0

Note: New APIs clearly marked in API docs with "@since 1.1"

- Overall the API has been updated to make use of Java 5, wherever possible the API has been generified and been made more type safe.
- Unchecked exceptions thrown from the API have been clearly specified
- All methods updated to indicate behavior on different/invalid inputs based on underlying domain objects
- `equals()` and `toString()` in various classes are updated to be inline with the domain object behavior

Table 19-1 New classes for application router support

<code>SipApplicationRouter</code>	application router interface
<code>SipApplicationRouterInfo</code>	application routing information
<code>SipApplicationRoutingDirective</code>	application routing directive
<code>SipApplicationRoutingRegion</code>	application routing region
<code>SipApplicationRoutingRegionType</code>	enum for routing region types
<code>SipRouteModifier</code>	enum used to influence route value

Table 19-2 New classes for application convergence

<code>SipSessionsUtil</code>	session management for converged apps
<code>ConvergedHttpSession</code>	extension to <code>HttpSession</code> for converged apps

Table 19-3 Misc new classes

<code>B2buaHelper</code>	support for B2BUA applications
<code>Parameterable</code>	represents SIP header field value with parameters
<code>ProxyBranch</code>	proxy branch information

Table 19-3 Misc new classes

<code>SipApplicationSessionAttributeListener</code>	get notifications of changes to the attribute lists of application sessions.
<code>SipApplicationSessionActivationListener</code>	new listener to support application session activation and passivation events
<code>SipApplicationSessionBindingEvent</code>	event for attribute binding/unbinding on <code>SipApplicationSession</code>
<code>SipApplicationSessionBindingListener</code>	listener for binding events
<code>SipServletContextEvent</code>	SIP Servlet specific context event.
<code>SipServletListener</code>	listener for <code>postInit()</code> event
<code>SipServletMessage.HeaderForm</code>	enum for header forms
<code>SipSession.State</code>	enum for session states
<code>UAMode</code>	enum for different UA modes, UAC or UAS.

Table 19-4 Changed or enhanced classes

<code>Address</code>	now implements <code>Parameterable</code>
<code>Proxy</code>	added new support for creating proxy branches
	added new support for setting outbound interface
	added new Path header support
	added Proxy timeout
	deprecated methods 'setStateful' and 'getStateful'
	deprecated sequential search attribute
<code>URI</code>	added new methods for setting and getting uri parameters
<code>SipURI</code>	moved parameter-related methods to <code>URI</code>
<code>TelURL</code>	now supports RFC 3966, also added new 'setPhoneNumber' method

Table 19-4 Changed or enhanced classes

SipSession	added new support for 'B2buaHelper' class
	added new 'isValid' method
	added method to query session state
	added new methods required for invalidation when ready mechanism
	added new method to support outbound interfaces
	added new method to support application composition
SipServletMessage	added new methods to support 'Parameterable' header types
	added ability to specify default use of compact or long header names in message
	added <code>pushRoute(Address uri)</code>
	<code>getAcceptLanguage</code> now returns null and <code>getAcceptLanguages</code> now returns an empty <code>Iterator</code> if no preferred locale was specified by the UA.
SipServletRequest	added support for new 'B2buaHelper' class
	added new method to retrieve previously popped route header
	added new methods to support new application composition
	added new method to support Path header
SipServletResponse	added new 'createPrack' method
	added new predefined response code constants
SipServlet	added new OUTBOUND_INTERFACES, SESSIONS context attributes
	added new methods: <ul style="list-style-type: none"> • <code>doPrack</code> - for SIP PRACK requests • <code>doUpdate</code> - for SIP UPDATE requests • <code>doRefer</code> - for SIP REFER requests • <code>doPublish</code> - for SIP PUBLISH requests

Table 19-4 Changed or enhanced classes

SipFactory	added new method to parse 'Parameterable' header types
SipApplicationSession	added new 'isValid' method
	added new 'encodeURL' method to support converged apps
	added new getExpirationTime()
	setExpires() now accepts 0 or negative argument
	getSessions(protocol) now throws IllegalArgumentException if the protocol is not understood
	added getSipSession(id) and getSession(String id, String protocol)
	deprecated the use of encodeURI() method
	added new methods required for invalidation when ready mechanism
ServletParseException	added support for nested exceptions
TooManyHopsException	added support for nested exceptions

B Definition of Initial Request

Since the determination that a request is an initial request determines when the application selection process is started, a well defined procedure for making this determination is necessary. This appendix provides a precise specification of what an initial request is.

The following procedure is used by a compliant container implementation to determine if a request is an initial request and will therefore require the application selection process to start:

1. Request Detection - Upon reception of a SIP message, determine if the message is a SIP request. If it is not a request, stop. The message is not an initial request.
2. Ongoing Transaction Detection - Employ methods of Section 17.2.3 in RFC 3261 to see if the request matches an existing transaction. If it does, stop. The request is not an initial request.
3. Examine Request Method. If it is CANCEL, BYE, PRACK, ACK, UPDATE or INFO, stop. The request is not an initial request for which application selection occurs.
4. Existing Dialog Detection - If the request has a tag in the To header field, the container computes the dialog identifier (as specified in section 12 of RFC 3261) corresponding to the request and compares it with existing dialogs. If it matches an existing dialog, stop. The request is not an initial request. The request is a subsequent request and must be routed to the application path associated with the existing dialog. If the request has a tag in the To header field, but the dialog identifier does not match any existing dialogs, the container must reject the request with a 481 (Call/Transaction Does Not Exist). Note: When this occurs, RFC 3261 says either the UAS has crashed or the request was misrouted. In the latter case, the misrouted request is best handled by rejecting the request. For the Sip Servlet environment, a UAS crash may mean either an application crashed or the container itself crashed. In either case, it is

impossible to route the request as a subsequent request and it is inappropriate to route it as an initial request. Therefore, the only viable approach is to reject the request.

5. Detection of Requests Sent to Encoded URIs - Requests may be sent to a container instance addressed to a URI obtained by calling the `encodeURI()` method of a `SipApplicationSession` managed by this container instance. When a container receives such a request, stop. This request is not an initial request. Refer to section [15.11.1 Session Targeting and Application Selection](#) for more information on how a request sent to an encoded URI is handled by the container.
6. Initial Request - The request is an initial request and the application invocation process commences. In this case and in this case only `SipServletRequest.isInitial()` MUST return true.

B.1 Retried Requests

The above procedure intentionally treats "retried" requests (as specified in Section 8.1.3.5 of RFC 3261) as initial requests. The fact that the CSeq value in such a request does not match the CSeq value in the original request ensures (even in the presence of race conditions where the original transaction is not yet gone) that the request would not be detected as a merged request. There are three good reasons for this approach:

1. Inconsistent Treatment of Retried Requests - Section 8.1.3.5 of RFC 3261 merely specifies that the UAC SHOULD use the same values for Call-ID, To, and From in the re-tried request as in the previous request. A UA could be noncompliant with RFC 3261 and use new values for Call-ID, To, and From headers. In this case, the request would surely be determined to be an initial request, and the application(s) that received the first request would not receive the retried request as a subsequent request.
2. Incorrect Application Handling of Retried Requests - Some of retried requests for 4XX responses in 8.1.3.5 could very well result in a set of applications being invoked that is different from the set invoked for the first request. Examples of 4xx responses where this is possible are: 415 (Unsupported Media Type) response, 416 (Unsupported URI Scheme), and 420 (Bad Extension). Treating a retried request as a subsequent request for these 4xx response codes would thus route the request to an application or servlet where either the application router would not have chosen the application and/or the triggering rule for the servlet does not match the retried request. This is undesirable.
3. Untractable Container Implementations - The determination of such retried requests requires that From tags, Call-IDs, and CSeq values for terminated transactions be nonetheless "remembered" for some period of time. The server transaction associated with a first request eventually reaches the Terminated state because the 4xx final response is sent to the UAC and

that response is subsequently acknowledged. Normally, when a transaction reaches the Terminated state, the container implementation would be free to release the resources associated with keeping track of the transaction. In fact, RFC 3261 dictates that the transaction be destroyed immediately upon the Terminated state being reached. However, because there is no specified time limit on how long a UAC may wait before sending a "retried" request, the container must keep such terminated transactions around. Although such a time limit could be configured, a retried request received after the time limit had expired would be treated as an initial request. In any case, keeping terminated transactions around would impose a significant burden on container implementations.

That retried requests are treated by the container as initial requests does unfortunately mean that some applications are going to be burdened with having to maintain the state required to detect such requests. This seems particularly true for applications that return the 401 (Unauthorized) or 407 (Proxy Authentication Required) expecting a retried request with the proper credentials included.

B.2 REGISTER Requests

REGISTER requests are also treated as initial requests, even when a REGISTER request refreshes an existing registration binding.

Definition of Initial Request

C Default Application Router

As an Application Router component is essential for the functioning of the container the following application router logic **SHOULD** be available with every container compliant with this specification. The container implementations **MAY** choose to provide a much richer Application Router component. For the purpose of this discussion the Application Router defined in this appendix is termed as Default Application Router (DAR).

The Application Router and the Container have a simple contract defined by the `SipApplicationRouter` interface.

The DAR **MUST** implement all the methods of that interface as described in this document.

C.1 The DAR Configuration File

The DAR works off a simple configuration text file which is modeled as a Java properties file:

- The properties file **MUST** be made available to the DAR and the location/content of this file **MUST** be accessible from a hierarchical URI which itself is to be supplied as a system property `"javax.servlet.sip.ar.dar.configuration"`
- The properties file has a simple format in which the name of the property is the SIP method and the value is a simple comma separated stringified value for the `SipApplicationRouterInfo` object.

```
INVITE: (sip-router-info-1), (sip-router-info-2)..  
SUBSCRIBE: (sip-router-info-3), (sip-router-info-4)..
```

- The properties file is first read by the DAR when the `init()` is first called on the DAR. The arguments passed in the `init()` are ignored.

- The properties file is refreshed each time `applicationDeployed()` or `applicationUndeployed()` is called. Similar to `init()`, the argument of these two invocations are ignored, these callbacks act just as a trigger to read the file afresh.

The sip-router-info data that goes in the properties file is a stringified version of the `SipApplicationRouterInfo` object. It consists of the following information :

- The name of the application as known to the container.
- The identity of the subscriber that the DAR returns. It can return any header in the SIP request using the DAR directive `DAR:SIP_HEADER` e.g `"DAR:From"` would return the SIP URI in From header. Or alternatively it can return any string.
- The routing region, one of the strings `"ORIGINATING"`, `"TERMINATING"` or `"NEUTRAL"`
- A SIP URI indicating the route as returned by the Application Router, it can be an empty string.
- A route modifier which can be any one of the strings `"ROUTE"`, `"ROUTE_BACK"` or `"NO_ROUTE"`
- A string representing `stateInfo`. As `stateInfo` is for Application Router's internal use only, what goes in this is up to the individual DAR implementations. As a hint the `stateInfo` could contain the index into the list of sip-router-info that was returned last.

Following is an example of the DAR configuration file:

```
INVITE: ("OriginatingCallWaiting", "DAR:From", "ORIGINATING", "",  
"NO_ROUTE", "0"), ("CallForwarding", "DAR:To", "TERMINATING", "",  
"NO_ROUTE", "1")
```

In this example, the DAR is setup to invoke two applications on INVITE request, one each in the originating and the terminating half. The applications are identified by their names as defined in the application deployment descriptors and used here. The subscriber identity returned in this case is the URI from the From and To header respectively for the two applications. The DAR does not return any route to the container and maintains the invocation state in the `stateInfo` as the index of the last application in the list.

C.2 The DAR Operation

The key interaction point between the Container and the Application Router is the method


```

SipApplicationRouterInfo getNextApplication
(SipServletRequest initialRequest,
 SipApplicationRoutingRegion region,
 SipApplicationRoutingDirective directive,
 SipTargetedRequestInfo targetedRequestInfo,
 Serializable stateInfo);

```

This method is invoked when an initial request is received by the container. When this method is invoked on DAR it will make use of the stateInfo and the initial request parameters and find out what SIP method is in the request. Next it will create the object `SipApplicationRouterInfo` from the sip-router-info information in the properties file, starting from the first in the list. The stateInfo could contain the index of the last sip-router-info returned so on next invocation of `getNextApplication` the DAR proceeds to the next sip-router- info in the list. The order of declaration of sip-router-info becomes the priority order of invocation.

As you would notice, this is a minimalist Application Router with no processing logic besides the declaration of the application order. It is expected that in real world deployments, the Application Router shall play an extremely important role in application orchestration and composition. It is likely to make use of complex rules and diverse data repositories. The DAR is intended to be a very simple implementation that is available as part of the reference implementation, and could be used instead of a real world Application Router.

Default Application Router

D References

- [3pcc] J. Rosenberg, J. Peterson, H. Schulzrinne and G. Camarillo, RFC 3725 - Best Current Practices for Third Party Call Control (3pcc) in the Session Initiation Protocol (SIP)
- [CPL] J. Lennox, X. Wu and H. Schulzrinne, RFC 3380 - CPL: A Language for User Control of Internet Telephony Services
- [IM] J. Rosenberg et al., RFC 3428 - Session Initiation Protocol (SIP) Extension for Instant Messaging.
- [JAF] B. Calder and B. Shannon, “JavaBeans Activation Framework Specification”, May 27, 1999.
- [JLS] “The Java Programming Language Specification”, <http://java.sun.com/docs/books/jls>.
- [JavaMail] Sun Microsystems, “JavaMail API Design Specification v1.2”, September 2000.
- [JAXP] R. Mordani, J. D. Davidson, and S. Boag, “Java API for XML Processing”, February 6, 2001.
- [privacy] C. Jennings, J. Peterson, and M. Watson, RFC 3325 - Private Extensions to the Session Initiation Protocol (SIP) for Asserted Identity within Trusted Networks.
- [refer] R. Sparks, RFC 3515 - The Session Initiation Protocol (SIP) Refer Method.
- [RFC 1738] T. Berners-Lee, L. Masinter, and M. McCahill, “Uniform Resource Locators (URL)”, RFC 1738, December 1994.
- [RFC 2278] N. Freed and J. Postel, “IANA Charset Registration Procedures”, RFC 2278, January 1998.
- [RFC 2327] M. Handley and V. Jacobson, “SDP: Session Description Protocol”, RFC 2327, April 1998.

References

- [3pcc] J. Rosenberg, J. Peterson, H. Schulzrinne and G. Camarillo, RFC 3725 - Best Current Practices for Third Party Call Control (3pcc) in the Session Initiation Protocol (SIP)
- [RFC 2976] S. Donovan, "The SIP INFO Method", RFC 2976, October 2000.
- [RFC 3261] J. Rosenberg, H. Schulzrinne, G. Camarillo, A. Johnston, J. Peterson, R. Sparks, M. Handley, and E. Schooler, "SIP: Session Initiation Protocol", RFC 3261, June 2002.
- [RFC 3262] J. Rosenberg and H. Schulzrinne, "Reliability of Provisional Responses in the Session Initiation Protocol (SIP)", RFC 3262, June 2002.
- [RFC 3265] A. B. Roach, "Session Initiation Protocol (SIP)-Specific Event Notification", RFC 3265, June 2002.
- [SERL] R. W. Steinfeldt and H. Smith, "SIP Service Execution Rule Language Framework and Requirements", Internet Engineering Task Force, May 2001. Work in progress.
- [Servlet API] D. Coward, "Java Servlet Specification, Version 2.3", September, 2001.
- [simple] Rosenberg et al., RFC 3856 - A Presence Event Package for the Session Initiation Protocol (SIP)
- [timer] S. Donovan and J. Rosenberg, RFC 4028 - Session Timers in the Session Initiation Protocol (SIP)
- [DFC1998] "Distributed feature composition: A virtual architecture for telecommunications services" Michael Jackson and Pamela Zave. IEEE Transactions on Software Engineering XXIV(10):831-847, October 1998.
- [CN2004] "Component coordination: A telecommunication case study" Pamela Zave and Healfdene H. Goguen and Thomas M. Smith. Computer Networks XXXV(5):645-664, August 2004.
- [SE2005] "Experience with component-based development of a telecommunication service" Gregory W. Bond and Eric Cheung and Healfdene H. Goguen and Karrie J. Hanson and Don Henderson and Gerald M. Karam and K. Hal Purdy and Thomas M. Smith and Pamela Zave. Proceedings of the Eighth International Symposium on Component-Based Software Engineering, May 2005.
- [JSR250] R. Mordani, Common Annotations for the Java™ Platform
- [RFC3327] D. Willis and B. Hoeneisen, Session Initiation Protocol (SIP) Extension Header Field for Registering Non-Adjacent Contacts.
- [RFC3966] H. Schulzrinne, The tel URI for Telephone Numbers.

- [3pcc] J. Rosenberg, J. Peterson, H. Schulzrinne and G. Camarillo, RFC 3725 - Best Current Practices for Third Party Call Control (3pcc) in the Session Initiation Protocol (SIP)
- [sipping-config
-framework] D. Petrie, A Framework for Session Initiation Protocol User Agent Profile Delivery.
- [RFC 3326] H. Schulzrinne, D. Oran and G. Camarillo, The Reason Header Field for the Session Initiation Protocol (SIP).
- [SIP Servlet
API] A. Kristensen, "SIP Servlet API, Version 1.0", February, 2003.

References

E Glossary

Address-of-Record

An address-of-record (AOR) is a SIP or SIPS URI that points to a domain with a location service that can map the URI to another URI where the user might be available. Typically, the location service is populated through registrations. An AOR is frequently thought of as the “public address” of the user.

Application developer

The producer of a SIP application. The output of an application developer is a set of servlet classes and supporting libraries and files (such as images, compressed archive files, etc.) for the servlet application. The application developer is typically an application domain expert. The developer is required to be aware of the servlet environment and its consequences when programming, including concurrency considerations, and create the SIP application accordingly.

Application assembler

Takes the output of the application developer and ensures that it is a deployable unit. Thus, the input of the application assembler is the servlet classes, JSP pages, HTML pages, and other supporting libraries and files for the servlet application. The output of the application assembler is a servlet application archive or a servlet application in an open directory structure.

Application path

A list of application instances to be invoked for incoming messages belonging to a particular dialog. This list is constructed as a side effect of the handling of the initial request and consists of the application acting as a UAC if the initial request was created within the container, the application acting as UAS if an application responded to the initial request, as well as all

applications that proxied with record-routing enabled. The application path is a logical concept and may or may not be explicitly represented by implementations.

Application Router

Application Router is a component which is essential for container's functioning. The contract between the Container and Application Router is an interface defined in this specification. On initial requests the Application Router is consulted by the Container to find which application to invoke, after any application completes processing of the request the Container again consults the Application Router to find which application to invoke next. The Application Router is thus the central place for determination of application invocation order.

Application session

Represents application instances. Application sessions acts as a store for application data and provides access to contained protocol sessions.

Back-To-Back User Agent

A back-to-back user agent (B2BUA) is a logical entity that receives a request and processes it as an user agent server (UAS). In order to determine how the request should be answered, it acts as an user agent client (UAC) and generates requests. Unlike a proxy server, it maintains dialog state and must participate in all requests sent on the dialogs it has established. Since it is a concatenation of a UAC and UAS, no explicit definitions are needed for its behavior.

Call

A call is an informal term that refers to some communication between peers generally set up for the purposes of a multimedia conversation.

Call leg

Another name for a dialog [RFC 2543]; not used in the revised SIP specification [RFC 3261].

Call Stateful

A proxy is call stateful if it retains state for a dialog from the initiating INVITE to the terminating BYE request. A call stateful proxy is always transaction stateful, but the converse is not necessarily true.

Client

A client is any network element that sends SIP requests and receives SIP responses. Clients may or may not interact directly with a human user. User agent clients and proxies are clients.

Committed message

A SIP message object that has been fully processed according to this specification and that should not be further modified, see [5.2 Implicit Transaction State](#).

Conference

A multimedia session that contains multiple participants.

Default servlet

The first servlet listed in the deployment descriptor.

Deployer

The deployer takes one or more servlet application archive files or other directory structures provided by an application developer and deploys the application into a specific operational environment. The operational environment includes a specific servlet container and SIP server. The deployer must resolve all the external dependencies declared by the developer. To perform his role, the deployer uses tools provided by the servlet container provider. The deployer is an expert in a specific operational environment. For example, the deployer is responsible for mapping the security roles defined by the application developer to the user groups and accounts that exist in the operational environment where the servlet application is deployed.

Dialog

A dialog is a peer-to-peer SIP relationship between two UAs that persists for some time. A dialog is established by SIP messages, such as a 2xx response to an INVITE request. A dialog is identified by a call identifier, local tag, and a remote tag. A dialog was formerly known as a call leg in RFC 2543. The baseline SIP specification defines only INVITE–BYE as a mechanism of establishing and terminating dialogs but allows extensions to define other methods capable of initiating dialogs. The SUBSCRIBE/NOTIFY methods defined by the event framework is an example of this [RFC 3265].

Downstream

A direction of message forwarding within a transaction that refers to the direction that requests flow from the user agent client to user agent server.

Final response

A response that terminates a SIP transaction, as opposed to a provisional response that does not. All 2xx, 3xx, 4xx, 5xx and 6xx responses are final.

Header

A header is a component of a SIP message that conveys information about the message. It is structured as a sequence of header fields.

Header field

A header field is a component of the SIP message header. It consists of one or more header field values separated by comma or having the same header field name.

Header field value:

A header field value is a singular value, which can be one of many in a header field.

Home Domain

The domain providing service to a SIP user. Typically, this is the domain present in the URI in the address-of-record of a registration.

Informational Response

Same as a provisional response.

Initial request

A request that is dispatched to applications based on rule matching rather than on an existing application path. Compare with subsequent request.

Initiator, calling party, caller

The party initiating a session (and dialog) with an INVITE request. A caller retains this role from the time it sends the initial INVITE that established a dialog until the termination of that dialog.

Invitation

An INVITE request.

Invitee, invited user, called party, callee

The party that receives an INVITE request for the purposes of establishing a new session. A callee retains this role from the time it receives the INVITE until the termination of the dialog established by that INVITE.

Location server

See Location service.

Location service

A location service is used by a SIP redirect or proxy server to obtain information about a callee's possible location(s). It contains a list of bindings of address-of-record keys to zero or more

contact addresses. The bindings can be created and removed in many ways; this specification defines a REGISTER method that updates the bindings.

Loose Routing

A proxy is said to be loose routing if it follows the procedures defined in this specification for processing of the Route header field. These procedures separate the destination of the request (present in the Request-URI) from the set of proxies that need to be visited along the way (present in the Route header field). A proxy compliant to these mechanisms is also known as a loose router.

Message

Data sent between SIP elements as part of the protocol. SIP messages are either requests or responses.

Method

The method is the primary function that a request is meant to invoke on a server. The method is carried in the request message itself. Example methods are INVITE and BYE.

Outbound proxy

A proxy that receives requests from a client, even though it may not be the server resolved by the Request-URI. Typically, a UA is manually configured with an outbound proxy, or can learn about one through auto-configuration protocols.

Parallel search

In a parallel search, a proxy issues several requests to possible user locations upon receiving an incoming request. Rather than issuing one request and then waiting for the final response before issuing the next request as in a sequential search, a parallel search issues requests without waiting for the result of previous requests.

Principal

A principal is an entity that can be authenticated by an authentication protocol. A principal is identified by a principal name and authenticated by using authentication data. The content and format of the principal name and the authentication data depend on the authentication protocol.

Protocol session

Common name for protocol specific session objects. A protocol session represents a point-to-point signaling relationship and serves as a repository for application data relating to that relationship. Examples include the `SipSession` and `HttpSession` interfaces defined by the SIP

and HTTP servlet specifications, respectively. A number of protocol sessions may belong to a single application session.

Provisional response

A response used by the server to indicate progress, but that does not terminate a SIP transaction. 1xx responses are provisional, other responses are considered final. Provisional responses are not sent reliably.

Proxy, proxy server

An intermediary entity that acts as both a server and a client for the purpose of making requests on behalf of other clients. A proxy server primarily plays the role of routing, which means its job is to ensure that a request is sent to another entity "closer" to the targeted user. Proxies are also useful for enforcing policy (for example, making sure a user is allowed to make a call). A proxy interprets, and, if necessary, rewrites specific parts of a request message before forwarding it.

Recursion

A client recurses on a 3xx response when it generates a new request to one or more of the URIs in the Contact header field in the response.

Redirect server

A redirect server is a user agent server that generates 3xx responses to requests it receives, directing the client to contact an alternate set of URIs.

Registrar

A registrar is a server that accepts REGISTER requests and places the information it receives in those requests into the location service for the domain it handles.

Regular transaction

A regular transaction is any transaction with a method other than INVITE, ACK, or CANCEL.

Request

A SIP message sent from a client to a server, for the purpose of invoking a particular operation.

Response

A SIP message sent from a server to a client, for indicating the status of a request sent from the client to the server.

Ringback

Ringback is the signaling tone produced by the calling party's application indicating that a called party is being alerted (ringing).

Role (development)

The actions and responsibilities taken by various parties during the development, deployment, and running of a servlet application. In some scenarios, a single party may perform several roles; in others, each role may be performed by a different party.

Role (security)

An abstract notion used by an application developer in an application that can be mapped by the Deployer to a user, or group of users, in a security policy domain.

Route set

A route set is a collection of ordered SIP or SIPS URI which represent a list of proxies that must be traversed when sending a particular request. A route set can be learned, through headers like Record-Route, or it can be configured.

Security policy domain

The scope over which security policies are defined and enforced by a security administrator of the security service. A security policy domain is also sometimes referred to as a realm.

Security technology domain

The scope over which the same security mechanism, such as Kerberos, is used to enforce a security policy. Multiple security policy domains can exist within a single technology domain.

Server

A server is a network element that receives requests in order to service them and sends back responses to those requests. Examples of servers are proxies, user agent servers, redirect servers, and registrars.

Sequential Search

In a sequential search, a proxy server attempts each contact address in sequence, proceeding to the next one only after the previous has generated a final response or has timed out. A 2xx or 6xx class final response always terminates a sequential search.

Servlet application archive

A single file that contains all of the components of a servlet application. This archive file is created by using standard JAR tools which allow any or all of the application components to be signed. Servlet application archive files are identified by the .sar extension. The SAR file layout is derived from the Web application archive (.war) file format but may contain servlets and deployment descriptors pertaining to different protocols, for example SIP and HTTP. With this specification either .sar or .war file formats can be used to package SIP Servlet applications.

Servlet Container Provider

A vendor that provides the runtime environment, namely the servlet container and possibly the SIP server, in which a servlet application runs as well as the tools necessary to deploy servlet applications. The expertise of the container provider is in HTTP-level programming. Since this specification does not specify the interface between the SIP server and the servlet container, it is left to the container provider to split the implementation of the required functionality between the container and the server.

Servlet definition

A unique name associated with a fully qualified class name of a class implementing the `Servlet` interface. A set of initialization parameters can be associated with a servlet definition.

Servlet mapping

A servlet definition that is associated by a servlet container with a URL path pattern. All requests to that path pattern are handled by the servlet associated with the servlet definition.

Session

From the SDP specification: “A multimedia session is a set of multimedia senders and receivers and the data streams flowing from senders to receivers. A multimedia conference is an example of a multimedia session.” [RFC 2327] (A session as defined for SDP can comprise one or more RTP sessions.) As defined, a callee can be invited several times, by different calls, to the same session. If SDP is used, a session is defined by the concatenation of the user name, session id, network type, address type and address elements in the origin field.

SIP application

A collection of SIP servlets and static resources which might include voice prompts, grammars, VoiceXML scripts, and other data. A SIP application may be packaged into an archive or exist in an open directory structure. All compatible servlet containers must accept a SIP application and perform a deployment of its contents into their runtime. This may mean that a container can run the application directly from a servlet application archive file or it may mean that it will move the contents of a servlet application into the appropriate locations for that particular container. SIP applications are to the SIP Servlet API what web applications are to the HTTP Servlet API.

SIP transaction

A SIP transaction occurs between a client and a server and comprises all messages from the first request sent from the client to the server up to a final (non-1xx) response sent from the server to the client. If the request is INVITE and the final response is a non-2xx, the transaction also

includes an ACK to the response. The ACK for a 2xx response to an INVITE request is a separate transaction.

SIP/web application, distributable

A SIP or web application that is written so that it can be deployed in a servlet container distributed across multiple Java virtual machines running on the same host or different hosts. The deployment descriptor for such an application uses the distributable element.

Stateful Proxy

A logical entity that maintains the client and server transaction state machines defined by this specification during the processing of a request. Also known as a transaction stateful proxy. The behavior of a stateful proxy is further defined in [10 Proxying](#). A (transaction) stateful proxy is not the same as a call stateful proxy.

Stateless Proxy

A logical entity that does not maintain the client or server transaction state machines defined in this specification when it processes requests. A stateless proxy forwards every request it receives downstream and every response it receives upstream. This specification deprecates the use of SIP Servlet application as transaction stateless proxy.

Strict routing

A proxy is said to be strict routing if it follows the Route processing rules of RFC 2543 and many prior Internet Draft versions of RFC 3261. That rule caused proxies to destroy the contents of the Request-URI when a Route header field was present. Strict routing behavior is not used in RFC 3261, in favor of a loose routing behavior. Proxies that perform strict routing are also known as strict routers.

Subsequent request

A request that is dispatched to applications, not by matching it against rules, but based on an existing application path created while processing an earlier initial request establishing the dialog.

System Administrator

The person responsible for the configuration and administration of the servlet container. The administrator is also responsible for overseeing the well-being of the deployed applications at run time. This specification does not define the contracts for system management and administration. The administrator typically uses runtime monitoring and management tools provided by the container provider and server vendors to accomplish these tasks.

System headers

Headers that are managed by the SIP servlet container and which servlets must not attempt to modify directly via calls to `setHeader` or `addHeader`. This includes Call-ID, From, To, CSeq, Via, Record-Route, Route, as well as Contact when used to specify a session signaling address, for example, in INVITEs and 200 response to INVITEs. System headers are discussed in section [5.4.2 System Headers](#).

TLS

Transport Layer Security. A layer four means of protecting TCP connections providing integrity, confidentiality, replay protection, and authentication.

Uniform resource locator (URL)

A compact string representation of information for location and access of resources via the Internet [RFC 1738]. SIP and SIPS URIs are syntactically similar to mailto URLs, that is, they are typically of the form `sip:user@host`. The user part is a user name or a telephone number. The host part is either a fully qualified domain name or a numeric IP address. SIP URIs are used within SIP messages to indicate the originator (From), current destination (Request-URI) and final recipient (To) of a SIP request, and to specify redirection addresses (Contact). When used as a hyperlink (for example in a Web page) a SIP URI indicates that the specified user or service can be contacted using SIP.

Upstream

The direction of message forwarding within a transaction that refers to the direction that responses flow from the user agent server back to the user agent client.

URL-encoded

A character string encoded according to RFC 1738 [RFC 1738, section 2.2].

User agent client (UAC)

A user agent client is a logical entity that creates a new request, and then uses the client transaction state machinery to send it. The role of UAC lasts only for the duration of that transaction. In other words, if a piece of software initiates a request, it acts as a UAC for the duration of that transaction. If it receives a request later, it assumes the role of a user agent server for the processing of that transaction.

User agent server (UAS)

A user agent server is a logical entity that generates a response to a SIP request. The response accepts, rejects, or redirects the request. This role lasts only for the duration of that transaction. In other words, if a piece of software responds to a request, it acts as a UAS for the duration of

that transaction. If it generates a request later, it assumes the role of a user agent client for the processing of that transaction.

User agent (UA)

A logical entity that can act as both a user agent client and user agent server.

Web application

A collection of servlets, JSP pages, HTML documents, and other web resources which might include image files, compressed archives, and other data. A web application may be packaged into an archive or exist in an open directory structure. All compatible servlet containers must accept a web application and perform a deployment of its contents into their runtime. This may mean that a container can run the application directly from a web application archive file or it may mean that it will move the contents of a web application into the appropriate locations for that particular container.

Glossary